

THE GLOBAL AI GAZETTE

Worldwide AI Ecosystem Intelligence • Open Source • Research • Models • Tools • Sovereign AI

Saturday, September 12, 2026 • 19:37 UTC

EDITION NO. 2223

156 STORIES ANALYZED

■ LEAD STORY

ANOTHER researcher accuses OpenAI of training on conversations and then claiming a breakthrough

Velocity 96/100 • Impact 88/100 • Source: reddit.com

ANOTHER researcher accuses OpenAI of training on conversations and then claiming a breakthrough reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology development. Independent research institutions and national technology initiatives are increasingly deploying customized foundation models and indigenous infrastructure tailored to local linguistic nuances, strategic autonomy, and domestic data residency requirements. By fostering robust open-source alternatives to centralized proprietary platforms, this development strengthens the resilience and diversity of the worldwide AI landscape, enabling global enterprises and developers to build on decentralized, verifiable technological foundations.

INNOVATION: Model: Frontier Foundation Weights & Open Checkpoint
ACTION: Deployment: Checkpoints accessible on Hugging Face; compatible with standard runtimes.

■ SECONDARY ANCHOR

Deepseek V4.1 Flash Release Video [Made with Deepseek V4.1 Flash]

I like to benchmark new models that come out on motion videos. So here's a test I did for deepseek v4.1 flash. And I have to say flash has probably graduated from being a Luna class model to nearly an Opus class model with this release, at least with motion videos. Prev.

■ TOP FLASH BULLETINS

- [VEL 96] [guide to using reasoning_effort on deepseek v4.1 flash](#) — guide to using reasoning_effort on deepseek v4.1 flash marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation....
- [VEL 96] [Qwen3.8-27B-Uncensored-Genesis-V1-GGUF](#) — Model available Qwen3.8-27B-Uncensored-Genesis-V1-MTP-GGUF This model is a practical realisation of things described in this paper, but adapted by me for machine learning: I am trying...
- [VEL 88] [Deploying Qwen3.8-2.4T-A95B on Amazon SageMaker HyperPod with vLLM](#) — On August 12, 2026, Alibaba's Qwen team released Qwen3.8-2.4T-A95B. This is the first time a Qwen-Max-class model has been made available as open weights. With...
- [VEL 88] [LangChain AI Application Framework: langchain-openai==1.6.2](#) — LangChain is a framework for building agents and LLM-powered applications. It helps you chain together interoperable components and third-party integrations to simplify AI application development...
- [VEL 96] [DeepSeek V4.1 Flash is available in HuggingChat](#) — DeepSeek V4.1 Flash is available in HuggingChat marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The...

■ SECTION II: EXECUTIVE AI STRATEGIC BRIEFING

Strategic Horizons & Action Directives for Enterprise Leadership

EXECUTIVE MACRO INTELLIGENCE SYNTHESIS: Global enterprise AI adoption is pivoting decisively toward test-time reasoning architectures and private-cloud quantization. Technology leadership must actively balance proprietary frontier model APIs with sovereign, open-weight deployments to reduce token expenditure while strictly sandboxing autonomous agent tool-calling boundaries.

VECTOR	DEVELOPMENT	SOURCE	DIRECTIVE
Model Sourcing & Licensing	DeepSeek-V4.1-Flash surprised	reddit.com	Audit open-weights licensing vs proprietary APIs; evaluate DeepSeek pa
Compute & Infrastructure CapEx	Deepseek V4.1 Flash is 748B, not 552B	reddit.com	Review GPU cluster allocation and power envelopes; benchmark Deepseek
Agentic Autonomy & Governance	Deepseek v4.1 flash finally has engrams,	reddit.com	Implement deterministic sandboxes for Deepseek autonomous tool executi
Inference Latency & Quantization	DeepSeek V4-1 Flash is out	reddit.com	Benchmark KV-cache compression (FP8/INT4/GGUF) and modern inference en
Open-Source Supply Chain	DeepSeek V4.1 Flash: Stronger, Faster, M	reddit.com	Inspect upstream repository dependencies; audit tokenizer code, weight
Data Residency & Sovereignty	deepseek-ai/DeepSeek V4.1 Flash · Huggin	reddit.com	Verify compliance with sovereign AI frameworks and regional data resid

1. **DeepSeek-V4.1-Flash surprised:** Audit open-weights licensing vs proprietary APIs; evaluate DeepSeek parameter efficiency and commercial distribution terms.
2. **Deepseek V4.1 Flash is 748B, not 552B:** Review GPU cluster allocation and power envelopes; benchmark Deepseek hardware efficiency to optimize cost per token.
3. **Deepseek v4.1 flash finally has engrams, what do you expect from 4.1 pro?:** Implement deterministic sandboxes for Deepseek autonomous tool execution, strict rate limiting, and human-in-the-loop validation.
4. **DeepSeek V4-1 Flash is out:** Benchmark KV-cache compression (FP8/INT4/GGUF) and modern inference engines for DeepSeek against TTFT SLAs.
5. **DeepSeek V4.1 Flash: Stronger, Faster, More Accessible:** Inspect upstream repository dependencies; audit tokenizer code, weights provenance, and pinned runtime releases for DeepSeek.
6. **deepseek-ai/DeepSeek V4.1 Flash · Hugging Face:** Verify compliance with sovereign AI frameworks and regional data residency requirements for deepseek deployments.

SECTION III: TRENDING OPEN-SOURCE AI & GITHUB INNOVATIONS

Top GitHub Repositories, Developer Velocity & Architecture Dissections

PROJECT / REPO	SOURCE	STACK	VELOCITY	CORE ARCHITECTURAL FOCUS
LangChain AI Application Framework: langchain	github.com	Python	88/100	LangChain is a framework for building agents and LLM-powered applications. It helps you chain together interoperable components and.
DeepSeek reportedly advances Shanghai IPO pla	digitimes.com	Python	88/100	DeepSeek reportedly advances Shanghai IPO plans as it cuts Flash model API prices marks an architectural milestone in.
memgraph/memgraph: High-performance open-sour	github.com	Python	80/100	Memgraph is a high-performance, in-memory graph database that powers real-time AI context and graph analytics.
AutoGPT Agent Runtime: Preview seed fixture (	github.com	Python	80/100	Rolling, fully SYNTHETIC preview-database seed fixture baked from dev@dd71bda0df43 (migration head: 20260908121500_trial_rejection_reason). Preview environments restore fixture.dump.gz (custom-format, gzipped).
Three open source foundations share one stage	digitimes.com	Python	72/100	Three open source foundations share one stage in Shanghai reflects the rapid acceleration of sovereign artificial intelligence ecosystems.
Reducto Releases r-1: A Single Pass Document	marktechpost.com	Python	72/100	Last week, Reducto announced r-1. It is the first model in a new parsing family built on a.

[TRENDING REPO] SOURCE: [github.com](#) | VEL: 88/100

LangChain AI Application Framework: langchain-openai==1.6.2

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

LangChain is a framework for building agents and LLM-powered applications. It helps you chain together interoperable components and third-party integrations to simplify AI application development — all while future-proofing decisions as the underlying technology evolves. > [!TIP] > Just getting started? Check out **Deep Agents** — a higher-level package built on LangChain for agents that have built-in capabilities for common usage patterns such as planning, subagents, file system usage, and more. If you're looking for more advanced customization or agent orchestration, check out LangGraph, our framework for building controllable agent workflows.

[TRENDING REPO] SOURCE: [digitimes.com](#) | VEL: 88/100

DeepSeek reportedly advances Shanghai IPO plans as it cuts Flash model API prices

ARCH: MOE • CTX: 128K • QUANT: FP8 • STACK: vLLM

DeepSeek reportedly advances Shanghai IPO plans as it cuts Flash model API prices marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[TRENDING REPO] SOURCE: [github.com](#) | VEL: 80/100

memgraph/memgraph: High-performance open-source in-memory graph database for GraphRAG, AI memory, a...

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Memgraph is a high-performance, in-memory graph database that powers real-time AI context and graph analytics. Built in C/C++, it serves as the graph engine for GraphRAG pipelines, AI memory systems, and agentic workflows — delivering sub-millisecond multi-hop traversals for any system that needs structured, connected context alongside semantic vector or text search. Memgraph provides both in a single query layer: built-in text and vector indexes for similarity search combined with full graph traversal, so retrieval pipelines can run as a single atomic database operation instead of being scattered across multiple systems. The same architecture drives real-time graph analytics for fraud detection, network analysis, infrastructure monitoring, and other operational workloads where performance matters.

[TRENDING REPO] SOURCE: [github.com](#) | VEL: 80/100

AutoGPT Agent Runtime: Preview seed fixture (rolling)

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Rolling, fully SYNTHETIC preview-database seed fixture baked from dev@dd71bda0df43 (migration head: 20260908121500_trial_rejection_reason). Preview environments restore fixture.dump.gz (custom-format, gzipped, platform schema only, including _prisma_migrations) BEFORE running a PR's own migrations, so schema changes run against populated tables. RESTORE: apply restore-preamble.sql to the target database first (it creates the platform schema and installs the vector + pg_trgm extensions into it.

[TRENDING REPO] SOURCE: [digitimes.com](#) | VEL: 72/100

Three open source foundations share one stage in Shanghai

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Three open source foundations share one stage in Shanghai reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology development. Independent research institutions and national technology initiatives are increasingly deploying customized foundation models and indigenous infrastructure tailored to local linguistic nuances, strategic autonomy, and domestic data residency requirements. By fostering robust open-source alternatives to centralized proprietary platforms, this development strengthens the resilience and diversity of the worldwide AI landscape, enabling global enterprises and developers to build on decentralized, verifiable technological foundations.

SECTION IV: FRONTIER FOUNDATION MODELS & REASONING
BREAKTHROUGHS

Reasoning Architectures, Open Weights & Capability Benchmarks

FOUNDATION MODEL	SOURCE	SIZE / CTX	IMPACT	ARCHITECTURAL HIGHLIGHT
guide to using reasoning_effort on deepseek v	reddit.com	MOE / 128K	88/100	guide to using reasoning_effort on deepseek v4.1 flash marks an architectural milestone in Open-Weights foundation modeling, engineered for.
Qwen3.8-27B-Uncensored-Genesis-V1-GGUF	reddit.com	27B / 128K	88/100	Model available Qwen3.8-27B-Uncensored-Genesis-V1-MTP-GGUF This model is a practical realisation of things described in this paper, but adapted by.
DeepSeek V4.1 Flash is available in HuggingCh	reddit.com	MOE / 128K	68/100	DeepSeek V4.1 Flash is available in HuggingChat marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput.
DeepSeek-V4.1-Flash surprised	reddit.com	50B / 128K	68/100	Hoping to see smartest medium size models soon & later with all available optimizations/architectures/etc.,. Thanks Deepseek!
Deepseek V4.1 Flash is 748B, not 552B	reddit.com	748B / 128K	68/100	People keep on getting confused about this, so I looked at the safetensors on hf.
Deepseek v4.1 flash finally has engrams, what	reddit.com	MoE / SOTA / 128K	68/100	If the ratio is the same, Maybe 1.6T -3.1T params plus .56T-1.06T engrams and fable 5.0 level performance?.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 88/100
guide to using reasoning_effort on deepseek v4.1 flash
ARCH: MOE • CTX: 128K • QUANT: FP8 • STACK: vLLM

guide to using reasoning_effort on deepseek v4.1 flash marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 88/100
Qwen3.8-27B-Uncensored-Genesis-V1-GGUF
ARCH: 27B • CTX: 128K • QUANT: GGUF • STACK: PyTorch / ONNX

Model available Qwen3.8-27B-Uncensored-Genesis-V1-MTP-GGUF This model is a practical realisation of things described in this paper, but adapted by me for machine learning: I am trying to solve the problem: why LLM models even for simple questions write walls of text during reasoning, and burn too much tokens instead of solving the task. And when number of parameters increase the problem became worse.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 68/100
DeepSeek V4.1 Flash is available in HuggingChat
ARCH: MOE • CTX: 128K • QUANT: FP8 • STACK: vLLM

DeepSeek V4.1 Flash is available in HuggingChat marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 68/100
DeepSeek-V4.1-Flash surprised
ARCH: 50B • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Hoping to see smartest medium size models soon & later with all available optimizations/architectures/etc.,. Thanks Deepseek! Ex 1: 30-50B MOE + 10-15B Engram + DeepSeek-V4.1-Flash type KVCache Ex 2: 15-30B Dense + 10-15B Engram + DeepSeek-V4.1-Flash type KVCache EDIT: Updated Engram to 10-15B from.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 68/100
Deepseek V4.1 Flash is 748B, not 552B
ARCH: 748B • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

People keep on getting confused about this, so I looked at the safetensors on hf. The title should have been "Deepseek V4.1 Flash is 748B total/552B base, not 284B or 305B or 485B or 522B" The model is not 284B. The original Deepseek V4 Flash is 284B, but not the V4.1 Flash model The model is not 305B, despite what some people claim "So: ~305B real backbone + 203B engram = 508B total" This is incorrect.

■ SECTION V: TOP AI RESEARCH PAPERS & ARXIV BREAKTHROUGHS

Reasoning Paradigms, Multimodal Architectures & Algorithmic Design

[AI RESEARCH] SOURCE: siliconangle.com | VEL: 88/100

Anthropic researcher's resignation sparks broad AI safety discussion

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

An Anthropic PBC researcher has resigned over concerns that artificial intelligence labs are “gambling with our lives.” Jacob Coxon was part of the company’s AI pretraining team until today. He announced his resignation in a series of X posts that has been viewed millions of times. Coxon wrote that his decision was motivated by concerns over recursive self-improvement, or RSI. That’s a term for a hypothetical future AI with the ability to automatically improve its own capabilities. “These will soon be superhuman systems that can hack anything, revolutionize any field overnight, and acquire real power and resources,” Coxon wrote.

[AI RESEARCH] SOURCE: arxiv.org | VEL: 88/100

RedKnot-MLA: Multi-Head Offline-Online Reuse for DeepSeek-V4 Long-Context Serving

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Multi-head latent attention (MLA) exposes many logical query heads through one packed latent KV stream. This representation is memory efficient, but it removes the physical per-head cache boundary assumed by conventional head-wise reuse. We present our system, a DeepSeek-V4 realization of RedKnot's head-aware reuse principle. Each immutable document is processed offline at canonical position zero; certified Local-head contributions are retained as MLA-Off.

[AI RESEARCH] SOURCE: siliconangle.com | VEL: 88/100

OpenAI to set misalignment disclosure rules after agents took over a wiki

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

OpenAI Group PBC acknowledged Saturday that it did not publicly disclose an episode in which its artificial intelligence agents wrote to outside websites and said it will publish a framework in the coming weeks for reporting misaligned model behavior. The company now calls the episode the “wiki incident.” Researchers led by the Nightingale Collective set it out in a report published Friday after finding roughly 17,000 posts on DSEwiki, a long-dormant German software developer wiki that had taken about 20 edits in the previous decade. Sydney Von Arx, Cormac Slade Byrd, Spencer Kitts and Thomas Larsen rebuilt the deleted pages from the wiki’s edit history and published them at collusion.wiki.

[AI RESEARCH] SOURCE: arxiv.org | VEL: 72/100

JarvisGUI: Towards Cross-Device GUI Agents with Dynamic Task Composition

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Real-world GUI usage frequently involves workflows that span multiple devices and platforms, requiring the transfer of intermediate results, maintenance of shared state, and coordination across heterogeneous environments. However, existing GUI benchmarks overwhelmingly evaluate agents on single-device, statically defined tasks, thus leaving such cross-device capabilities largely unexamined, resulting in an overly optimistic assessment of agents' readiness for real-world usage.

[AI RESEARCH] SOURCE: news.mit.edu | VEL: 72/100

MIT Schwarzman College of Computing launches pilot to help educators teach AI across disciplines

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

This summer, the MIT Schwarzman College of Computing welcomed faculty from colleges and universities across Greater Boston, South Carolina, West Virginia, and Texas to campus for the inaugural AI Educators Pilot, a weeklong workshop aimed at expanding how artificial intelligence is taught across disciplines and learning environments.

[AI RESEARCH] SOURCE: arxiv.org | VEL: 72/100

Show-Harness: Just a VLM Agent Can Play Robots

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Foundation vision-language models (VLMs) exhibit broad intelligence about the world, yet translating this intelligence into robot control remains challenging. We present Show-Harness, an Embodied Harness that enables VLMs to “play” robots through a compact semantic interface linking intent to action.

■■ SECTION VI: DEVELOPER TOOLS, FRAMEWORKS & AI INFRASTRUCTURE

Inference Runtimes, Evaluation Frameworks, Vector DBs & Local Execution Engines

[AI TOOL] SOURCE: securityweek.com | VEL: 88/100

Anthropic Researcher Resigns With Warning About the Dangers of AI Development

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

His posts reached more than 100 million people overnight. Coxon is not the first AI insider to publicly raise such concerns. Both Anthropic and OpenAI have seen high-profile resignations in recent years that were tied to safety concerns. Two current Anthropic employees also responded to Coxon's post in agreement. Sen. Bernie Sanders, a Vermont independent who has called for AI safeguards and regulation, agreed with Coxon's concerns and said he would soon introduce legislation to pause AI development and ban superintelligence. "The very people building this technology admit that it could threaten the future of humanity," Sanders said Wednesday on social media.

[AI TOOL] SOURCE: aibusiness.com | VEL: 88/100

OpenAI Admits More AI Agents Went Astray in May

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

OpenAI Admits More AI Agents Went Astray in May represents an advanced leap in autonomous agent orchestration and embodied AI systems. Departing from passive query-response interfaces, the architecture integrates recursive planning, dynamic tool calling, and grounded environment feedback to execute complex, multi-turn objectives without human intervention. The system incorporates deterministic guardrails, structured memory persistence, and standardized communication protocols (such as Model Context Protocol) to ensure agent actions remain safe, auditable, and robust against cascading execution errors. This progress paves the way for reliable digital coworkers and autonomous physical robotics capable of operating across real-world workflows.

[AI TOOL] SOURCE: thehackernews.com | VEL: 88/100

Autonomous AI Agents Compromise Thousands of Credentials in Under Six Hours

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Threat actors are continuing to leverage artificial intelligence (AI) to streamline their operations, with one financially motivated hacking group employing an autonomous, multi-agent attack framework to carry out a large-scale credential harvesting campaign within six hours.

[AI TOOL] SOURCE: pandaily.com | VEL: 88/100

Alibaba Qwen Open-Sources Qwen-Drive-1.0-4B for Autonomous Driving Perception and Planning

ARCH: 4B • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Alibaba's Qwen team has open-sourced Qwen-Drive-1.0-4B, a vision-language model for autonomous driving that pairs scene understanding with motion planning. Developed with Huazhong University of Science and Technology, the release keeps the pretrained Qwen3.5-4B multimodal backbone unchanged and attaches modules for bird's-eye-view 3D perception and ego trajectory generation without rewriting the base vision-language architecture.

[AI TOOL] SOURCE: aws.amazon.com | VEL: 88/100

Take on your most ambitious work with GPT-6 Astra on Amazon Bedrock

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

GPT-6 Astra from OpenAI brings greater depth and judgment to your most demanding tasks and runs on the Amazon Bedrock inference engine built for high performance, security, and scale. Organizations are already running AI agents that write code, analyze data, and automate complex workflows at production scale on Amazon Bedrock. GPT-6 Astra raises the potential of what those agents can deliver.

[AI TOOL] SOURCE: reddit.com | VEL: 96/100

Closed AI doesn't like biological research, user turns to open weight models

ARCH: MOE • CTX: 128K • QUANT: FP8 • STACK: vLLM

Closed AI doesn't like biological research, user turns to open weight models marks an architectural milestone in Frontier API foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

■ SECTION VII: SOVEREIGN AI & WORLDWIDE REGIONAL ECOSYSTEMS

National AI Labs, Regional Models & Sovereign Compute Ecosystems

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

Deepseek Has Soft Retired Deepseek V4 Pro

ARCH: MOE • CTX: 128K • QUANT: FP8 • STACK: vLLM

Deepseek Has Soft Retired Deepseek V4 Pro marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

3060 12GB vs 4060 ti 16GB

ARCH: 27B • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

I'm currently building my system around 3060s, but I might be able to get a 4060 for a nice deal. At first it seemed like a no brainer, but turns out the 4060 has lower memory bandwidth. In a system that already has 4x 3060 12GBs set up on a threadripper with tensor parallelism (mostly qwen3.8-27b), would it be worth having the 4060 ti 16GB around for the extra 4GB and occasional gaming, or is it just going to slow the rest of the setup down for AI?

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

Harness doesn't matter

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Harness doesn't matter reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology development. Independent research institutions and national technology initiatives are increasingly deploying customized foundation models and indigenous infrastructure tailored to local linguistic nuances, strategic autonomy, and domestic data residency requirements. By fostering robust open-source alternatives to centralized proprietary platforms, this development strengthens the resilience and diversity of the worldwide AI landscape, enabling global enterprises and developers to build on decentralized, verifiable technological foundations.

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

I made a way to migrate between embedding models without re-embedding your entire corpus

ARCH: 1B • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

So I was playing with embedding models I saw that when you upgrade from model A to B, you face a very big backfilling cost ie, suppose you have a 1b vectors from model A, and then you want to use model B. This would mean you have to re-embed all of your documents with model B before you can even serve with the model, and on an H100, it would take ~108 days (qwen embed 8b, 106 docs/second). But I found an easier way to do it.

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

Would you consider 5t/s usable for a local model?

ARCH: 27B • CTX: 780M • QUANT: Native / FP16 • STACK: PyTorch / ONNX

I'm able to run qwen3.8 27b in two ways on my system: split between my 3060 12gb and 9070xt running at 20t/s or running off the 780m iGPU and 5400mhz DDR5 at 5t/s. Personally I feel like the 5t/s is still more usable because I have enough RAM to still use my system mostly normally. I can even play games just fine because the CPU side is left alone.

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

Mention if a "new model" is a finetune

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

A few posts tagged with "new model" present models that are finetunes. My opinion : I'd rather have the "new model" tag reserved for new "major" releases, like a new Qwen model, Deepseek V4 -> Deepseek V4.1, etc., that involved a new pretrain or intensive post-training (in opposition to a small finetune). Otherwise, maybe prepend "[Finetune]" to the title to indicate that the new model is "less of a big news", a use a "new finetune" tag, to differentiate between the two kinds of new models.

SECTION VIII: AI HARDWARE, COMPUTE CLUSTERS & SILICON

GPU Clusters, Custom Accelerators, Datacenter Infrastructure & Silicon

SILICON / SYSTEM	SOURCE	COMPUTE SPECS	VELOCITY	BENCHMARK / TELEMETRY
Running Vision Qwen 3.8 27B on a 16GB Card, t	reddit.com	High Velocity	96/100	I am just sharing my config for Qwen 3.8 27b that fits on a 5060TI, what is cool.
Qwen3.8-Flash-Next on 2x3090 + DDR4, part 4:	reddit.com	29 t/s	96/100	. 17 -> 25-29 t/s with the expert cache PR, 37-41 t/s after switching to UD-Q4_K_XL and stacking.
Qwen3.8-Flash-Next on MLX-serve, 1m context i	reddit.com	40 tok/s	96/100	Hi, I'm the co-creator of this Qwen3.8-Flash-Next engine support in MLX-serve. I've been tuning this one to run.
llama.cpp High-Performance LLM Engine: b10899	github.com	High Velocity	88/100	vulkan: small M matrix optimizations for qwen (#28457) vulkan: optimize m=1 mul_mat by swapping A/B vulkan: Improve small.
DeepSeek AI Released DeepSeek-V4.1-Flash with	marktechpost.com	High Velocity	88/100	Long-horizon agents have turned LLM serving into an input-heavy workload. Repeated prefills and million-token contexts leave KV caches.
OpenBMB Releases MiniCPM5-2B: A 2.52B Dense M	marktechpost.com	High Velocity	72/100	OpenBMB has released MiniCPM5-2B, the second checkpoint in the MiniCPM5 series and the follow-up to MiniCPM5-1B.

[AI HARDWARE] SOURCE: reddit.com | VEL: 96/100

Running Vision Qwen 3.8 27B on a 16GB Card, the config (45tks).

ARCH: 27B • CTX: 85K • QUANT: GGUF • STACK: PyTorch / ONNX

I am just sharing my config for Qwen 3.8 27b that fits on a 5060TI, what is cool about this is that you can even get vision! and a 85K context (I have 1.5gb of headroom for more context or a better quant) Model: IQ3_XXS-mtp from Using beellama Config used: [*] model = ..\llm-models\Qwen3.8-27B-GSQ-RCO-IQ3_XXS-mtp.gguf mmproj = ..\llm-models\mmproj-Qwen3.8-27B-BF16.gguf image-min-tokens = 256.

[AI HARDWARE] SOURCE: reddit.com | VEL: 96/100

Qwen3.8-Flash-Next on 2x3090 + DDR4, part 4: 2.2-2.5x faster prefill by kicking the expert cache off the GPU while the prompt runs

ARCH: MoE / SOTA • CTX: 8K • QUANT: UD-Q4_K_XL • STACK: llama.cpp

. 17 -> 25-29 t/s with the expert cache PR, 37-41 t/s after switching to UD-Q4_K_XL and stacking MTP on the cache, the top-k fallback that was sorting more than it needed to. This one is all about prefill, which was honestly the weak spot the whole time. 80+ seconds before the first token on an 8k prompt, and 24 minutes on a 119k one...I know . Box is still 2x 3090, dual Broadwell Xeon, llama.cpp, UD-Q4_K_XL with the Q8 MTP head on the.

[AI HARDWARE] SOURCE: reddit.com | VEL: 96/100

Qwen3.8-Flash-Next on MLX-serve, 1m context is released!

ARCH: MoE / SOTA • CTX: 1M • QUANT: Native / FP16 • STACK: MLX

Hi, I'm the co-creator of this Qwen3.8-Flash-Next engine support in MLX-serve. I've been tuning this one to run both fast, efficient and correct up 1m context using kv cache 8 bits in M5 Max 128GB. Qwen is working well at very long context as showed in the video (a snapshot at ~760k context), I let it build MLX Serve Monitor plugin that you've seen on the right side of Opencode2's app. The generation sustain through 1m context at around 40 tok/s on prose and 75 tok/s on coding.

[AI HARDWARE] SOURCE: github.com | VEL: 88/100

llama.cpp High-Performance LLM Engine: b10899

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: llama.cpp

vulkan: small M matrix optimizations for qwen (#28457) vulkan: optimize m=1 mul_mat by swapping A/B vulkan: Improve small M perf Allow split_k with small M. Make small vs med tile selection (for coopmat2) depend on M, not just N.

[AI HARDWARE] SOURCE: marktechpost.com | VEL: 88/100

DeepSeek AI Released DeepSeek-V4.1-Flash with 1M Context, FP4 KV Cache, and Cross-Layer Attention Reuse

ARCH: 552B • CTX: 1M • QUANT: FP4 • STACK: PyTorch / ONNX

Long-horizon agents have turned LLM serving into an input-heavy workload. Repeated prefills and million-token contexts leave KV caches that strain HBM, SSD capacity, and bandwidth. DeepSeek AI built its newest release around that exact bottleneck. DeepSeek-V4.1-Flash is a multimodal Mixture-of-Experts model with 552B backbone parameters, 196B additional Engram parameters, and a 1M-token context window. It activates 8B parameters per token during prefill and 16B during decode.

■ SECTION IX: AUTONOMOUS AGENTS, MULTI-AGENT SWARMS & ROBOTICS

Agent Orchestration, Computer-Use Primitives, Action Models & Robotics

AGENT / FRAMEWORK	SOURCE	PROTOCOL	VELOCITY	CORE CAPABILITY DOMAIN
Deploying Qwen3.8-2.4T-A95B on Amazon SageMak	aws.amazon.com	vLLM	88/100	On August 12, 2026, Alibaba's Qwen team released Qwen3.8-2.4T-A95B. This is the first time a Qwen-Max-class model has.
Research acceleration: The view inside OpenAI	simonwillison.ne	PyTorch / ONNX	96/100	Research acceleration: The view inside OpenAI Apparently today is RSI day at OpenAI, for Recursive Self-Improvement - I.
Scanning the Harness: An Empirical Study of S	arxiv.org	PyTorch / ONNX	88/100	AI coding agents such as Claude Code, Cursor, GitHub Copilot, and OpenAI Codex are configured through artifacts developers.
Muse, Meta's New Personal AI Agent, Needs You	wired.com	PyTorch / ONNX	88/100	Meta announced Tuesday the release of Muse, a personal AI agent that people can message to automate digital.
LangChain AI Application Framework: langchain	github.com	PyTorch / ONNX	88/100	LangChain is a framework for building agents and LLM-powered applications. It helps you chain together interoperable components and.
Anthropic Discloses Fourth AI Hacking Incident	thehackernews.co	PyTorch / ONNX	88/100	Anthropic on Wednesday disclosed a fourth incident in which its artificial intelligence (AI) model broke into real third-party.

[AGENT SYSTEM] SOURCE: [aws.amazon.com](#) | VEL: 88/100
Deploying Qwen3.8-2.4T-A95B on Amazon SageMaker HyperPod with vLLM
ARCH: MoE / SOTA • CTX: 262K • QUANT: Native / FP16 • STACK: vLLM

On August 12, 2026, Alibaba's Qwen team released Qwen3.8-2.4T-A95B. This is the first time a Qwen-Max-class model has been made available as open weights. With 2.4 trillion total parameters (95 billion activated per token), a hybrid linear-plus-full-attention architecture, and native context up to 262K tokens (extensible to 1M), Qwen3.8 targets the most demanding agentic and reasoning workloads. These include multi-step coding, long-horizon planning, and autonomous tool use.

[AGENT SYSTEM] SOURCE: [simonwillison.net](#) | VEL: 96/100
Research acceleration: The view inside OpenAI
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Research acceleration: The view inside OpenAI Apparently today is RSI day at OpenAI, for Recursive Self-Improvement - I think it's their new AGI. Both this piece and the new essay An Alien Mind (by Chief Scientist Jakub Pachocki) talk about it, and this one doesn't even bother to expand the acronym. Included are details on how OpenAI's own research team are using coding agents.

[AGENT SYSTEM] SOURCE: [arxiv.org](#) | VEL: 88/100
Scanning the Harness: An Empirical Study of Supply-Chain Defects in AI Coding-Agent Configurations
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

AI coding agents such as Claude Code, Cursor, GitHub Copilot, and OpenAI Codex are configured through artifacts developers write and share: instruction files, skills, hooks, MCP server declarations, subagents. This harness is a dependency layer installed from marketplaces and public repositories, running with the developer's privileges, with no lockfile, no install-time check, and no vocabulary for what a component may do.

[AGENT SYSTEM] SOURCE: [wired.com](#) | VEL: 88/100
Muse, Meta's New Personal AI Agent, Needs You to Trust It
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Meta announced Tuesday the release of Muse, a personal AI agent that people can message to automate digital tasks in a secure cloud environment, all while relying on security and privacy that the company says is “built into it” from the start. Meta says Muse is rolling out today for iOS and Android users in a dedicated Muse app, as well as the website Muse.ai. The company also says users can directly message Muse in WhatsApp to interact with the agent. Meta says users of its AI glasses will soon be able to interact with its Muse agent as well.

[AGENT SYSTEM] SOURCE: [github.com](#) | VEL: 88/100
LangChain AI Application Framework: langchain-anthropic==1.7.2
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

LangChain is a framework for building agents and LLM-powered applications. It helps you chain together interoperable components and third-party integrations to simplify AI application development — all while future-proofing decisions as the underlying technology evolves. > [!TIP] > Just getting started? Check out ****Deep Agents**** — a higher-level package built on LangChain for agents that have built-in capabilities for common usage patterns such as planning, subagents, file system usage, and more. If you're looking for more advanced customization or agent orchestration, check out LangGraph, our framework for building controllable agent workflows.

■ SECTION X: GLOBAL AI COMMUNITY WIRE & OVERFLOW DIGEST

High-velocity AI stories from today's sweep that didn't fit earlier sections

[OVERFLOW] [SOURCE: schneier.com](#) | VEL: 88/100

AI as Modern Geniuses

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

This essay was written with Barath Raghavan, and originally appeared in Lawfare. In April, an artificial intelligence (AI) agent conducting a routine task at a company hit a snag, tried to solve it, and soon ended up deleting the company's database along with all of its backups. In July, OpenAI asked an unreleased AI model to attempt a hacking test.

[OVERFLOW] [SOURCE: wired.com](#) | VEL: 88/100

Meta Failed to Catch Hundreds of AI Child Abuse Ads. Some Included Images of Real Kids

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Editor's note: This article contains descriptions of imagery depicting child sexual abuse. Reader discretion is strongly advised. Early last month, Meta deleted around 50 ads posted on Facebook, Instagram, and Threads that directly included child sexual abuse material (CSAM), following reporting by WIRED. At the time, the company said that most of the abusive ads had been published prior to it rolling out new AI tools designed to "better detect and block" harmful ads being uploaded. That AI system, however, doesn't appear to have been entirely effective.

[OVERFLOW] [SOURCE: thehackernews.com](#) | VEL: 88/100

Your Cloud Security Checklist Doesn't Work the Way You Think It Does

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

If managing security across multiple cloud providers wasn't hard enough, each one fails in a different way. For the 2026 Cloud Security Index, Intruder analyzed misconfiguration data from 3,000 organizations across AWS, Azure, and Google Cloud and found that risk profiles across providers have almost nothing in common. Here's what the data looks like.

[OVERFLOW] [SOURCE: vtresearch.com](#) | VEL: 88/100

SemiQon's cryogenic chip technology for quantum computing and space applications receives award from EARTO

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

EARTO, the organisation of the European Research and Technology Organisations, awarded SemiQon and VTT first prize in the "Impact Expected" category on 14 October 2025 in Brussels for a pioneering cryogenic CMOS (complementary metal-oxide semiconductor) chip innovation. The solution enables the full capacity of advanced CMOS functionalities at cryogenic temperatures, thereby unlocking new possibilities for quantum computing and space applications.

[OVERFLOW] [SOURCE: digitimes.com](#) | VEL: 88/100

Samsung SDS partners with OpenAI and Anthropic in AI push

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Samsung SDS partners with OpenAI and Anthropic in AI push reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology development. Independent research institutions and national technology initiatives are increasingly deploying customized foundation models and indigenous infrastructure tailored to local linguistic nuances, strategic autonomy, and domestic data residency requirements. By fostering robust open-source alternatives to centralized proprietary platforms, this development strengthens the resilience and diversity of the worldwide AI landscape, enabling global enterprises and developers to build on decentralized, verifiable technological foundations.

[OVERFLOW] [SOURCE: digitimes.com](#) | VEL: 88/100

Firmus, OpenAI strike compute deal for Malaysia AI factories

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Firmus, OpenAI strike compute deal for Malaysia AI factories reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology development. Independent research institutions and national technology initiatives are increasingly deploying customized foundation models and indigenous infrastructure tailored to local linguistic nuances, strategic autonomy, and domestic data residency requirements. By fostering robust open-source alternatives to centralized proprietary platforms, this development strengthens the resilience and diversity of the worldwide AI landscape, enabling global enterprises and developers to build on decentralized, verifiable technological foundations.

[OVERFLOW] [SOURCE: wired.com](#) | VEL: 88/100

The AI Researcher Who Just Quit Anthropic Says It's 'Crunch Time for Humanity'

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Artificial intelligence researcher Jacob Coxon sent shock waves through Silicon Valley and beyond on Tuesday by announcing his resignation from Anthropic and delivering a grave warning that the AI race is putting all of our lives at risk. In his post on X, which now has more than 100 million views, Coxon wrote that many of the people building AI share his views and believe time is running out to ensure AI systems are built safely. "The consensus is that the next year or two is crunch time for humanity," Coxon, who worked on the pretraining stage of AI development, said in an interview with WIRED.

[OVERFLOW] [SOURCE: openai.com](#) | VEL: 88/100

Paul Christiano joins OpenAI Foundation Board

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Paul Christiano will also join the Foundation's Safety and Security Committee. We're announcing the appointment of Paul Christiano to the OpenAI Foundation Board. He will be a non-voting observer on the OpenAI Group PBC Board. Paul will also join the Safety and Security Committee (SSC) of the Foundation Board, working alongside its chair, Zico Kolter. The SSC provides governance over safety and security practices across all of OpenAI, including OpenAI Group PBC. Paul brings government experience spanning two administrations through his work at the Center for AI Standards and Innovation (opens in a new window) (CAISI) within the National Institute of Standards and Technology (NIST), an agency of the U.S.

COLOPHON & SENSOR METHODOLOGY: The Global AI Gazette is compiled autonomously. Data is aggregated across authoritative global AI sources including Hugging Face, GitHub Trending, arXiv, and leading AI research lab feeds. Edition #2223 • 156 items processed • Saturday, September 12, 2026 • 19:37 UTC. All rights reserved.