

THE GLOBAL AI GAZETTE

Worldwide AI Ecosystem Intelligence • Open Source • Research • Models • Tools • Sovereign AI

Thursday, September 10, 2026 • 19:55 UTC

EDITION NO. 2213

239 STORIES ANALYZED

■ LEAD STORY

guide to using reasoning_effort on deepseek v4.1 flash

Velocity 96/100 • Impact 88/100 • Source: reddit.com

guide to using reasoning_effort on deepseek v4.1 flash marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

INNOVATION: Architecture: Reasoning LLM & Multi-Step Chain-of-Thought
ACTION: Deployment: Checkpoints accessible on Hugging Face; compatible with standard runtimes.

■ SECONDARY ANCHOR

Deepseek V4.1 Flash Release Video [Made with Deepseek V4.1 Flash]

I like to benchmark new models that come out on motion videos. So here's a test I did for deepseek v4.1 flash. And I have to say flash has probably graduated from being a Luna class model to nearly an Opus class model with this release, at least with motion videos. Prev.

■ TOP FLASH BULLETINS

[VEL 96] ANOTHER researcher accuses OpenAI of training on conversations and then claiming a breakthrough — ANOTHER researcher accuses OpenAI of training on conversations and then claiming a breakthrough reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology...

[VEL 96] Qwen3.8-27B-Uncensored-Genesis-V1-GGUF — Model available Qwen3.8-27B-Uncensored-Genesis-V1-MTP-GGUF This model is a practical realisation of things described in this paper, but adapted by me for machine learning: I am trying...

[VEL 88] LangChain AI Application Framework: langchain-openai==1.6.2 — LangChain is a framework for building agents and LLM-powered applications. It helps you chain together interoperable components and third-party integrations to simplify AI application development...

[VEL 88] When LLM Agents Negotiate: Private Information and Dynamic Bargaining in Supply Chains — As LLM agents move from decision support to autonomous procurement, firms need to know whether delegated negotiators create value, divide it predictably, and avoid money-losing...

[VEL 88] Social Chain of Thought: A Multi-Agent Architecture Grounded in Medical Differential Diagnosis Methodology — Medical diagnostic reasoning is a high-impact use case for LLMs that carries significant implications for the health and wellbeing of users. When OpenAI (2026) reports...

SECTION II: EXECUTIVE AI STRATEGIC BRIEFING

Strategic Horizons & Action Directives for Enterprise Leadership

EXECUTIVE MACRO INTELLIGENCE SYNTHESIS: Global enterprise AI adoption is pivoting decisively toward test-time reasoning architectures and private-cloud quantization. Technology leadership must actively balance proprietary frontier model APIs with sovereign, open-weight deployments to reduce token expenditure while strictly sandboxing autonomous agent tool-calling boundaries.

VECTOR	DEVELOPMENT	SOURCE	DIRECTIVE
Model Sourcing & Licensing	HealthBench-Psych: A Mental Health Subse	arxiv.org	Audit open-weights licensing vs proprietary APIs; evaluate OpenAI para
Compute & Infrastructure CapEx	Sycophantic Anchors: Localizing and Quan	arxiv.org	Review GPU cluster allocation and power envelopes; benchmark Sycophant
Agentic Autonomy & Governance	vLLM Inference Engine: v0.29.0rc3	github.com	Implement deterministic sandboxes for vLLM autonomous tool execution,
Inference Latency & Quantization	Claude Fable 5.1 made me a really nice a	simonwillison.n	Benchmark KV-cache compression (FP8/INT4/GGUF) and modern inference en
Open-Source Supply Chain	DeepSeek V4.1 Flash is available in Hugg	reddit.com	Inspect upstream repository dependencies; audit tokenizer code, weight
Data Residency & Sovereignty	Running Vision Qwen 3.8 27B on a 16GB Ca	reddit.com	Verify compliance with sovereign AI frameworks and regional data resid

1. **HealthBench-Psych: A Mental Health Subset of OpenAI's HealthBench:** Audit open-weights licensing vs proprietary APIs; evaluate OpenAI parameter efficiency and commercial distribution terms.
2. **Sycophantic Anchors: Localizing and Quantifying User Agreement in Reasoning Models:** Review GPU cluster allocation and power envelopes; benchmark Sycophantic Anchors: Localiz hardware efficiency to optimize cost per token.
3. **vLLM Inference Engine: v0.29.0rc3:** Implement deterministic sandboxes for vLLM autonomous tool execution, strict rate limiting, and human-in-the-loop validation.
4. **Claude Fable 5.1 made me a really nice animated pelican:** Benchmark KV-cache compression (FP8/INT4/GGUF) and modern inference engines for Claude against TTFT SLAs.
5. **DeepSeek V4.1 Flash is available in HuggingChat:** Inspect upstream repository dependencies; audit tokenizer code, weights provenance, and pinned runtime releases for DeepSeek.
6. **Running Vision Qwen 3.8 27B on a 16GB Card, the config (45tks).:** Verify compliance with sovereign AI frameworks and regional data residency requirements for Qwen deployments.

SECTION III: TRENDING OPEN-SOURCE AI & GITHUB INNOVATIONS

Top GitHub Repositories, Developer Velocity & Architecture Dissections

PROJECT / REPO	SOURCE	STACK	VELOCITY	CORE ARCHITECTURAL FOCUS
LangChain AI Application Framework: langchain	github.com	Python	88/100	LangChain is a framework for building agents and LLM-powered applications. It helps you chain together interoperable components and.
DeepSeek-V3 Open Model Codebase: v1.0.0	github.com	Python	96/100	We present DeepSeek-V3, a strong Mixture-of-Experts (MoE) language model with 671B total parameters with 37B activated for each.
vLLM Inference Engine: v0.28.0rc2: [Spec Deco	github.com	Python	88/100	Easy, fast, and cheap LLM serving for everyone Documentation Blog Paper Twitter/X User.
vLLM Inference Engine: v0.28.0	github.com	Python	88/100	v0.28.0 Highlights This release features 584 commits from 270 contributors (76 new)! Kimi-K3 performance push: a major optimization.
vLLM Inference Engine: v0.28.1rc0	github.com	Python	88/100	Easy, fast, and cheap LLM serving for everyone Documentation Blog Paper Twitter/X User.
vLLM Inference Engine: v0.29.0rc1: [Bugfix] H	github.com	Python	88/100	Easy, fast, and cheap LLM serving for everyone Documentation Blog Paper Twitter/X User.

[TRENDING REPO] SOURCE: [github.com](#) | VEL: 88/100
LangChain AI Application Framework: langchain-openai==1.6.2
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX
LangChain is a framework for building agents and LLM-powered applications. It helps you chain together interoperable components and third-party integrations to simplify AI application development — all while future-proofing decisions as the underlying technology evolves. > [!TIP] > Just getting started? Check out ****Deep Agents**** — a higher-level package built on LangChain for agents that have built-in capabilities for common usage patterns such as planning, subagents, file system usage, and more. If you're looking for more advanced customization or agent orchestration, check out LangGraph, our framework for building controllable agent workflows.

[TRENDING REPO] SOURCE: [github.com](#) | VEL: 96/100
DeepSeek-V3 Open Model Codebase: v1.0.0
ARCH: MOE • CTX: 788M • QUANT: Native / FP16 • STACK: PyTorch / ONNX
We present DeepSeek-V3, a strong Mixture-of-Experts (MoE) language model with 671B total parameters with 37B activated for each token. To achieve efficient inference and cost-effective training, DeepSeek-V3 adopts Multi-head Latent Attention (MLA) and DeepSeekMoE architectures, which were thoroughly validated in DeepSeek-V2. Furthermore, DeepSeek-V3 pioneers an auxiliary-loss-free strategy for load balancing and sets a multi-token prediction training objective for stronger performance. We pre-train DeepSeek-V3 on 14.8 trillion diverse and high-quality tokens, followed by Supervised Fine-Tuning and Reinforcement Learning stages to fully harness its capabilities. Comprehensive evaluations reveal that DeepSeek-V3 outperforms other open-source models and achieves performance comparable to leading closed-source models. Despite its excellent performance, DeepSeek-V3 requires only 2.788M H800 GPU hours for its full training..

[TRENDING REPO] SOURCE: [github.com](#) | VEL: 88/100
vLLM Inference Engine: v0.28.0rc2: [Spec Decode] DFlash2: local convolution + candidate selector (#52816)
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: vLLM
Easy, fast, and cheap LLM serving for everyone | Documentation | Blog | Paper | Twitter/X | User Forum | Developer Slack | ■ We have built a vLLM website to help you get started with vLLM. Please visit [vllm.ai](#) to learn more. For events, please visit [vllm.ai/events](#) to join us.

[TRENDING REPO] SOURCE: [github.com](#) | VEL: 88/100
vLLM Inference Engine: v0.28.0
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: vLLM
v0.28.0 Highlights This release features 584 commits from 270 contributors (76 new)! Kimi-K3 performance push: a major optimization effort for Kimi-K3 across the stack — Decode Context Parallel (DCP) support (#50484), fused FlashKDA decode and prefill kernels (#50654, #51311, #52458), SiTU activation support for MegaMoE (#50510), GEMM-RS for sequence parallelism (#52079), combined all-gathers with 1.5~3x kernel-level speedup (#51070), an adaptive speculative token budget delivering ~60% better.
[TRENDING REPO] SOURCE: [github.com](#) | VEL: 88/100
vLLM Inference Engine: v0.28.1rc0
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: vLLM
Easy, fast, and cheap LLM serving for everyone | Documentation | Blog | Paper | Twitter/X | User Forum | Developer Slack | ■ We have built a vLLM website to help you get started with vLLM. Please visit [vllm.ai](#) to learn more. For events, please visit [vllm.ai/events](#) to join us.

■ SECTION IV: FRONTIER FOUNDATION MODELS & REASONING
BREAKTHROUGHS

Reasoning Architectures, Open Weights & Capability Benchmarks

FOUNDATION MODEL	SOURCE	SIZE / CTX	IMPACT	ARCHITECTURAL HIGHLIGHT
ANOTHER researcher accuses OpenAI of training	reddit.com	MoE / SOTA / 128K	88/100	ANOTHER researcher accuses OpenAI of training on conversations and then claiming a breakthrough reflects the rapid acceleration of.
Qwen3.8-27B-Uncensored-Genesis-V1-GGUF	reddit.com	27B / 128K	88/100	Model available Qwen3.8-27B-Uncensored-Genesis-V1-MTP-GGUF This model is a practical realisation of things described in this paper, but adapted by.
Sycophantic Anchors: Localizing and Quantifyi	arxiv.org	1.5B / 128K	88/100	Reasoning models frequently agree with incorrect user suggestions -- a behavior known as sycophancy. However, it is unclear.
Claude Fable 5.1 made me a really nice animat	simonwillison.ne	MoE / SOTA / 128K	88/100	Today is Claude Fable (and Mythos) 5.1 day. Anthropic say that Fable 5.1 "sets a new standard for.
DeepSeek V4.1 Flash is available in HuggingCh	reddit.com	MOE / 128K	68/100	DeepSeek V4.1 Flash is available in HuggingChat marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput.
Deepseek v4.1 flash finally has engrams, what	reddit.com	MoE / SOTA / 128K	68/100	If the ratio is the same, Maybe 1.6T -3.1T params plus .56T-1.06T engrams and fable 5.0 level performance?.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 88/100

ANOTHER researcher accuses OpenAI of training on conversations and then claiming a breakthrough

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

ANOTHER researcher accuses OpenAI of training on conversations and then claiming a breakthrough reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology development. Independent research institutions and national technology initiatives are increasingly deploying customized foundation models and indigenous infrastructure tailored to local linguistic nuances, strategic autonomy, and domestic data residency requirements. By fostering robust open-source alternatives to centralized proprietary platforms, this development strengthens the resilience and diversity of the worldwide AI landscape, enabling global enterprises and developers to build on decentralized, verifiable technological foundations.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 88/100

Qwen3.8-27B-Uncensored-Genesis-V1-GGUF

ARCH: 27B • CTX: 128K • QUANT: GGUF • STACK: PyTorch / ONNX

Model available Qwen3.8-27B-Uncensored-Genesis-V1-MTP-GGUF This model is a practical realisation of things described in this paper, but adapted by me for machine learning: I am trying to solve the problem: why LLM models even for simple questions write walls of text during reasoning, and burn too much tokens instead of solving the task. And when number of parameters increase the problem became worse.

[FRONTIER MODEL] SOURCE: arxiv.org | IMPACT: 88/100

Sycophantic Anchors: Localizing and Quantifying User Agreement in Reasoning Models

ARCH: 1.5B • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Reasoning models frequently agree with incorrect user suggestions -- a behavior known as sycophancy. However, it is unclear where in the reasoning trace this agreement originates and how strong the commitment is. We introduce \emph{sycophantic anchors} -- sentences identified via counterfactual analysis that commit models to user agreement. Across four reasoning models spanning three architecture families (Llama, Qwen, Falcon-hybrid) and 1.5B--8B parameters, we analyze over 200,000 counterfactua.

[FRONTIER MODEL] SOURCE: simonwillison.net | IMPACT: 88/100

Claude Fable 5.1 made me a really nice animated pelican

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Today is Claude Fable (and Mythos) 5.1 day. Anthropic say that Fable 5.1 "sets a new standard for coding, knowledge work, and long-running problem-solving tasks". Their announcement spends a notable amount of time on scientific research, boasting of a 52.6% score on the brand new Terminal-Bench-Science 0.1 benchmark (first announced on August 27th), up from 24.7% for Fable 5, 29.0% for Opus 5 and 22.4% for GPT-5.6 Sol.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 68/100

DeepSeek V4.1 Flash is available in HuggingChat

ARCH: MOE • CTX: 128K • QUANT: FP8 • STACK: vLLM

DeepSeek V4.1 Flash is available in HuggingChat marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

■ SECTION V: TOP AI RESEARCH PAPERS & ARXIV BREAKTHROUGHS

Reasoning Paradigms, Multimodal Architectures & Algorithmic Design

[AI RESEARCH] SOURCE: [arxiv.org](#) | VEL: 88/100

HealthBench-Psych: A Mental Health Subset of OpenAI's HealthBench

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

General-purpose health benchmarks increasingly anchor claims about LLM medical performance, but they are not always resolved by clinical specialty, making domain-specific performance hard to isolate. Mental health is of acute public-health concern as millions of people turn to LLMs for psychological support, and most existing evaluations are bespoke academic benchmarks that are difficult to integrate into developer workflows. We introduce HealthBench-Psych and HealthBench-Psych-Hard.

[AI RESEARCH] SOURCE: [arxiv.org](#) | VEL: 88/100

A Human Audit of OpenAI's AI-Generated Mathematical Proofs

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

We assess 18 chapter-specific reviews of the ten mathematical results announced by OpenAI on 1 August 2026, alongside review standards, Lean formalizations, subsequent research, and mathematical references. The article audits this review record without claiming a complete reconstruction of all ten proofs. No confirmed substantive mathematical error in a principal result remains in the examined assessments, although review depth varies and some dependencies remain partly checked.

[AI RESEARCH] SOURCE: [arxiv.org](#) | VEL: 88/100

The Phenomenology of Machine: A Comprehensive Analysis of the Sentience of the OpenAI-o1 Model Integrating Functionalism, Consciousness Theories, Active Inference, and AI Architectures

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

This paper explores the hypothesis that the OpenAI-o1 model—a transformer-based AI trained with reinforcement learning from human feedback (RLHF)—displays characteristics of consciousness during its training and inference phases. Adopting functionalism, which argues that mental states are defined by their functional roles, we assess the possibility of AI consciousness.

[AI RESEARCH] SOURCE: [arxiv.org](#) | VEL: 88/100

Diagram-MMU: A Multi-Modal Benchmark for Scientific Diagrams

ARCH: MoE / SOTA • CTX: 7K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Multimodal Large Language Models (MLLMs) have been growing the capability for scientific writing and collaboration. For example, OpenAI Prism is a free workspace for scientific writing and collaboration. One important feature in Prism is turning scientific diagrams directly into LaTeX TikZ code. In this paper, we build a benchmark, Diagram-MMU, a multi-modal benchmark designed to assess MLLMs' ability for scientific diagram parsing and understanding. Diagram-MMU features 3.7k curated diagrams.

[AI RESEARCH] SOURCE: [arxiv.org](#) | VEL: 88/100

Can LLMs Reason in a Legally Meaningful Manner? A Small-scale Study on European Court of Human Rights Cases

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Reasoning has become a standard technique and feature for contemporary LLMs; however, its application and quality in the context of demanding legal-oriented tasks, such as legal case forecasting, remain under explored. We investigate how LLMs reason in the context of legal case forecasting, using legal cases from the European Court of Human Rights (ECtHR) as a testbed. We evaluate OpenAI GPT 5.4, a recent top-tier LLM, by exploring alternative prompting strategies that are more or less suggestiv.

[AI RESEARCH] SOURCE: [arxiv.org](#) | VEL: 88/100

On Train-Test Class Overlap and Detection for Image Retrieval

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

How important is it for training and evaluation sets to not have class overlap in image retrieval? We revisit Google Landmarks v2 clean, the most popular training set, by identifying and removing class overlap with Revisited Oxford and Paris [34], the most popular evaluation set. By comparing the original and the new RGLDv2-clean on a benchmark of reproduced state-of-the-art methods, our findings are striking.

■ ■ SECTION VI: DEVELOPER TOOLS, FRAMEWORKS & AI INFRASTRUCTURE

Inference Runtimes, Evaluation Frameworks, Vector DBs & Local Execution Engines

[AI TOOL] SOURCE: simonwillison.net | VEL: 88/100

Claude's new system prompt really doesn't want to reproduce song lyrics

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Anthropic publish the system prompts for their Claude consumer applications (Claude.ai and the Claude mobile apps - sadly not for Claude Cowork or Claude Code). I love that they do this, and that they share not just the current prompts but historic changes to their prompts as well. They used to keep all of the prompts on a single page, but when I checked today I noticed they had re-arranged those prompts into an index page and then a page per model - here's the page for Haiku 4.5 for example.

[AI TOOL] SOURCE: simonwillison.net | VEL: 88/100

llm-gemini 0.34

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Release: llm-gemini 0.34 New model gemini-3.8-flash for Gemini 3.8 Flash, with low, medium and high thinking levels. #146 Fixed async responses failing to record the resolved model version. Thanks, Charlie Tonneslan. #137 Google released Gemini 3.8 Flash (and 3.8 Flash Cyber, but that's available to "trusted defenders" only) today. Here are the pelicans for high, medium, and low. This is high: For comparison, here are the same pelicans generated using Gemini 3.7 Flash.

[AI TOOL] SOURCE: github.com | VEL: 88/100

Ollama Local Inference Runtime: v0.33.3

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: Ollama

The official Ollama Docker image `ollama/ollama` is available on Docker Hub. You'll be prompted to run a model or connect Ollama to your existing agents or applications such as `Claude Code`, `OpenClaw`, `OpenCode`, `Codex`, `Copilot`, and more. Supported integrations include Claude Code, Codex, Copilot CLI, DeepSeek Harness, Droid, and OpenCode.

[AI TOOL] SOURCE: github.com | VEL: 80/100

requarks/wiki: Wiki.js | Next Generation Open Source Wiki...

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

A modern, lightweight and powerful wiki app built on NodeJS ****Official Website**** - ****Documentation**** - Requirements - Installation - Demo - Changelog - Feature Requests - Chat with us on Discord - Translations ***(We need your help!)*** - Special Thanks - Contribute Follow our Twitter feed to learn about upcoming updates and new releases!

[AI TOOL] SOURCE: github.com | VEL: 80/100

windmill-labs/windmill: Open-source developer platform to power your entire infra and turn scripts into ...

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Open-source developer platform for internal code: APIs, background jobs, workflows and UIs. Self-hostable alternative to Retool, Pipedream, Superblocks and a simplified Temporal with autogenerated UIs and custom UIs to trigger workflows and scripts as internal apps. Scripts are turned into sharable UIs automatically, and can be composed together into flows or used into richer apps built with low-code. Supported languages: Python, TypeScript, Go, Bash, SQL, GraphQL, PowerShell, Rust, and more.

[AI TOOL] SOURCE: github.com | VEL: 80/100

grokability/snipe-it: A free open source IT asset/license management system...

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Snipe-IT - Open Source Asset Management System This is a FOSS project for asset management in IT Operations. Knowing who has which laptop, when it was purchased in order to depreciate it correctly, handling software licenses, etc. Snipe-IT is actively developed and we release quite frequently. (Check out the live demo here.).

■ SECTION VII: SOVEREIGN AI & WORLDWIDE REGIONAL ECOSYSTEMS

National AI Labs, Regional Models & Sovereign Compute Ecosystems

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

Running qwen 3.8 27B iq3 xxs on RTX 3060.

ARCH: 27B • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Getting anywhere from 10 - 20 tps. Thinking Off . Took about 4 mins and 7 mins. Running on about "IQ3_S - 3.4375 bpw" 37.03.960.932 I slot print_timing: id 0 | task 2665 | prompt processing, n_tokens = 12516, progress = 0.98, t = 37.08 s / 337.55 tokens per second 37.04.777.411 I slot print_timing: id 0 | task 2665 | prompt processing, n_tokens = 12768, progress = 1.00, t = 37.78 s / 337.94 tokens per second 37.13.746.606 I slot print_timing: id 0 | task 2665 | n_gen = 100, tg = 11.34 t/s.

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

Adding emotion control tags to Qwen3-TTS

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

I fine-tuned Qwen3-TTS with inline transcript control tags rather in lieu of a separate instruction parameter and thought the community here might be interested in the result. The basic premise was to use Qwen's CustomVoice model as a teacher prompted with natural language directions ("Speak with intense anger") pairing the resulting audio with a discrete control token for the student to learn ("Anger").

[SOVEREIGN [GLOBAL]] SOURCE: [github.com](#) | VEL: 72/100

llama.cpp High-Performance LLM Engine: b10896

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: llama.cpp

spec: fix failed to decode mtmd chunk with DFlash (#28587) speculative: fix failed to decode mtmd chunk with DFlash When using DFlash w/ vision models, the drafter memory fails to allocate new tokens because images report a fixed offset. Stop copying them to allow the drafter to continue. address PR feedback limit M-RoPE skip to images only, allow audio to pass through.

[SOVEREIGN [GLOBAL]] SOURCE: [github.com](#) | VEL: 72/100

llama.cpp High-Performance LLM Engine: b10897

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: llama.cpp

ci : Update WoA CUDA 13.4 release to use 13.4.1 GA redistributables (#28687) Move Windows ARM64 CUDA 13.4 builds from the Developer Preview archives to the 13.4.1 GA redistributables Website: Attestations: macOS/iOS: macOS Apple Silicon (arm64) macOS Apple Silicon (arm64, KleidiAI enabled) DISABLED macOS Intel (x64) iOS XCFramework Linux: Ubuntu x64 (CPU) Ubuntu arm64 (CPU) Ubuntu s390x (CPU) Ubuntu x64 (Vulkan) Ubuntu.

[SOVEREIGN [GLOBAL]] SOURCE: [github.com](#) | VEL: 72/100

llama.cpp High-Performance LLM Engine: b10900

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: llama.cpp

vulkan: use add_alloc_dep to enable topk_moe fusion for prefill (#28422) Website: Attestations: macOS/iOS: macOS Apple Silicon (arm64) macOS Apple Silicon (arm64, KleidiAI enabled) DISABLED macOS Intel (x64) iOS XCFramework Linux: Ubuntu x64 (CPU) Ubuntu arm64 (CPU) Ubuntu s390x (CPU) Ubuntu x64 (Vulkan) Ubuntu arm64 (Vulkan) Ubuntu x64 (ROCm 10.0) Ubuntu x64 (OpenVINO) Ubuntu x64 (SYCL FP32) Ubuntu x64 (SYCL FP16).

[SOVEREIGN [GLOBAL]] SOURCE: [github.com](#) | VEL: 72/100

Hugging Face Transformers: Release 5.17.0

ARCH: 780B • CTX: 1M • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Release v5.17.0 New Model additions HYV4 Hy4-Preview is a 780B-parameter mixture-of-experts language model that activates 49B parameters per token. Each MoE layer holds 256 routed experts plus one always-active shared expert and routes every token to 8 of them. The context window is 1M tokens. The architecture combines four features: Multi-head Latent Attention (MLA) compresses keys and values into a low-rank latent (kv_lora_rank) that kv_b_proj expands back to one key/value per query head.

SECTION VIII: AI HARDWARE, COMPUTE CLUSTERS & SILICON

GPU Clusters, Custom Accelerators, Datacenter Infrastructure & Silicon

SILICON / SYSTEM	SOURCE	COMPUTE SPECS	VELOCITY	BENCHMARK / TELEMETRY
vLLM Inference Engine: v0.29.0rc3	github.com	High Velocity	88/100	Easy, fast, and cheap LLM serving for everyone Documentation Blog Paper Twitter/X User.
Running Vision Qwen 3.8 27B on a 16GB Card, t	reddit.com	High Velocity	96/100	I am just sharing my config for Qwen 3.8 27b that fits on a 5060TI, what is cool.
Qwen3.8-Flash-Next on 2x3090 + DDR4, part 4:	reddit.com	29 t/s	96/100	. 17 -> 25-29 t/s with the expert cache PR, 37-41 t/s after switching to UD-Q4_K_XL and stacking.
Qwen3.8-Flash-Next on MLX-serve, 1m context i	reddit.com	40 tok/s	96/100	Hi, I'm the co-creator of this Qwen3.8-Flash-Next engine support in MLX-serve. I've been tuning this one to run.
llama.cpp High-Performance LLM Engine: b10899	github.com	High Velocity	88/100	vulkan: small M matrix optimizations for qwen (#28457) vulkan: optimize m=1 mul_mat by swapping A/B vulkan: Improve small.
Ollama Local Inference Runtime: v0.34.0-rc5	github.com	High Velocity	88/100	The official Ollama Docker image `ollama/ollama` is available on Docker Hub. You'll be prompted to run a model.

[AI HARDWARE] SOURCE: github.com | VEL: 88/100

vLLM Inference Engine: v0.29.0rc3

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: vLLM

Easy, fast, and cheap LLM serving for everyone | Documentation | Blog | Paper | Twitter/X | User Forum | Developer Slack | ■ We have built a vLLM website to help you get started with vLLM. Please visit vllm.ai to learn more. For events, please visit vllm.ai/events to join us.

[AI HARDWARE] SOURCE: reddit.com | VEL: 96/100

Running Vision Qwen 3.8 27B on a 16GB Card, the config (45tks).

ARCH: 27B • CTX: 85K • QUANT: GGUF • STACK: PyTorch / ONNX

I am just sharing my config for Qwen 3.8 27b that fits on a 5060TI, what is cool about this is that you can even get vision! and a 85K context (I have 1.5gb of headroom for more context or a better quant) Model: IQ3_XXS-mtp from Using be llama Config used: [*] model = ..\llm-models\Qwen3.8-27B-GSQ-RCO-IQ3_XXS-mtp.gguf mmp proj = ..\llm-models\mmp proj-Qwen3.8-27B-BF16.gguf image-min-tokens = 256.

[AI HARDWARE] SOURCE: reddit.com | VEL: 96/100

Qwen3.8-Flash-Next on 2x3090 + DDR4, part 4: 2.2-2.5x faster prefill by kicking the expert cache off the GPU while the prompt runs

ARCH: MoE / SOTA • CTX: 8K • QUANT: UD-Q4_K_XL • STACK: llama.cpp

. 17 -> 25-29 t/s with the expert cache PR, 37-41 t/s after switching to UD-Q4_K_XL and stacking MTP on the cache, the top-k fallback that was sorting more than it needed to. This one is all about prefill, which was honestly the weak spot the whole time. 80+ seconds before the first token on an 8k prompt, and 24 minutes on a 119k one...I know . Box is still 2x 3090, dual Broadwell Xeon, llama.cpp, UD-Q4_K_XL with the Q8 MTP head on the.

[AI HARDWARE] SOURCE: reddit.com | VEL: 96/100

Qwen3.8-Flash-Next on MLX-serve, 1m context is released!

ARCH: MoE / SOTA • CTX: 1M • QUANT: Native / FP16 • STACK: MLX

Hi, I'm the co-creator of this Qwen3.8-Flash-Next engine support in MLX-serve. I've been tuning this one to run both fast, efficient and correct up 1m context using kv cache 8 bits in M5 Max 128GB. Qwen is working well at very long context as showed in the video (a snapshot at ~760k context), I let it build MLX Serve Monitor plugin that you've seen on the right side of Opencode2's app. The generation sustain through 1m context at around 40 tok/s on prose and 75 tok/s on coding.

[AI HARDWARE] SOURCE: github.com | VEL: 88/100

llama.cpp High-Performance LLM Engine: b10899

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: llama.cpp

vulkan: small M matrix optimizations for qwen (#28457) vulkan: optimize m=1 mul_mat by swapping A/B vulkan: Improve small M perf Allow split_k with small M. Make small vs med tile selection (for coopmat2) depend on M, not just N.

■ SECTION IX: AUTONOMOUS AGENTS, MULTI-AGENT SWARMS & ROBOTICS

Agent Orchestration, Computer-Use Primitives, Action Models & Robotics

AGENT / FRAMEWORK	SOURCE	PROTOCOL	VELOCITY	CORE CAPABILITY DOMAIN
When LLM Agents Negotiate: Private Information	arxiv.org	PyTorch / ONNX	88/100	As LLM agents move from decision support to autonomous procurement, firms need to know whether delegated negotiators create.
Social Chain of Thought: A Multi-Agent Architect	arxiv.org	PyTorch / ONNX	88/100	Medical diagnostic reasoning is a high-impact use case for LLMs that carries significant implications for the health and.
Enhancing Virtual Agents through SLMs and Edge	arxiv.org	PyTorch / ONNX	88/100	Embodied intelligent virtual agents are expected to operate as persistent, adaptive, and context-aware entities within complex virtual and.
Research acceleration: The view inside OpenAI	simonwillison.ne	PyTorch / ONNX	96/100	Research acceleration: The view inside OpenAI Apparently today is RSI day at OpenAI, for Recursive Self-Improvement - I.
AutoGPT Agent Runtime: autogpt-platform-beta-	github.com	PyTorch / ONNX	96/100	■ Release autogpt-platform-beta-v0.7.3 Date: January 2025 ■ What's New? New Features #14003 - Replace the email system with.
LangChain AI Application Framework: langchain	github.com	PyTorch / ONNX	88/100	LangChain is a framework for building agents and LLM-powered applications. It helps you chain together interoperable components and.

[AGENT SYSTEM] SOURCE: [arxiv.org](#) | VEL: 88/100

When LLM Agents Negotiate: Private Information and Dynamic Bargaining in Supply Chains

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

As LLM agents move from decision support to autonomous procurement, firms need to know whether delegated negotiators create value, divide it predictably, and avoid money-losing contracts. We study this in a canonical supply chain bargaining problem: a buyer with private demand information negotiates a quantity-payment contract with an uninformed seller. We benchmark nine LLMs from OpenAI, Google, and Alibaba against a validated Perfect Bayesian Equilibrium across 9,840 LLM-to-LLM negotiations.

[AGENT SYSTEM] SOURCE: [arxiv.org](#) | VEL: 88/100

Social Chain of Thought: A Multi-Agent Architecture Grounded in Medical Differential Diagnosis Methodology

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Medical diagnostic reasoning is a high-impact use case for LLMs that carries significant implications for the health and wellbeing of users. When OpenAI (2026) reports that more than 5% of ChatGPT messages globally are healthcare-related, the transparency of these systems becomes a serious design concern. This is especially true for complex cases, where differential diagnosis often requires integrating multiple forms of specialist reasoning.

[AGENT SYSTEM] SOURCE: [arxiv.org](#) | VEL: 88/100

Enhancing Virtual Agents through SLMs and Edge-Computing: An Exploratory Evaluation of Think and Memory Processes

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Embodied intelligent virtual agents are expected to operate as persistent, adaptive, and context-aware entities within complex virtual and Metaverse worlds. However, implementing cognitively capable agents in such environments is conceptually and technologically challenging.

[AGENT SYSTEM] SOURCE: [simonwillison.net](#) | VEL: 96/100

Research acceleration: The view inside OpenAI

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Research acceleration: The view inside OpenAI Apparently today is RSI day at OpenAI, for Recursive Self-Improvement - I think it's their new AGI. Both this piece and the new essay An Alien Mind (by Chief Scientist Jakub Pachocki) talk about it, and this one doesn't even bother to expand the acronym. Included are details on how OpenAI's own research team are using coding agents.

[AGENT SYSTEM] SOURCE: [github.com](#) | VEL: 96/100

AutoGPT Agent Runtime: autogpt-platform-beta-v0.7.3

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

■ Release autogpt-platform-beta-v0.7.3 Date: January 2025 ■ What's New? New Features #14003 - Replace the email system with the Briefing/Alert/Verdict/Ops design (by @Torantulino) #14156 - Add selectable ask_question options and Needs-You prompting rule to Copilot (by @Abhi1992002) #14143 - Show live compaction progress in the Copilot chat (by @0ubbe) #14099 - Add expert team tools and copilot chat polish behind flags (by @Abhi1992002) #14165 - Fire Google Ads conversions across the signup-to-pa.

■ SECTION X: GLOBAL AI COMMUNITY WIRE & OVERFLOW DIGEST

High-velocity AI stories from today's sweep that didn't fit earlier sections

[OVERFLOW] [SOURCE: reddit.com](#) | VEL: 96/100

I made a way to migrate between embedding models without re-embedding your entire corpus

ARCH: 1B • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

So I was playing with embedding models I saw that when you upgrade from model A to B, you face a very big backfilling cost. Ie, suppose you have a 1b vectors from model A, and then you want to use model B. This would mean you have to re-embed all of your documents with model B before you can even serve with the model, and on an H100, it would take ~108 days (qwen embed 8b, 106 docs/second). But I found an easier way to do it.

[OVERFLOW] [SOURCE: simonwillison.net](#) | VEL: 96/100

Quoting Jakub Pachocki

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

The strongest argument I see for continuing to train much smarter models quickly is the need to build defensive systems against the dangers posed by other AI. [...] We will need powerful, aligned AI for defense; to secure infrastructure, to protect against rogue agents in real time, and to invent entirely new protective measures. This will be a primary focus of OpenAI's deployment efforts.

[OVERFLOW] [SOURCE: simonwillison.net](#) | VEL: 96/100

Introducing ChatGPT Images 2.5

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Introducing ChatGPT Images 2.5 OpenAI's image generation models are apparently used "more than 3 billion images across ChatGPT Images and the GPT-4o Image models in the API". This latest release improves their instruction-following ability across multiple turns, responds faster, and "is better at preserving the subjects in your reference photos". There are two new model IDs in the API: gpt-image-2.5-sunburst and gpt-image-2.5-flare.

[OVERFLOW] [SOURCE: simonwillison.net](#) | VEL: 96/100

On the Navier–Stokes Millennium Prize Problem

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

On the Navier–Stokes Millennium Prize Problem Impressive result from OpenAI, who used an unreleased model to produce a resolution to the Navier–Stokes existence and smoothness problem, one of the seven Millennium Prize Problems that have been subject to a \$1,000,000 prize since May 24th, 2000.

[OVERFLOW] [SOURCE: reddit.com](#) | VEL: 96/100

Would you consider 5t/s usable for a local model?

ARCH: 27B • CTX: 780M • QUANT: Native / FP16 • STACK: PyTorch / ONNX

I'm able to run qwen3.8 27b in two ways on my system: split between my 3060 12gb and 9070xt running at 20t/s or running off the 780m iGPU and 5400mhz DDR5 at 5t/s. Personally I feel like the 5t/s is still more usable because I have enough RAM to still use my system mostly normally. I can even play games just fine because the CPU side is left alone.

[OVERFLOW] [SOURCE: reddit.com](#) | VEL: 96/100

Mention if a "new model" is a finetune

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

A few posts tagged with "new model" present models that are finetunes. My opinion : I'd rather have the "new model" tag reserved for new "major" releases, like a new Qwen model, Deepseek V4 -> Deepseek V4.1, etc., that involved a new pretrain or intensive post-training (in opposition to a small finetune). Otherwise, maybe prepend "[Finetune]" to the title to indicate that the new model is "less of a big news", a use a "new finetune" tag, to differentiate between the two kinds of new models.

[OVERFLOW] [SOURCE: reddit.com](#) | VEL: 96/100

Don't let FOMO win if you're interested in local llm from a hobby/learning aspect

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: vllm

Just a reminder for those out there itching to get into local llms - don't let FOMO or "gear acquisition syndrom" take over. No matter the hobby, it's so easy to get stuck in a trap where we buy more trying to do more only to realize we've lost the fun in it all or even the notion of learning. Obviously, if you're into writing llama or vllm or hardware drivers or whatever - you got to do what you got to do.

[OVERFLOW] [SOURCE: github.com](#) | VEL: 96/100

Open WebUI: v0.9.5

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Added ■■■ Redirect-based SSRF protection. All outbound HTTP requests now block 3xx redirects by default via a new AIOHTTP_CLIENT_ALLOW_REDIRECTS environment variable, preventing redirect-based SSRF where a public URL silently redirects to internal addresses (RFC 1918, loopback, cloud-metadata endpoints). Affected call sites include web fetch, image loading, OAuth discovery, tool server execution, and code interpreter login. #24491 ■■■ Iframe content security policy.

COLOPHON & SENSOR METHODOLOGY: The Global AI Gazette is compiled autonomously. Data is aggregated across authoritative global AI sources including Hugging Face, GitHub Trending, arXiv, and leading AI research lab feeds. Edition #2213 • 239 items processed • Thursday, September 10, 2026 • 19:55 UTC. All rights reserved.