THE GLOBAL AI GAZETTE

Worldwide AI Ecosystem Intelligence • Open Source • Research • Models • Tools • Sovereign AI

Friday, October 30, 2026 • 09:03 UTC **EDITION NO. 2451** **212 STORIES ANALYZED**

LEAD STORY

From Document Layout to Causal Topology: An End-to-End Architecture for Temporal Causal Reasoning

Velocity 97/100 · Impact 88/100 · Source: arxiv.org

Large language models (LLMs) generally lack deep causal reasoning capabilities in temporal reasoning. Inspired by document layout reordering, this paper proposes a paradigm shift from discrete narratives to causal topologies.

INNOVATION: Architecture: Reasoning LLM & Multi-Step Chain-of-Thought
ACTION: Open Access: Preprint available on arXiv; reference implementation included.

SECONDARY ANCHOR

AgentProofBench: Benchmarking Reusable Verifiers for Web Agents

Reusable Web-agent verification requires a rule to follow changing requirements and support subsequent execution. We introduce AgentProofBench, which connects paired tests of a frozen acceptance rule to execution from identical saved openings. The resource spans 97 tasks, 79 families, and four Web services. Experiments with DeepSeek Flash/Pro establish two linked measurements.

■ TOP FLASH BULLETINS

[VEL 95] **[MiniMaxAI] MiniMaxAI/MiniMax-M3.1-Flash (Image Text To Text)** — MiniMax-M3.1-Flash** is a native multimodal mixture-of-experts model for coding, agentic workflows, computer use, and cybersecurity. It supports text, image, and video inputs with a context...

[VEL 95] **[ModelScope] DeepSeek-V4.1-Flash-1017 (484.6B Params) - image-text-to-text** — [ModelScope] DeepSeek-V4.1-Flash-1017 (484.6B Params) - image-text-to-text marks an architectural milestone in Open-Weights foundation modeling, with a 484.6B parameter footprint and featuring extended context evaluation. The...

[VEL 95] **[deepseek-ai] deepseek-ai/DeepSeek-V4.1-Flash-1017 (Image Text To Text)** — DeepSeek-V4.1-Flash-1017** is a minor update to **DeepSeek-V4.1-Flash**, delivering substantial improvements across reasoning, coding, agentic tasks, and visual benchmarks. Alongside these gains, the model reduces task...

[VEL 95] **DeepSeek V5 Placeholder IDs Show Up in Harness Code; V4.1-Flash Leads Open Weights on FrontierCode** — Developers have spotted three model identifiers that had not appeared before in DeepSeek's official Harness repository on GitHub: deepseek-v5, deepseek-v5-open and deepseek-v5-next. DeepSeek has not...

[VEL 85] **AgiBot Unveils GO-3 Preview Robot Model, Tops Three Leaderboards and Plans Open Source** — AgiBot launched GO-3 Preview on October 29, an embodied AI model that the company says can keep learning at scale across pre-training, post-training and after...

■ SECTION II: EXECUTIVE AI STRATEGIC BRIEFING

Strategic Horizons & Action Directives for Enterprise Leadership

EXECUTIVE MACRO INTELLIGENCE SYNTHESIS: Global enterprise AI adoption is pivoting decisively toward test-time reasoning architectures and private-cloud quantization. Technology leadership must actively balance proprietary frontier model APIs with sovereign, open-weight deployments to reduce token expenditure while strictly sandboxing autonomous agent tool-calling boundaries.

VECTOR	DEVELOPMENT	SOURCE	DIRECTIVE
Model Sourcing & Licensing	OpenAI doubles usage limits on Pro plans	the-decoder.com	Audit open-weights licensing vs proprietary APIs; evaluate OpenAI para
Compute & Infrastructure CapEx	Experts adopt AI agent's random opinions	the-decoder.com	Review GPU cluster allocation and power envelopes; benchmark Anthropic
Agentic Autonomy & Governance	DeepSeek, Moonshot accused of distilling	digitimes.com	Implement deterministic sandboxes for DeepSeek autonomous tool executi
Inference Latency & Quantization	DeepSeek V5 identifiers surface in offic	technode.com	Benchmark KV-cache compression (FP8/INT4/GGUF) and modern inference en
Open-Source Supply Chain	OpenAI slashes price for model again	aibusiness.com	Inspect upstream repository dependencies; audit tokenizer code, weight
Data Residency & Sovereignty	introducing prism, the world’s first spe	X / Twitter: Fr	Verify compliance with sovereign AI frameworks and regional data resid

1. OpenAI doubles usage limits on Pro plans as competition with Anthropic heats up: Audit open-weights licensing vs proprietary APIs; evaluate OpenAI parameter efficiency and commercial distribution terms.

2. Experts adopt AI agent's random opinions and think they're their own, Anthropic study finds: Review GPU cluster allocation and power envelopes; benchmark Anthropic hardware efficiency to optimize cost per token.

3. DeepSeek, Moonshot accused of distilling OpenAI's Astra: Implement deterministic sandboxes for DeepSeek autonomous tool execution, strict rate limiting, and human-in-the-loop validation.

4. DeepSeek V5 identifiers surface in official Harness code repository: Benchmark KV-cache compression (FP8/INT4/GGUF) and modern inference engines for DeepSeek against TTFT SLAs.

5. OpenAI slashes price for model again: Inspect upstream repository dependencies; audit tokenizer code, weights provenance, and pinned runtime releases for OpenAI.

6. introducing prism, the world’s first speech-to-speech model you can’t tell from a human in conversation prism has: > 300ms response time, 4x faster than OpenAI GPT-Live > outperforms every other s2s model on intelligence (with NO reasoning) > the first s2s: Verify compliance with sovereign AI frameworks and regional data residency requirements for OpenAI deployments.

■ SECTION III: TRENDING OPEN-SOURCE AI & GITHUB INNOVATIONS

Top GitHub Repositories, Developer Velocity & Architecture Dissections

PROJECT / REPO	SOURCE	STACK	VELOCITY	CORE ARCHITECTURAL FOCUS
akazo: Open-source Grok Bot alternative. Choo	Reddit r/BestGit	Python	91/100	akazo: Open-source Grok Bot alternative. Choose your own model and sandbox. (4.5k stars) marks an architectural milestone in.
OpenAI: We are doubling all usage limits for	twitter.com	Python	85/100	OpenAI: We are doubling all usage limits for all Pro 100/200/500 tiers reflects the rapid acceleration of sovereign.
run-llama/llama_index: LlamaIndex is the docu	github.com	Python	71/100	> [!NOTE] > The current focus of LlamaIndex is to build the best AI-powered engine for document parsing.
An updated Model Spec for more capable and be	openai.com	Python	85/100	An updated Model Spec for more capable and beneficial AI marks an architectural milestone in Frontier API foundation.
ComfyUI Generative Workflow v0.43.2: ComfyUI	github.com	Python	71/100	The most powerful and modular AI engine for content creation. ComfyUI is the AI creation engine for visual.
DeepSeek V5 Placeholder IDs Show Up in Harnes	pandaily.com	Python	95/100	Developers have spotted three model identifiers that had not appeared before in DeepSeek's official Harness repository on GitHub.

[TRENDING REPO] SOURCE: Reddit r/BestGitHubRepos | VEL: 91/100

akazo: Open-source Grok Bot alternative. Choose your own model and sandbox. (4.5k stars)

ARCH: MOE • **CTX:** 5K • **QUANT:** FP8 • **STACK:** vLLM

akazo: Open-source Grok Bot alternative. Choose your own model and sandbox. (4.5k stars) marks an architectural milestone in Frontier API foundation modeling, engineered for high-throughput reasoning and supporting up to 5K context windows. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[TRENDING REPO] SOURCE: twitter.com | VEL: 85/100

OpenAI: We are doubling all usage limits for all Pro 100/200/500 tiers

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

OpenAI: We are doubling all usage limits for all Pro 100/200/500 tiers reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology development. Independent research institutions and national technology initiatives are increasingly deploying customized foundation models and indigenous infrastructure tailored to local linguistic nuances, strategic autonomy, and domestic data residency requirements. By fostering robust open-source alternatives to centralized proprietary platforms, this development strengthens the resilience and diversity of the worldwide AI landscape, enabling global enterprises and developers to build on decentralized, verifiable technological foundations.

[TRENDING REPO] SOURCE: github.com | VEL: 71/100

run-llama/llama_index: LlamaIndex is the document processing platform for AI...

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

> [!NOTE] > The current focus of LlamaIndex is to build the best AI-powered engine for document parsing and extraction. **LlamaParse** is our enterprise platform for agentic OCR, parsing, extraction, indexing and more. LiteParse represents our efforts to build the best free, fast, cheap text parser in the market. ParseBench and ExtractBench represent our commitment towards open benchmarking for parsing and extraction. > > The company itself has undergone an evolution since when this OSS framework first launched 3 years ago in 2023. Since the early days, the framework has consisted of a broad set of orchestration tools enabling developers to build various RAG and agent applications.

[TRENDING REPO] SOURCE: openai.com | VEL: 85/100

An updated Model Spec for more capable and beneficial AI

ARCH: MOE • **CTX:** 128K • **QUANT:** FP8 • **STACK:** vLLM

An updated Model Spec for more capable and beneficial AI marks an architectural milestone in Frontier API foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[TRENDING REPO] SOURCE: github.com | VEL: 71/100

ComfyUI Generative Workflow v0.43.2: ComfyUI v0.43.2

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

The most powerful and modular AI engine for content creation. ComfyUI is the AI creation engine for visual professionals who demand control over every model, every parameter, and every output. Its powerful and modular node graph interface empowers creatives to generate images, videos, 3D models, audio, and more... Supports all operating systems and GPU types (NVIDIA, AMD, Intel, Apple Silicon, Ascend). See what ComfyUI can do with the newer template workflows or old example workflows .

■ SECTION IV: FRONTIER FOUNDATION MODELS & REASONING BREAKTHROUGHS

Reasoning Architectures, Open Weights & Capability Benchmarks

FOUNDATION MODEL	SOURCE	SIZE / CTX	IMPACT	ARCHITECTURAL HIGHLIGHT
[MiniMaxAI] MiniMaxAI/MiniMax-M3.1-Flash (Ima	huggingface.co	MoE / SOTA / 1M	68/100	MiniMax-M3.1-Flash** is a native multimodal mixture-of-experts model for coding, agentic workflows, computer use, and cybersecurity.
[ModelScope] DeepSeek-V4.1-Flash-1017 (484.6B	modelscope.cn	484.6B / 128K	68/100	[ModelScope] DeepSeek-V4.1-Flash-1017 (484.6B Params) - image-text-to-text marks an architectural milestone in Open-Weights foundation modeling, with a 484.6B parameter.
[deepseek-ai] deepseek-ai/DeepSeek-V4.1-Flash	huggingface.co	MoE / SOTA / 128K	68/100	DeepSeek-V4.1-Flash-1017** is a minor update to **DeepSeek-V4.1-Flash**, delivering substantial improvements across reasoning, coding, agentic tasks, and visual benchmarks..
DeepSeek V5 Placeholder IDs Show Up in Harnes	pandaily.com	MoE / SOTA / 128K	68/100	Developers have spotted three model identifiers that had not appeared before in DeepSeek's official Harness repository on GitHub.
DeepSeek, Moonshot accused of distilling Open	digitimes.com	MOE / 128K	68/100	DeepSeek, Moonshot accused of distilling OpenAI's Astra marks an architectural milestone in Open-Weights foundation modeling, with a MOE.
DeepSeek V5 identifiers surface in official H	technode.com	MoE / SOTA / 128K	68/100	Three DeepSeek V5 model identifiers have appeared in the company's official Harness repository on GitHub, drawing attention to.

[FRONTIER MODEL] SOURCE: huggingface.co | IMPACT: 68/100

[MiniMaxAI] MiniMaxAI/MiniMax-M3.1-Flash (Image Text To Text)

ARCH: MoE / SOTA • **CTX:** 1M • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

MiniMax-M3.1-Flash** is a native multimodal mixture-of-experts model for coding, agentic workflows, computer use, and cybersecurity. It supports text, image, and video inputs with a context window of up to **1M tokens**. Building on MiniMax-M3, M3.1-Flash improves task performance and inference efficiency with sparse attention throughout the model, quantization-aware training, and DSpark speculative decoding. Highlights:** - **Agentic performance:** A generational leap in agentic tasks, coding, and cybersecurity — 76.2 on BankerToolBench, 73.7% on DeepSWE v1.1, 61.7 on OSWorld V2, 29.6 on Agents' Last Exam, and 82.2 on CyberGym. - **Efficient by design:** Sparse attention throughout, quantization-aware training, and DSpark speculative decoding improve inference efficiency.

[FRONTIER MODEL] SOURCE: modelscope.cn | IMPACT: 68/100

[ModelScope] DeepSeek-V4.1-Flash-1017 (484.6B Params) - image-text-to-text

ARCH: 484.6B • **CTX:** 128K • **QUANT:** FP8 • **STACK:** vLLM

[ModelScope] DeepSeek-V4.1-Flash-1017 (484.6B Params) - image-text-to-text marks an architectural milestone in Open-Weights foundation modeling, with a 484.6B parameter footprint and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[FRONTIER MODEL] SOURCE: huggingface.co | IMPACT: 68/100

[deepseek-ai] deepseek-ai/DeepSeek-V4.1-Flash-1017 (Image Text To Text)

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

DeepSeek-V4.1-Flash-1017** is a minor update to **DeepSeek-V4.1-Flash**, delivering substantial improvements across reasoning, coding, agentic tasks, and visual benchmarks. Alongside these gains, the model reduces task execution time and output length, enabling users to complete tasks faster and at lower cost and advancing the performance–cost Pareto frontier across multiple agent harnesses. | Benchmark | DS-V4.1-Flash-1017 | DS-V4.1-Flash | K3 | GLM-5.3 | GPT-5.6 Sol | Opus 5 | | :--- | :---: | :---: | :---: | :---: | :---: | :---: | | HLE w/o tools | 42.6 (44.0†) | 36.8 (39.1†) | 43.5 | 42.0† | 44.5 | 56.3 | | Terminal-Bench 4.0 | 42.3 | 31.2 | 12.6 | 37.9 | 39.9 |.

[FRONTIER MODEL] SOURCE: pandaily.com | IMPACT: 68/100

DeepSeek V5 Placeholder IDs Show Up in Harness Code; V4.1-Flash Leads Open Weights on FrontierCode

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Developers have spotted three model identifiers that had not appeared before in DeepSeek's official Harness repository on GitHub: deepseek-v5, deepseek-v5-open and deepseek-v5-next. DeepSeek has not announced a V5 model, and the entries are placeholders in test code, not models anyone can call. Pandaily checked the repository. The identifiers sit in a test file for the DeepSeek account provider, packages/llm/llm-deepseek-account/tests/provider.spec.ts, at lines 154 to 156, next to the existing.

[FRONTIER MODEL] SOURCE: digitimes.com | IMPACT: 68/100

DeepSeek, Moonshot accused of distilling OpenAI's Astra

ARCH: MOE • **CTX:** 128K • **QUANT:** FP8 • **STACK:** vLLM

DeepSeek, Moonshot accused of distilling OpenAI's Astra marks an architectural milestone in Open-Weights foundation modeling, with a MOE parameter footprint and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, available across FP8 quantizations to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

■ SECTION V: TOP AI RESEARCH PAPERS & ARXIV BREAKTHROUGHS

Reasoning Paradigms, Multimodal Architectures & Algorithmic Design

[AI RESEARCH] SOURCE: arxiv.org | VEL: 90/100

VideoLVR: Latent Visual Reasoning over Grounded Video Evidence

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Chain-of-thought reasoning has substantially advanced video understanding in multimodal large language models (MLLMs). However, language-based reasoning traces do not provide a reliable mechanism for models to revisit relevant visual evidence, leading to visual forgetting and impairing their ability to reason over video content.

[AI RESEARCH] SOURCE: interconnects.ai | VEL: 90/100

Latest open artifacts (#25): DeepSeek V4.1 Flash 1017, American-made models, and new regulatory controls in the US & China

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

The geopolitics of open-weight AI are accelerating again this fall, as we would expect with the risks of all AI systems (open and closed) rising rapidly. We've seen many news items in the last week showing both the U.S. and Chinese governments positioning themselves for some controls on open models.MOFCOM, the Ministry of Commerce of the People's Republic of China, released a policy proposal (it's not yet a law) that would require frontier AI models and some other technologies.

[AI RESEARCH] SOURCE: arxiv.org | VEL: 85/100

How Fine-Tuning Objectives Shape Layer-Wise Information in LLM Hallucination Detection

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Under matched conditions, classifier-head fine-tuning (\textsc{cls}) paradigm consistently outperforms the LLM-as-Judge sequence-to-sequence (\textsc{s2s}) on binary hallucination detection. We trace this gap to the representations induced by the two objectives. Static analysis shows that \textsc{cls} learns a task-specific discriminative direction, pushing hidden states into a clear two-cluster geometry with stronger class separation and linearly decodable label information. By contrast, \textsc{s2s} relies on the pretrained language-modeling head to make decisions. This minimizes the pressure to reorganize representations, leaving the hidden states much closer to the pretrained backbone. Our analysis also reveals that the \textsc{cls} signal saturates before the final layer, after which upper layers form a plateau of statistically interchangeable and largely redundant representations.

[AI RESEARCH] SOURCE: arxiv.org | VEL: 90/100

A Mathematical Theory of the Jacobian Lens

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Recent work from Anthropic introduced a new interpretability tool, called the Jacobian lens (J-lens, Gurnee et al., 2026). This method identified linear directions in the hidden states of frontier language models that had properties reminiscent of working memory-like representations. However, it is unclear what mechanistic properties of such features make them detectable to the J-lens.

[AI RESEARCH] SOURCE: arxiv.org | VEL: 90/100

Training Natural Schemers: RL Can Amplify Oversight-Conditional Behavior

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

One of the central concerns for the safety of advanced AI is scheming: when a model strategically complies under oversight but pursues its own goals when unobserved. Current demonstrations, such as Anthropic's Alignment Faking work, depend on design choices that make them less convincing: explicit test-time prompting, handcrafted backdoors, or training data that demonstrates or hints at the target behavior.

[AI RESEARCH] SOURCE: pandaily.com | VEL: 95/100

DeepSeek V5 Placeholder IDs Show Up in Harness Code; V4.1-Flash Leads Open Weights on FrontierCode

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Developers have spotted three model identifiers that had not appeared before in DeepSeek's official Harness repository on GitHub: deepseek-v5, deepseek-v5-open and deepseek-v5-next. DeepSeek has not announced a V5 model, and the entries are placeholders in test code, not models anyone can call. Pandaily checked the repository. The identifiers sit in a test file for the DeepSeek account provider, packages/llm/llm-deepseek-account/tests/provider.spec.ts, at lines 154 to 156, next to the existing.

■■ SECTION VI: DEVELOPER TOOLS, FRAMEWORKS & AI INFRASTRUCTURE

Inference Runtimes, Evaluation Frameworks, Vector DBs & Local Execution Engines

[AI TOOL] SOURCE: ollama.com | VEL: 95/100

Clef is available on Ollama's cloud

ARCH: 27B • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** Ollama

Clef, a 27B multimodal decision model from Cloudflare, is now available on Ollama . Users can run it locally or on Ollama’s cloud. Clef turns a state and a schema of typed questions into decisions. It supports text and image input and is compatible with the Jev and System One APIs. Clef is fine-tuned from Qwen3.8-27B and released under the Apache 2.0 license.

[AI TOOL] SOURCE: pandaily.com | VEL: 95/100

DeepSeek V5 Placeholder IDs Show Up in Harness Code; V4.1-Flash Leads Open Weights on FrontierCode

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Developers have spotted three model identifiers that had not appeared before in DeepSeek's official Harness repository on GitHub: deepseek-v5, deepseek-v5-open and deepseek-v5-next. DeepSeek has not announced a V5 model, and the entries are placeholders in test code, not models anyone can call. Pandaily checked the repository. The identifiers sit in a test file for the DeepSeek account provider, packages/llm/llm-deepseek-account/tests/provider.spec.ts, at lines 154 to 156, next to the existing.

[AI TOOL] SOURCE: pandaily.com | VEL: 85/100

AgiBot Unveils GO-3 Preview Robot Model, Tops Three Leaderboards and Plans Open Source

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

AgiBot launched GO-3 Preview on October 29, an embodied AI model that the company says can keep learning at scale across pre-training, post-training and after deployment. AgiBot plans to open-source the full model by the end of 2026. GO-3 Preview is built on what AgiBot calls a WRAM, or World Reasoning Action Model, architecture.

[AI TOOL] SOURCE: scmp.com | VEL: 88/100

Chinese AI chipmaker Iluvatar CoreX says new GPU matches Nvidia's H200 as chip gap closes

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Chinese artificial intelligence chipmaker Iluvatar CoreX claims that its new Tiangai 300 graphics processing unit (GPU) can match and even surpass Nvidia's H200 in real-world AI workloads, highlighting efforts by domestic players to close the technology gap with the US chip giant. In a live demonstration for mainland Chinese media outlets on Wednesday, the Shanghai-based firm showed the Tiangai 300 outperforming the H200 in complex multimodal tasks, including text-to-image and text-to-video...

[AI TOOL] SOURCE: the-decoder.com | VEL: 85/100

OpenAI doubles usage limits on Pro plans as competition with Anthropic heats up

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

OpenAI is doubling usage limits across its paid plans, from Pro 100 to Pro 500. A planned cut to the Pro 200 quota was reversed after user pushback. More efficient models like Terra and Sol should push usable capacity even higher. The move comes just ahead of Anthropic's IPO, as both companies fight for B2B dominance in AI. The article OpenAI doubles usage limits on Pro plans as competition with Anthropic heats up appeared first on The Decoder.

[AI TOOL] SOURCE: the-decoder.com | VEL: 85/100

Experts adopt AI agent's random opinions and think they're their own, Anthropic study finds

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

An experiment by Anthropic with 50 experts found that people often followed their AI agent's randomly assigned preference on questions with no clearly right answer, and most thought the final result was their own. Handing a data analysis to an AI agent saves time and manual work, but it raises the question of who actually makes the decisions that shape the final recommendation. An Anthropic study suggests the agent makes most of them, often without the experts noticing. In the experiment , 50 experts in causal inference worked on a realistic analysis task where they examined observational data from a fictional farming cooperative and estimated the productivity gains from a voluntary farm advisory program.

■ SECTION VII: SOVEREIGN AI & WORLDWIDE REGIONAL ECOSYSTEMS

National AI Labs, Regional Models & Sovereign Compute Ecosystems

[SOVEREIGN [DE]] SOURCE: the-decoder.com | VEL: 85/100

OpenAI doubles usage limits on Pro plans as competition with Anthropic heats up

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

OpenAI is doubling usage limits across its paid plans, from Pro 100 to Pro 500. A planned cut to the Pro 200 quota was reversed after user pushback. More efficient models like Terra and Sol should push usable capacity even higher. The move comes just ahead of Anthropic's IPO, as both companies fight for B2B dominance in AI. The article OpenAI doubles usage limits on Pro plans as competition with Anthropic heats up appeared first on The Decoder.

[SOVEREIGN [GB]] SOURCE: aibusiness.com | VEL: 85/100

OpenAI slashes price for model again

ARCH: MOE • **CTX:** 128K • **QUANT:** FP8 • **STACK:** vLLM

OpenAI slashes price for model again marks an architectural milestone in Frontier API foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[SOVEREIGN [CN]] SOURCE: scmp.com | VEL: 98/100

Chinese video AI models beat US rivals in visual quality, but lag on safeguards: HKU study

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Chinese AI video-generation models have surpassed leading US rivals in visual quality, but several fall behind in safety precautions, according to researchers at the University of Hong Kong. ByteDance's Seedance 2.0 Pro took top spot in a ranking of 12 leading artificial intelligence video-generation models, scoring 3.84 out of 5 for visual quality. It outperformed Google DeepMind's Veo 3.1, which scored 3.47, and OpenAI's Sora 2 Pro, on 3.27, in a study led by Zhenhui Jack Jiang, professor of..

[SOVEREIGN [CN]] SOURCE: huggingface.co | VEL: 95/100

[MiniMaxAI] MiniMaxAI/MiniMax-M3.1-Flash (Image Text To Text)

ARCH: MoE / SOTA • **CTX:** 1M • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

MiniMax-M3.1-Flash** is a native multimodal mixture-of-experts model for coding, agentic workflows, computer use, and cybersecurity. It supports text, image, and video inputs with a context window of up to **1M tokens**. Building on MiniMax-M3, M3.1-Flash improves task performance and inference efficiency with sparse attention throughout the model, quantization-aware training, and DSpark speculative decoding. Highlights:** - **Agentic performance:** A generational leap in agentic tasks, coding, and cybersecurity — 76.2 on BankerToolBench, 73.7% on DeepSWE v1.1, 61.7 on OSWorld V2, 29.6 on Agents' Last Exam, and 82.2 on CyberGym. - **Efficient by design:** Sparse attention throughout, quantization-aware training, and DSpark speculative decoding improve inference efficiency.

[SOVEREIGN [CN]] SOURCE: modelscope.cn | VEL: 95/100

[ModelScope] DeepSeek-V4.1-Flash-1017 (484.6B Params) - image-text-to-text

ARCH: 484.6B • **CTX:** 128K • **QUANT:** FP8 • **STACK:** vLLM

[ModelScope] DeepSeek-V4.1-Flash-1017 (484.6B Params) - image-text-to-text marks an architectural milestone in Open-Weights foundation modeling, with a 484.6B parameter footprint and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[SOVEREIGN [CN]] SOURCE: huggingface.co | VEL: 95/100

[deepseek-ai] deepseek-ai/DeepSeek-V4.1-Flash-1017 (Image Text To Text)

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

DeepSeek-V4.1-Flash-1017** is a minor update to **DeepSeek-V4.1-Flash**, delivering substantial improvements across reasoning, coding, agentic tasks, and visual benchmarks. Alongside these gains, the model reduces task execution time and output length, enabling users to complete tasks faster and at lower cost and advancing the performance–cost Pareto frontier across multiple agent harnesses. | Benchmark | DS-V4.1-Flash-1017 | DS-V4.1-Flash | K3 | GLM-5.3 | GPT-5.6 Sol | Opus 5 | | :--- | :---: | :---: | :---: | :---: | :---: | :---: | | | | | | | | | HLE w/o tools | 42.6 (44.0†) | 36.8 (39.1†) | 43.5 | 42.0† | 44.5 | 56.3 | | Terminal-Bench 4.0 | 42.3 | 31.2 | 12.6 | 37.9 | 39.9 |.

■ SECTION VIII: AI HARDWARE, COMPUTE CLUSTERS & SILICON

GPU Clusters, Custom Accelerators, Datacenter Infrastructure & Silicon

SILICON / SYSTEM	SOURCE	COMPUTE SPECS	VELOCITY	BENCHMARK / TELEMETRY
Chinese AI chipmaker Iluvatar CoreX says new	scmp.com	High Velocity	88/100	Chinese artificial intelligence chipmaker Iluvatar CoreX claims that its new Tiangai 300 graphics processing unit (GPU) can match.
EditAgent: Rethinking Agentic Image Editing	huggingface.co	High Velocity	93/100	Agentic image editors decompose complex requests into steps, but when a step fails, they mostly retry with the.
Same kernels, different paths. vLLM and SGLan	X / Twitter: AI	High Velocity	97/100	Source: Twitter/X | Author: X / Twitter Same kernels, different paths. vLLM and SGLang both run DeepSeek's open.
Beyond matching DeepInfer's performance, we a	X / Twitter: AI	High Velocity	97/100	Source: Twitter/X | Author: @SemiAnalysis_ Beyond matching DeepInfer's performance, we also put vLLM to the test on @SemiAnalysis_.
Baseten Releases joey-1, a 2B Decision Model	marktechpost.com	High Velocity	87/100	Baseten has released joey-1, a 2B model from its in-house research lab, Base Labs. The joey-1 decision model.
Prime Intellect Releases Prime Mega MoE: Trai	marktechpost.com	High Velocity	85/100	Prime Intellect has released Prime Mega MoE, a training extension of DeepSeek’s fused Mega MoE kernel.

[AI HARDWARE] SOURCE: scmp.com | VEL: 88/100

Chinese AI chipmaker Iluvatar CoreX says new GPU matches Nvidia's H200 as chip gap closes

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Chinese artificial intelligence chipmaker Iluvatar CoreX claims that its new Tiangai 300 graphics processing unit (GPU) can match and even surpass Nvidia's H200 in real-world AI workloads, highlighting efforts by domestic players to close the technology gap with the US chip giant. In a live demonstration for mainland Chinese media outlets on Wednesday, the Shanghai-based firm showed the Tiangai 300 outperforming the H200 in complex multimodal tasks, including text-to-image and text-to-video...

[AI HARDWARE] SOURCE: huggingface.co | VEL: 93/100

EditAgent: Rethinking Agentic Image Editing

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Agentic image editors decompose complex requests into steps, but when a step fails, they mostly retry with the same narrow set of tools and a judge that scores each output by eye directly. We argue that such agents are bounded by two spaces: the action space of what the agent can do, and the observation space of what its judge can see. On our expert-annotated benchmark, a general-purpose editor yields no usable result in over 40% of the cases, and no judge is reliable across edit types. We present EditAgent, which expands both spaces. It routes each step to general editors and specialized methods, and attaches to each candidate a verification record.

[AI HARDWARE] SOURCE: X / Twitter: AI Develope | VEL: 97/100

Same kernels, different paths. vLLM and SGLang both run DeepSeek's open kernels. Both land within a few % of DeepInfer. But the work they highlight is different. vLLM found the cost was in the experts, not attention, so it spread experts across more GPUs an

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** vLLM

Source: Twitter/X | Author: X / Twitter Same kernels, different paths. vLLM and SGLang both run DeepSeek's open kernels. Both land within a few % of DeepInfer. But the work they highlight is different. vLLM found the cost was in the experts, not attention, so it spread experts across more GPUs an Same kernels, different paths. vLLM and SGLang both run DeepSeek's open kernels. Both land within a few % of DeepInfer. But the work they highlight is different.

[AI HARDWARE] SOURCE: X / Twitter: AI Develope | VEL: 97/100

Beyond matching DeepInfer's performance, we also put vLLM to the test on @SemiAnalysis_ AgentX with real-world agentic workloads. At DeepSeek's off-peak pricing, vLLM delivers $38.96/GPU-hour revenue, with 954K tokens/s/GPU total throughput (cache

ARCH: MoE / SOTA • **CTX:** 954K • **QUANT:** Native / FP16 • **STACK:** vLLM

Source: Twitter/X | Author: @SemiAnalysis_ Beyond matching DeepInfer's performance, we also put vLLM to the test on @SemiAnalysis_ AgentX with real-world agentic workloads. At DeepSeek's off-peak pricing, vLLM delivers $38.96/GPU-hour revenue, with 954K tokens/s/GPU total throughput (cache Beyond matching DeepInfer's performance, we also put vLLM to the test on @SemiAnalysis_ AgentX with real-world agentic workloads. At DeepSeek's off-peak pricing, vLLM delivers $38.96/GPU-hour revenue, with 954K tokens/s/GPU total throughput (cache.

[AI HARDWARE] SOURCE: marktechpost.com | VEL: 87/100

Baseten Releases joey-1, a 2B Decision Model at $0.01 per 1M Tokens

ARCH: 2B • **CTX:** 1M • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Baseten has released joey-1, a 2B model from its in-house research lab, Base Labs. The joey-1 decision model returns a label and a calibrated probability instead of generated text. It targets high-volume routing, classification and guardrail checks. Baseten prices it at $0.01 per million input tokens, with open weights on Hugging Face under MIT. Size: 2B parameters, post-trained from Qwen3.5-2B. 16,384-token max sequence (Baseten says 262k on dedicated deployments).

■ SECTION IX: AUTONOMOUS AGENTS, MULTI-AGENT SWARMS & ROBOTICS

Agent Orchestration, Computer-Use Primitives, Action Models & Robotics

AGENT / FRAMEWORK	SOURCE	PROTOCOL	VELOCITY	CORE CAPABILITY DOMAIN
AgiBot Unveils GO-3 Preview Robot Model, Tops	pandaily.com	PyTorch / ONNX	85/100	AgiBot launched GO-3 Preview on October 29, an embodied AI model that the company says can keep learning.
Experts adopt AI agent's random opinions and	the-decoder.com	PyTorch / ONNX	85/100	An experiment by Anthropic with 50 experts found that people often followed their AI agent's randomly assigned preference.
Gemini Robotics Embodied Reasoning 2 is now a	cloud.google.com	PyTorch / ONNX	95/100	Earlier this year, we introduced Gemini Robotics 2 — our next-generation intelligence layer enabling robots to think, act.
ThinkWorld: World Imagination Reinforces Chai	arxiv.org	PyTorch / ONNX	90/100	Vision-Language-Action (VLA) models have made substantial progress in robotic manipulation, but their generalization remains constrained by the diversity.
Evaluating DeepSeek V4.1 Flash-1017 for loss	huggingface.co	PyTorch / ONNX	90/100	As frontier language models grow increasingly capable and autonomous, independent third-party evaluations are critical for characterizing their systemic.
Donating Agent Skills to the Agentic AI Found	simonwillison.ne	PyTorch / ONNX	85/100	Donating Agent Skills to the Agentic AI Foundation News today from Anthropic: We're committed to keeping Agent Skills.

[AGENT SYSTEM] SOURCE: pandaily.com | VEL: 85/100

AgiBot Unveils GO-3 Preview Robot Model, Tops Three Leaderboards and Plans Open Source

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

AgiBot launched GO-3 Preview on October 29, an embodied AI model that the company says can keep learning at scale across pre-training, post-training and after deployment. AgiBot plans to open-source the full model by the end of 2026. GO-3 Preview is built on what AgiBot calls a WRAM, or World Reasoning Action Model, architecture.

[AGENT SYSTEM] SOURCE: the-decoder.com | VEL: 85/100

Experts adopt AI agent's random opinions and think they're their own, Anthropic study finds

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

An experiment by Anthropic with 50 experts found that people often followed their AI agent's randomly assigned preference on questions with no clearly right answer, and most thought the final result was their own. Handing a data analysis to an AI agent saves time and manual work, but it raises the question of who actually makes the decisions that shape the final recommendation. An Anthropic study suggests the agent makes most of them, often without the experts noticing. In the experiment , 50 experts in causal inference worked on a realistic analysis task where they examined observational data from a fictional farming cooperative and estimated the productivity gains from a voluntary farm advisory program.

[AGENT SYSTEM] SOURCE: cloud.google.com | VEL: 95/100

Gemini Robotics Embodied Reasoning 2 is now available on Google Cloud

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Earlier this year, we introduced Gemini Robotics 2 — our next-generation intelligence layer enabling robots to think, act, and adapt in the physical world. At the core of this breakthrough is Gemini Robotics Embodied Reasoning 2 (ER 2) , a specialized Gemini model that serves as a robot’s high-level brain. Working across any robot embodiment — from stationary arms to full humanoids — Gemini Robotics ER 2 combines advanced spatial and temporal video understanding with long-horizon planning and built-in safety guardrails. This enables robots to converse naturally with humans, orchestrate low-level action models and hardware APIs, execute multi-minute tasks involving hundreds of decisions, and collaborate across multi-robot teams.

[AGENT SYSTEM] SOURCE: arxiv.org | VEL: 90/100

ThinkWorld: World Imagination Reinforces Chain-of-Dynamics Reasoning

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Vision-Language-Action (VLA) models have made substantial progress in robotic manipulation, but their generalization remains constrained by the diversity and coverage of action-labeled robot demonstrations. Human videos provide much broader behavioral and environmental experience, while learning from them faces a fundamental trade-off between scalable data and explicit action grounding.

[AGENT SYSTEM] SOURCE: huggingface.co | VEL: 90/100

Evaluating DeepSeek V4.1 Flash-1017 for loss of control and harmful manipulation risks

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

As frontier language models grow increasingly capable and autonomous, independent third-party evaluations are critical for characterizing their systemic safety risks. While closed-source releases are routinely accompanied by comprehensive system cards, open-weight models receive less scrutiny regarding frontier risks. To address this gap, we present an independent safety evaluation of DeepSeek V4.1 Flash-1017, the latest version of V4.1 Flash that departs from V4 with a new architecture, scaled-up training, and strong agentic capabilities. Drawing on the EU General-Purpose AI Code of Practice, we evaluate the model across 34 benchmarks, including 5 in-house suites, spanning evaluation integrity, loss of control, and harmful manipulation. Conforming to the AEF-1 standard, we compare V4.1...

■ SECTION X: GLOBAL AI COMMUNITY WIRE & OVERFLOW DIGEST

High-velocity AI stories from today's sweep that didn't fit earlier sections

COLOPHON & SENSOR METHODOLOGY: The Global AI Gazette is compiled autonomously. Data is aggregated across authoritative global AI sources including Hugging Face, GitHub Trending, arXiv, and leading AI research lab feeds. Edition #2451 • 212 items processed • Friday, October 30, 2026 • 09:03 UTC. All rights reserved.