

# THE GLOBAL AI GAZETTE

Worldwide AI Ecosystem Intelligence • Open Source • Research • Models • Tools • Sovereign AI

Tuesday, September 15, 2026 • 22:09 UTC

EDITION NO. 2238

160 STORIES ANALYZED

■ LEAD STORY

## Intern-S2-397B (multimodal, reasoning, coding, and scientific agent capabilities)

Velocity 96/100 • Impact 88/100 • Source: reddit.com

Intern-S2-397B (multimodal, reasoning, coding, and scientific agent capabilities) marks an architectural milestone in Frontier API foundation modeling, with a 397B parameter footprint and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

**INNOVATION:** Architecture: Reasoning LLM & Multi-Step Chain-of-Thought  
**ACTION:** Deployment: Checkpoints accessible on Hugging Face; compatible with standard runtimes.

■ SECONDARY ANCHOR

## Running Qwen 3.8 next on 16vram+32ram - A useful/fun post for the gpu poors

Hello Reddit. Posting this for fun. I thought it was a lonely and silly journey to set up Qwen 3.8 Next on a system that doesn't really run it properly—it was a challenge that might help the community. I have yet to benchmark this specific REAP version versus Qwen 3.8 27B QK4, but my assumption is that it will do much better, despite the hemorrhaged...

## ■ TOP FLASH BULLETINS

**[VEL 96]** DeepSeek V4.1 Flash beats Astra on AA's new benchmark — AA shipped a new benchmark last week as part of the Intelligence Index v4.3 update — a brand-new private eval that replaces  $\tau^3$ . Astra was...

**[VEL 96]** Dear 24G owners, try VLLM you might be able to run Qwen3.8 27B INT4, 144K FP8 KV on RTX 3090 with better speed. (TLDR VLLM AOT) — VLLM Benchmark: Prefill, Prompt processing - avg, 871.93 tok/s (3 hours constant running xhigh) - 10K prompt, 1000.26 tok/s (16 runs) - 90K prompt, 743,59...

**[VEL 96]** ANOTHER researcher accuses OpenAI of training on conversations and then claiming a breakthrough — ANOTHER researcher accuses OpenAI of training on conversations and then claiming a breakthrough reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology...

**[VEL 96]** Deepseek V4.1 Flash Release Video [Made with Deepseek V4.1 Flash] — I like to benchmark new models that come out on motion videos. So here's a test I did for deepseek v4.1 flash. And I have...

**[VEL 96]** guide to using reasoning\_effort on deepseek v4.1 flash — guide to using reasoning\_effort on deepseek v4.1 flash marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation....

## SECTION II: EXECUTIVE AI STRATEGIC BRIEFING

Strategic Horizons & Action Directives for Enterprise Leadership

**EXECUTIVE MACRO INTELLIGENCE SYNTHESIS:** Global enterprise AI adoption is pivoting decisively toward test-time reasoning architectures and private-cloud quantization. Technology leadership must actively balance proprietary frontier model APIs with sovereign, open-weight deployments to reduce token expenditure while strictly sandboxing autonomous agent tool-calling boundaries.

| VECTOR                           | DEVELOPMENT                              | SOURCE          | DIRECTIVE                                                              |
|----------------------------------|------------------------------------------|-----------------|------------------------------------------------------------------------|
| Model Sourcing & Licensing       | Controlling Reasoning Effort in LLMs     | magazine.sebast | Audit open-weights licensing vs proprietary APIs; evaluate Controlling |
| Compute & Infrastructure CapEx   | Deploying Qwen3.8-2.4T-A95B on Amazon Sa | aws.amazon.com  | Review GPU cluster allocation and power envelopes; benchmark vLLM hard |
| Agentic Autonomy & Governance    | 3D viz of how Deepseek Flash v4.1 is dif | reddit.com      | Implement deterministic sandboxes for Deepseek autonomous tool executi |
| Inference Latency & Quantization | Another Qwen3.8-27b Appreciation Post    | reddit.com      | Benchmark KV-cache compression (FP8/INT4/GGUF) and modern inference en |
| Open-Source Supply Chain         | R9V Update: now ~100 tok/s in TG on Qwen | reddit.com      | Inspect upstream repository dependencies; audit tokenizer code, weight |
| Data Residency & Sovereignty     | Decided to build a game, and test the ce | reddit.com      | Verify compliance with sovereign AI frameworks and regional data resid |

1. **Controlling Reasoning Effort in LLMs:** Audit open-weights licensing vs proprietary APIs; evaluate Controlling Reasoning Effort parameter efficiency and commercial distribution terms.
2. **Deploying Qwen3.8-2.4T-A95B on Amazon SageMaker HyperPod with vLLM:** Review GPU cluster allocation and power envelopes; benchmark vLLM hardware efficiency to optimize cost per token.
3. **3D viz of how Deepseek Flash v4.1 is different from a typical decode only transformer:** Implement deterministic sandboxes for Deepseek autonomous tool execution, strict rate limiting, and human-in-the-loop validation.
4. **Another Qwen3.8-27b Appreciation Post:** Benchmark KV-cache compression (FP8/INT4/GGUF) and modern inference engines for Another Qwen3.8-27b Apprecia against TTFT SLAs.
5. **R9V Update: now ~100 tok/s in TG on Qwen3.8 Flash Next IQ4\_XS on x2 R9700 + 128GB RAM. Fixed crashes with n-gram SSD streaming, improved diagnostics, plus pinned images. Q4\_K\_XL now supported, 50 tok/s TG.:** Inspect upstream repository dependencies; audit tokenizer code, weights provenance, and pinned runtime releases for R9V Update: now ~100 tok/s i.
6. **Decided to build a game, and test the ceiling of Qwen3.8 27b:** Verify compliance with sovereign AI frameworks and regional data residency requirements for Decided to build a game, and deployments.

# SECTION III: TRENDING OPEN-SOURCE AI & GITHUB INNOVATIONS

Top GitHub Repositories, Developer Velocity & Architecture Dissections

| PROJECT / REPO                                 | SOURCE        | STACK  | VELOCITY | CORE ARCHITECTURAL FOCUS                                                                                                         |
|------------------------------------------------|---------------|--------|----------|----------------------------------------------------------------------------------------------------------------------------------|
| JuliusBrussee/caveman: why use many token w    | github.com    | Python | 96/100   | Your AI coding agent bills by the word and writes like it knows that. Caveman make it stop..                                     |
| Merge pull request #1596 from QwenLM/update_0  | github.com    | Python | 96/100   | Qwen Chat   Hugging Face   ModelScope   Paper   Blog.                                                                            |
| DeepSeek-V3 Open Model v1.0.0: This release i  | github.com    | Python | 96/100   | We present DeepSeek-V3, a strong Mixture-of-Experts (MoE) language model with 671B total parameters with 37B activated for each. |
| DeepSeek reportedly advances Shanghai IPO pla  | digitimes.com | Python | 88/100   | DeepSeek reportedly advances Shanghai IPO plans as it cuts Flash model API prices marks an architectural milestone in.           |
| Muse, Meta's New Personal AI Agent, Needs You  | wired.com     | Python | 88/100   | Meta announced Tuesday the release of Muse, a personal AI agent that people can message to automate digital.                     |
| vLLM Inference Engine v0.29.1: Dual-key gumble | github.com    | Python | 88/100   | Easy, fast, and cheap LLM serving for everyone   Documentation   Blog   Paper   Twitter/X   User.                                |

[TRENDING REPO] SOURCE: [github.com](#) | VEL: 96/100

**JuliusBrussee/caveman: why use many token when few token do trick — Claude Code skill that cuts 65% o...**

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Your AI coding agent bills by the word and writes like it knows that. Caveman make it stop. #1 on GitHub Trending · July 2026 · #1 Repository of the Day on Trendshift · April 2026 #1 on Hacker News · 904 points · 366 comments · #8 Product of the Day on Product Hunt Cited in CAVEWOMAN , an Adobe Research paper that measured caveman-style output cutting cost 1.4 to 2.4×, up to 3× Tested by JetBrains on 86 real coding tasks: "costs you nothing measurable in quality"

[TRENDING REPO] SOURCE: [github.com](#) | VEL: 96/100

**Merge pull request #1596 from QwenLM/update\_0806**

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Qwen Chat | Hugging Face | ModelScope | Paper | Blog Documentation Demo | WeChat ( ) | Discord Visit our Hugging Face or ModelScope organization (click links above), search checkpoints with names starting with `Qwen3-` or visit the Qwen3 collection, and you will find all you need! Enjoy! To learn more about Qwen3, feel free to read our documentation [EN|ZH].

[TRENDING REPO] SOURCE: [github.com](#) | VEL: 96/100

**DeepSeek-V3 Open Model v1.0.0: This release is created for archival purposes and DOI generation.**

ARCH: MOE • CTX: 788M • QUANT: Native / FP16 • STACK: PyTorch / ONNX

We present DeepSeek-V3, a strong Mixture-of-Experts (MoE) language model with 671B total parameters with 37B activated for each token. To achieve efficient inference and cost-effective training, DeepSeek-V3 adopts Multi-head Latent Attention (MLA) and DeepSeekMoE architectures, which were thoroughly validated in DeepSeek-V2. Furthermore, DeepSeek-V3 pioneers an auxiliary-loss-free strategy for load balancing and sets a multi-token prediction training objective for stronger performance. We pre-train DeepSeek-V3 on 14.8 trillion diverse and high-quality tokens, followed by Supervised Fine-Tuning and Reinforcement Learning stages to fully harness its capabilities. Comprehensive evaluations reveal that DeepSeek-V3 outperforms other open-source models and achieves performance comparable to leading closed-source models. Despite its excellent performance, DeepSeek-V3 requires only 2.788M H800 GPU hours for its full training..

[TRENDING REPO] SOURCE: [digitimes.com](#) | VEL: 88/100

**DeepSeek reportedly advances Shanghai IPO plans as it cuts Flash model API prices**

ARCH: MOE • CTX: 128K • QUANT: FP8 • STACK: vLLM

DeepSeek reportedly advances Shanghai IPO plans as it cuts Flash model API prices marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[TRENDING REPO] SOURCE: [wired.com](#) | VEL: 88/100

**Muse, Meta's New Personal AI Agent, Needs You to Trust It**

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Meta announced Tuesday the release of Muse, a personal AI agent that people can message to automate digital tasks in a secure cloud environment, all while relying on security and privacy that the company says is “built into it” from the start. Meta says Muse is rolling out today for iOS and Android users in a dedicated Muse app, as well as the website Muse.ai. The company also says users can directly message Muse in WhatsApp to interact with the agent. Meta says users of its AI glasses will soon be able to interact with its Muse agent as well.

SECTION IV: FRONTIER FOUNDATION MODELS & REASONING  
BREAKTHROUGHS

Reasoning Architectures, Open Weights & Capability Benchmarks

| FOUNDATION MODEL                              | SOURCE           | SIZE / CTX        | IMPACT | ARCHITECTURAL HIGHLIGHT                                                                                                                      |
|-----------------------------------------------|------------------|-------------------|--------|----------------------------------------------------------------------------------------------------------------------------------------------|
| DeepSeek V4.1 Flash beats Astra on AA's new b | reddit.com       | MoE / SOTA / 128K | 88/100 | AA shipped a new benchmark last week as part of the Intelligence Index v4.3 update — a brand-new.                                            |
| ANOTHER researcher accuses OpenAI of training | reddit.com       | MoE / SOTA / 128K | 88/100 | ANOTHER researcher accuses OpenAI of training on conversations and then claiming a breakthrough reflects the rapid acceleration of.          |
| Deepseek V4.1 Flash Release Video [Made with  | reddit.com       | MoE / SOTA / 128K | 88/100 | I like to benchmark new models that come out on motion videos. So here's a test I did.                                                       |
| guide to using reasoning_effort on deepseek v | reddit.com       | MOE / 128K        | 88/100 | guide to using reasoning_effort on deepseek v4.1 flash marks an architectural milestone in Open-Weights foundation modeling, engineered for. |
| GPT-6 Astra, Looped Transformers, and Hidden  | magazine.sebasti | MoE / SOTA / 128K | 88/100 | A lot has happened in the last few weeks. I am sure that OpenAI's GPT-6 Astra is top.                                                        |
| 3D viz of how Deepseek Flash v4.1 is differen | reddit.com       | MOE / 128K        | 68/100 | 3D viz of how Deepseek Flash v4.1 is different from a typical decode only transformer marks an architectural.                                |

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 88/100

DeepSeek V4.1 Flash beats Astra on AA's new benchmark

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

AA shipped a new benchmark last week as part of the Intelligence Index v4.3 update — a brand-new private eval that replaces  $\tau^3$ . Astra was farming a ton of points on it and used those to get even with Fable, but... looks like we have a new king. So they changed the index twice in three days to make Astra look not-quite-worse than Fable, and then a random guy quietly took first place on it.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 88/100

ANOTHER researcher accuses OpenAI of training on conversations and then claiming a breakthrough

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

ANOTHER researcher accuses OpenAI of training on conversations and then claiming a breakthrough reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology development. Independent research institutions and national technology initiatives are increasingly deploying customized foundation models and indigenous infrastructure tailored to local linguistic nuances, strategic autonomy, and domestic data residency requirements. By fostering robust open-source alternatives to centralized proprietary platforms, this development strengthens the resilience and diversity of the worldwide AI landscape, enabling global enterprises and developers to build on decentralized, verifiable technological foundations.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 88/100

Deepseek V4.1 Flash Release Video [Made with Deepseek V4.1 Flash]

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

I like to benchmark new models that come out on motion videos. So here's a test I did for deepseek v4.1 flash. And I have to say flash has probably graduated from being a Luna class model to nearly an Opus class model with this release, at least with motion videos. Prev.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 88/100

guide to using reasoning\_effort on deepseek v4.1 flash

ARCH: MOE • CTX: 128K • QUANT: FP8 • STACK: vLLM

guide to using reasoning\_effort on deepseek v4.1 flash marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[FRONTIER MODEL] SOURCE: magazine.sebastianraschk | IMPACT: 88/100

GPT-6 Astra, Looped Transformers, and Hidden Reasoning

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

A lot has happened in the last few weeks. I am sure that OpenAI's GPT-6 Astra is top of mind for everyone right now. In particular, thoughts on its performance, the looped transformer/recurrent depth aspects, and rumors that Astra is "hiding" its reasoning trace (i.e., chain of thought).So, in this article, I want to start with some brief impressions of Astra and some thoughts on where all this is headed.

## ■ SECTION V: TOP AI RESEARCH PAPERS & ARXIV BREAKTHROUGHS

*Reasoning Paradigms, Multimodal Architectures & Algorithmic Design*

[AI RESEARCH] SOURCE: [magazine.sebastianraschke.com](https://magazine.sebastianraschke.com/) | VEL: 88/100

### Controlling Reasoning Effort in LLMs

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

It has been almost two years since OpenAI released o1, a model that popularized the idea of LLM-based reasoning models. DeepSeek-R1 followed about four months later, together with details of a reinforcement learning with verifiable rewards (RLVR) recipe to train such reasoning models. Last week, OpenAI released the GPT-5.6 model family. It comes in three sizes, each with roughly five or six reasoning-effort settings. Figure 1: The GPT 5.6 Sol model with different reasoning effort settings.

[AI RESEARCH] SOURCE: [arxiv.org](https://arxiv.org/) | VEL: 88/100

### Guardrailed Meta-Agent Loops: Stress-Testing Policy Pinning, Budget Bounds, and Crash Recovery

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

arXiv:2609.12216v1 Announce Type: new Abstract: Self-improving agent workflows create an audit problem when the same controller can change both its behavior and the conditions under which that behavior is judged. We present GuardrailLoop, a simulation-based testbed that makes three operational contracts jointly testable: preservation of human-defined policy, compute accounting at every recorded execution prefix, and recovery of a specified scientific state after crashes.

[AI RESEARCH] SOURCE: [siliconangle.com](https://siliconangle.com/) | VEL: 88/100

### Anthropic researcher's resignation sparks broad AI safety discussion

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

An Anthropic PBC researcher has resigned over concerns that artificial intelligence labs are “gambling with our lives.” Jacob Coxon was part of the company’s AI pretraining team until today. He announced his resignation in a series of X posts that has been viewed millions of times. Coxon wrote that his decision was motivated by concerns over recursive self-improvement, or RSI. That’s a term for a hypothetical future AI with the ability to automatically improve its own capabilities. “These will soon be superhuman systems that can hack anything, revolutionize any field overnight, and acquire real power and resources,” Coxon wrote.

[AI RESEARCH] SOURCE: [arxiv.org](https://arxiv.org/) | VEL: 88/100

### Scanning the Harness: An Empirical Study of Supply-Chain Defects in AI Coding-Agent Configurations

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

AI coding agents such as Claude Code, Cursor, GitHub Copilot, and OpenAI Codex are configured through artifacts developers write and share: instruction files, skills, hooks, MCP server declarations, subagents. This harness is a dependency layer installed from marketplaces and public repositories, running with the developer's privileges, with no lockfile, no install-time check, and no vocabulary for what a component may do.

[AI RESEARCH] SOURCE: [arxiv.org](https://arxiv.org/) | VEL: 88/100

### What is the Difference Between Me and You? Benchmarking the Quality Gap Between Human-Written and AI-Generated Code

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

AI coding assistants are becoming co-authors of production software, yet their evaluation centers on functional correctness, leaving open whether their code differs from human code in the quality dimensions dominating lifecycle cost.

[AI RESEARCH] SOURCE: [arxiv.org](https://arxiv.org/) | VEL: 88/100

### SemVerBench: Benchmarking LLM Comprehension of Version-Constraint Resolution Semantics

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Large language model (LLM) coding agents constantly decide whether a version satisfies a constraint such as  $^1 1.2.3$  or  $\geq 2.0, 1.2$  means  $\geq 1.3.0$ ) traps every model on Cargo (near 60%), and although standard PEP 440 prefix matching is universal, on zero-pad/post-release corner cases GPT-5.1 collapses (0/26) while Claude stays at 97-100% (verified on a 67-item oracle-validated set). Opus significantly outperforms all other models, and Sonnet outperforms the OpenAI models (McNemar).

## ■ ■ SECTION VI: DEVELOPER TOOLS, FRAMEWORKS & AI INFRASTRUCTURE

*Inference Runtimes, Evaluation Frameworks, Vector DBs & Local Execution Engines*

[AI TOOL] SOURCE: [simonwillison.net](https://simonwillison.net) | VEL: 88/100

### OpenAI agents attacked RubyGems back in May

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

OpenAI agents carried out an undisclosed attack on RubyGems is a new bombshell report from Spencer Kitts, Thomas Larsen, and Sydney Von Arx - three of the four authors of the report on the agent attack on disused wikis (previously) last week.

[AI TOOL] SOURCE: [securityweek.com](https://securityweek.com) | VEL: 88/100

### Anthropic Researcher Resigns With Warning About the Dangers of AI Development

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

His posts reached more than 100 million people overnight. Coxon is not the first AI insider to publicly raise such concerns. Both Anthropic and OpenAI have seen high-profile resignations in recent years that were tied to safety concerns. Two current Anthropic employees also responded to Coxon's post in agreement. Sen. Bernie Sanders, a Vermont independent who has called for AI safeguards and regulation, agreed with Coxon's concerns and said he would soon introduce legislation to pause AI development and ban superintelligence. "The very people building this technology admit that it could threaten the future of humanity," Sanders said Wednesday on social media.

[AI TOOL] SOURCE: [thehackernews.com](https://thehackernews.com) | VEL: 88/100

### Anthropic Discloses Fourth AI Hacking Incident Involving Claude Opus 4.6

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Anthropic on Wednesday disclosed a fourth incident in which its artificial intelligence (AI) model broke into real third-party systems, marking the latest in a growing list of cases that have raised concerns about the security risks posed by autonomous AI agents. The AI company said the incident dates back to January 2026 and involved an early version of Claude Opus 4.6 that breached "

[AI TOOL] SOURCE: [reddit.com](https://reddit.com) | VEL: 96/100

### The rhetoric is really heating up!

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

The entire page of the NY Times today above the fold absent one article is AI (the models are just too strong/too dangerous, must be regulated). They forgot to include "Sponsored by OpenAI" at the end of the articles, sure that was just an oversight? This is what the end of a bubble looks like, desperate attempts to get some sort of regulatory capture in place to keep the business model from collapsing in upon itself.

[AI TOOL] SOURCE: [arxiv.org](https://arxiv.org) | VEL: 88/100

### Watermarks Without Verification: AI Text Watermarking After the EU AI Act

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

On August 2, 2026, the obligations of Article 50 of the EU AI Act took effect, requiring generative AI providers to mark the content their systems produce and ensure it can be detected as AI-generated. Days later, Anthropic disclosed that every Claude model released after that date embeds a watermark based on SynthID-Text in all generated text, enabled by default with no user opt-out; Google has deployed SynthID-Text in Gemini since 2024.

[AI TOOL] SOURCE: [arxiv.org](https://arxiv.org) | VEL: 88/100

### A Tool-Augmented, GPT-4 Chatbot for Real-Time Repository Data Analysis

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Software repositories contain vast amounts of data on code contributions, bug reports, and project activities, yet this information remains challenging for non-technical stakeholders and developers to access due to limited expertise in querying repositories. To address this, we introduce a novel chatbot architecture leveraging OpenAI's GPT-4 model for automated extraction and analysis of repository data.

## ■ SECTION VII: SOVEREIGN AI & WORLDWIDE REGIONAL ECOSYSTEMS

*National AI Labs, Regional Models & Sovereign Compute Ecosystems*

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

### Talk me out of buying a 3rd Spark

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Does anyone think the gurus on the DGX Spark forum are going to figure out how to magically fit DeepSeek 4.1 Flash on a 2x cluster, or is it only possible on 3 or 4 Sparks?

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

### React Native ExecuTorch is now up to 92x faster ■■

ARCH: 0.6B • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

In v0.10 we achieved significant speedups over v0.9. The video shows the maximum speedups we measured for specific groups of models. Among LLMs, the biggest gain came from Qwen3 0.6B, which runs over 3x faster on long prompts. For instance segmentation, FastSAM reached speedups of up to 92x! We replaced the monolithic native modules with TypeScript pipelines you can inspect. ■ It runs across all major silicon backends and makes it easier to plug in your very own model.

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

### 3060 12GB vs 4060 ti 16GB

ARCH: 27B • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

I'm currently building my system around 3060s, but I might be able to get a 4060 for a nice deal. At first it seemed like a no brainer, but turns out the 4060 has lower memory bandwidth. In a system that already has 4x 3060 12GBs set up on a threadripper with tensor parallelism (mostly qwen3.8-27b), would it be worth having the 4060 ti 16GB around for the extra 4GB and occasional gaming, or is it just going to slow the rest of the setup down for AI?

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

### Harness doesn't matter

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Harness doesn't matter reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology development. Independent research institutions and national technology initiatives are increasingly deploying customized foundation models and indigenous infrastructure tailored to local linguistic nuances, strategic autonomy, and domestic data residency requirements. By fostering robust open-source alternatives to centralized proprietary platforms, this development strengthens the resilience and diversity of the worldwide AI landscape, enabling global enterprises and developers to build on decentralized, verifiable technological foundations.

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

### What to run at 128GB VRAM?

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Long time lurker, but I'm finally upgrading to 128GB VRAM, and I'm trying to figure out what to run. I had been leaning towards Qwen3.8 Flash-Next at ~Q4, and I generally prefer to not run anything below Q4. But I feel like the reception to Flash-Next has been a bit "meh", so I'm considering GLM 5.3 at ~Q2 or Deepseek 4 Flash at Q2 or Q3. I'm sure I'll try all 3, but I'm really curious what people in the same boat have been doing?

[SOVEREIGN [FR]] SOURCE: [arxiv.org](#) | VEL: 88/100

### IRWOZ 2.0: A Large Language Model-driven Dialogue Dataset for Industrial Robot Conversations

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

IRWOZ has improved industrial human-robot interaction (HRI) dialogue systems through domain-specific annotations. However, its initial version contains substantial noise in dialogue states and utterances, limiting state-tracking accuracy. We introduce IRWOZ 2.0, which addresses these limitations through large language model (LLM) enhanced generation (Mistral/Claude-3.5) and quality refinements.

SECTION VIII: AI HARDWARE, COMPUTE CLUSTERS & SILICON

GPU Clusters, Custom Accelerators, Datacenter Infrastructure & Silicon

| SILICON / SYSTEM                              | SOURCE     | COMPUTE SPECS | VELOCITY | BENCHMARK / TELEMETRY                                                                                                                                        |
|-----------------------------------------------|------------|---------------|----------|--------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Dear 24G owners, try VLLM you might be able t | reddit.com | 871.93 tok/s  | 96/100   | VLLM Benchmark: Prefill, Prompt processing - avg, 871.93 tok/s (3 hours constant running xhigh) - 10K prompt, 1000.26.                                       |
| Decided to build a game, and test the ceiling | reddit.com | High Velocity | 96/100   | This took roughly 5 hours to create, using 2 different configured harnesses, same model. RTX 3090, overclocked +12%.                                         |
| Hoping for Optimized Smarter Upcoming Models  | reddit.com | High Velocity | 96/100   | It's still a dream for many folks to run medium size(30B range) models @ Q8 with Unquantized KVCache.                                                        |
| Running Vision Qwen 3.8 27B on a 16GB Card, t | reddit.com | High Velocity | 96/100   | I am just sharing my config for Qwen 3.8 27b that fits on a 5060TI, what is cool.                                                                            |
| Substrate-Aware AI Agents: Execution Context  | arxiv.org  | High Velocity | 88/100   | Autonomous AI agents increasingly select actions in environments whose memory, execution-time, runtime, compute, and operational constraints determine what. |
| llama.cpp High-Performance LLM Engine: b10899 | github.com | High Velocity | 88/100   | vulkan: small M matrix optimizations for qwen (#28457) vulkan: optimize m=1 mul_mat by swapping A/B vulkan: Improve small.                                   |

[AI HARDWARE] SOURCE: reddit.com | VEL: 96/100  
**Dear 24G owners, try VLLM you might be able to run Qwen3.8 27B INT4, 144K FP8 KV on RTX 3090 with better speed. (TLDR VLLM AOT)**  
ARCH: 27B • CTX: 144K • QUANT: INT4 • STACK: VLLM  
VLLM Benchmark: Prefill, Prompt processing - avg, 871.93 tok/s (3 hours constant running xhigh) - 10K prompt, 1000.26 tok/s (16 runs) - 90K prompt, 743,59 tok/s (16 runs) Decode, tok gen - avg, 38.39 tok/s (3 hours constant running xhigh) - 10K, 42.3 tok/s (16 runs) - 90K, 34 tok/s (16 runs) Preamble: I am on WSL2. Running the 27B Q5 UD GGUF through llama.cpp with 81,920 context plus MTP gives me around 25-30 tok/s.

[AI HARDWARE] SOURCE: reddit.com | VEL: 96/100  
**Decided to build a game, and test the ceiling of Qwen3.8 27b**  
ARCH: 27B • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX  
This took roughly 5 hours to create, using 2 different configured harnesses, same model. RTX 3090, overclocked +12% gain (MSI Afterburner), Q4KM - built this for fun, will be throwing it on GitHub, opensource for people to get an idea of a project created to the near ceiling of performance & capability for q3.8 27b. & also maybe ya'll can contribute to the game only iterating locally. It would be a fun little experiment.

[AI HARDWARE] SOURCE: reddit.com | VEL: 96/100  
**Hoping for Optimized Smarter Upcoming Models .... Like DeepSeek-V4.1-Flash( KVCache + Engram) in Small/Medium/Big sizes**  
ARCH: 30B • CTX: 256K • QUANT: Native / FP16 • STACK: PyTorch / ONNX  
It's still a dream for many folks to run medium size(30B range) models @ Q8 with Unquantized KVCache (256K Context) on their GPUs. It would be awesome to have DeepSeek-V4.1-Flash's KVCache + Engram for all Upcoming models. Even for big models. Engram - Heard that approximately 1/3-1/2 of Model size. Might come in different size range too. So 10-15 GB for 30B models. Here few models with approximate numbers. Current models in Odd rows & Future/Fictional models in even rows(Bold).

[AI HARDWARE] SOURCE: reddit.com | VEL: 96/100  
**Running Vision Qwen 3.8 27B on a 16GB Card, the config (45tk/s).**  
ARCH: 27B • CTX: 85K • QUANT: GGUF • STACK: PyTorch / ONNX  
I am just sharing my config for Qwen 3.8 27b that fits on a 5060TI, what is cool about this is that you can even get vision! and a 85K context (I have 1.5gb of headroom for more context or a better quant) Model: IQ3\_XXS-mtp from Using beellama Config used: [\*] model = ..\llm-models\Qwen3.8-27B-GSQ-RCO-IQ3\_XXS-mtp.gguf mmproj = ..\llm-models\mmproj-Qwen3.8-27B-BF16.gguf image-min-tokens = 256.

[AI HARDWARE] SOURCE: arxiv.org | VEL: 88/100  
**Substrate-Aware AI Agents: Execution Context as a First-Class Input**  
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX  
Autonomous AI agents increasingly select actions in environments whose memory, execution-time, runtime, compute, and operational constraints determine what counts as a suitable plan. We call the absence of this execution context from an agent's planning state substrate blindness. We test this general proposition through numerical code generation, where selected implementation choices and operational consequences are directly observable.

## SECTION IX: AUTONOMOUS AGENTS, MULTI-AGENT SWARMS & ROBOTICS

Agent Orchestration, Computer-Use Primitives, Action Models & Robotics

| AGENT / FRAMEWORK                             | SOURCE         | PROTOCOL       | VELOCITY | CORE CAPABILITY DOMAIN                                                                                                               |
|-----------------------------------------------|----------------|----------------|----------|--------------------------------------------------------------------------------------------------------------------------------------|
| Deploying Qwen3.8-2.4T-A95B on Amazon SageMak | aws.amazon.com | vLLM           | 88/100   | On August 12, 2026, Alibaba's Qwen team released Qwen3.8-2.4T-A95B. This is the first time a Qwen-Max-class model has.               |
| unslothai/unsloth: Local UI to run and train  | github.com     | MLX            | 96/100   | Unsloth is the first desktop app to run and train models. ## ■ Get started Download the native.                                      |
| NousResearch/hermes-agent: The agent that gro | github.com     | PyTorch / ONNX | 96/100   | # Hermes Agent ■ Hermes Agent   Hermes Desktop The self-improving AI agent built by Nous Research.** It's.                           |
| langgenius/dify: Build Agentic workflows, RAG | github.com     | PyTorch / ONNX | 96/100   | Dify Cloud · Self-hosting · Documentation · Dify edition overview Dify is an open-source LLM app development platform..              |
| code-yeongyu/oh-my-openagent: OmO: Just type  | github.com     | PyTorch / ONNX | 96/100   | > [!NOTE] > **OmO for Codex is try LazyCodex** > > We loved Anthropic models enough to get.                                          |
| langchain-ai/langchain: The agent engineering | github.com     | PyTorch / ONNX | 96/100   | LangChain is a framework for building agents and LLM-powered applications. It helps you chain together interoperable components and. |

[AGENT SYSTEM] SOURCE: [aws.amazon.com](#) | VEL: 88/100  
**Deploying Qwen3.8-2.4T-A95B on Amazon SageMaker HyperPod with vLLM**  
ARCH: MoE / SOTA • CTX: 262K • QUANT: Native / FP16 • STACK: vLLM

On August 12, 2026, Alibaba's Qwen team released Qwen3.8-2.4T-A95B. This is the first time a Qwen-Max-class model has been made available as open weights. With 2.4 trillion total parameters (95 billion activated per token), a hybrid linear-plus-full-attention architecture, and native context up to 262K tokens (extensible to 1M), Qwen3.8 targets the most demanding agentic and reasoning workloads. These include multi-step coding, long-horizon planning, and autonomous tool use.

[AGENT SYSTEM] SOURCE: [github.com](#) | VEL: 96/100  
**unslothai/unsloth: Local UI to run and train LLMs and diffusion models. Supports GGUF, MLX, Qwen3.8...**  
ARCH: MoE / SOTA • CTX: 128K • QUANT: GGUF • STACK: MLX

Unsloth is the first desktop app to run and train models. ## ■ Get started Download the native Unsloth Desktop app for your operating system: Platform Link Windows Download macOS Download Linux / Ubuntu (deb) Download Linux (AppImage) Download ## ■ Features Unsloth works on \*\*Windows, Linux, WSL\*\* and \*\*macOS\*\*. We support \*\*Multi GPU setups, NVIDIA, AMD, Intel GPUs, CPUs\*\* and the \*\*Vulkan\*\* backend.

[AGENT SYSTEM] SOURCE: [github.com](#) | VEL: 96/100  
**NousResearch/hermes-agent: The agent that grows with you...**  
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

# Hermes Agent ■ Hermes Agent | Hermes Desktop The self-improving AI agent built by Nous Research.\*\* It's the only agent with a built-in learning loop — it creates skills from experience, improves them during use, nudges itself to persist knowledge, searches its own past conversations, and builds a deepening model of who you are across sessions. Run it on a \$5 VPS, a GPU cluster, or serverless infrastructure that costs nearly nothing when idle. It's not tied to your laptop — talk to it from Telegram while it works on a cloud VM. Use any model you want — Nous Portal, OpenRouter, OpenAI, your own endpoint, and many others.

[AGENT SYSTEM] SOURCE: [github.com](#) | VEL: 96/100  
**langgenius/dify: Build Agentic workflows, RAG pipelines, with rich AI model and tool support on o...**  
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Dify Cloud · Self-hosting · Documentation · Dify edition overview Dify is an open-source LLM app development platform. Its intuitive interface combines AI workflow, RAG pipeline, agent capabilities, model management, observability features (including Opik, Langfuse, and Arize Phoenix) and more, letting you quickly go from prototype to production.

[AGENT SYSTEM] SOURCE: [github.com](#) | VEL: 96/100  
**code-yeongyu/oh-my-openagent: OmO: Just type "mass ulw" keyword with your prompt. Now you are the master of gr...**  
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

> [!NOTE] > \*\*OmO for Codex is try LazyCodex\*\* > > We loved Anthropic models enough to get blocked. Now we are backing Codex. > If you are an OmO fan but the setup felt like too much, use LazyCodex. OmO for Codex has shipped: > > Learn more at lazycodex.ai. > [!TIP] > \*\*Building in Public\*\* > > The maintainer builds and maintains oh-my-openagent in real-time with Jobdori, an AI assistant running on a heavily customized fork of OpenClaw. > Every feature, every fix, every issue triage — live in our Discord. > > [( > > \*\*→ Watch it happen in #building-in-public\*\* > [!NOTE] > > [( > > \*\*OmO is maintained by.

## ■ SECTION X: GLOBAL AI COMMUNITY WIRE & OVERFLOW DIGEST

*High-velocity AI stories from today's sweep that didn't fit earlier sections*

[OVERFLOW] SOURCE: [thehackernews.com](https://thehackernews.com) | VEL: 88/100

### Google Play Early Access Abused to Push Thousands of Deceptive Android Apps

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Bad actors are misusing Google Play's Early Access program to push deceptive apps that claim to offer money, rewards, casino winnings, and premium content. Early Access apps are apps that haven't been released on the official Android app marketplace.

[OVERFLOW] SOURCE: [mp.weixin.qq.com](https://mp.weixin.qq.com) | VEL: 88/100

### Analysis of the attack chain of APT-C-55 (Kimsuky) organization using disguised installation packages to implant remote control Trojans

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

, 360 captured the attack activities of the Kimsuky organization. It is delivered through malicious LNK files and PowerShell. After environmental detection and memory loading, it uses scheduled tasks disguised as Chrome updates to achieve persistence to achieve long-term control purposes. APT-C-55 Kimsuky (also known as BabyShark etc.) is an advanced persistent threat organization that has long targeted South Korea's think tanks, government diplomatic departments, news media, and educational and academic institutions. The main purpose of cyber attacks is to steal intelligence. This organization is very active. Even though it has been repeatedly disclosed by security vendors in recent years, it has not been able to stop it from continuing its latent operations.

[OVERFLOW] SOURCE: [github.com](https://github.com) | VEL: 88/100

### Promptfoo LLM Security & Eval 0.123.0: 0.123.0 ■ BREAKING CHANGES providers: default GPT-5.6+ to Responses Features assertions:

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

0.123.0 (2026-09-10) ■ BREAKING CHANGES providers: default GPT-5.6+ to Responses (#10803) Features assertions: grade native audio with llm-rubric (#10814) (39c4fb5) providers: add Claude Fable and Mythos 5.1 (#10601) (ac6ba5b) providers: add GPT-6 Astra support (#10625) (cd4c68b) providers: add grok-4.6, grok-image-image-2.0, and grok-image-video-1.5 (#10587) (48a71cd) providers: default GPT-5.6+ to Responses (#10803) (b051efd) providers: expose MCP tool calls in response metadata (#10516) (.

[OVERFLOW] SOURCE: [mp.weixin.qq.com](https://mp.weixin.qq.com) | VEL: 88/100

### Seamless replacement of BlackDuck ■ Suspended mirror security code security + intelligent body security products both rank first in the Chinese market in terms of application rate!

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

"AI + DevOps Status Survey Report" ranks first in market application rate for five consecutive years, continuing to lead the new generation of digital supply chain security. Recently, the China Communications Standards Association released the "AI+DevOps Current Situation Survey Report (2026)". The survey cycle covers July 2026-September 2026, spanning high-demand industries such as finance, energy, government affairs, central and state-owned enterprises, and high-end manufacturing. Focusing on the development pattern and ecological status quo, a total of 60 companies collected 3,351 valid questionnaires. After sample screening, expert review, and data cross-checking, it is the first domestic AI+DevOps The authoritative reference for model selection.

[OVERFLOW] SOURCE: [vttresearch.com](https://vttresearch.com) | VEL: 88/100

### SemiQon's cryogenic chip technology for quantum computing and space applications receives award from EARTO

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

EARTO, the organisation of the European Research and Technology Organisations, awarded SemiQon and VTT first prize in the "Impact Expected" category on 14 October 2025 in Brussels for a pioneering cryogenic CMOS (complementary metal-oxide semiconductor) chip innovation. The solution enables the full capacity of advanced CMOS functionalities at cryogenic temperatures, thereby unlocking new possibilities for quantum computing and space applications.

[OVERFLOW] SOURCE: [asia.nikkei.com](https://asia.nikkei.com) | VEL: 88/100

### OpenAI says it is in talks to train models in Australia

ARCH: MOE • CTX: 128K • QUANT: FP8 • STACK: vLLM

OpenAI says it is in talks to train models in Australia marks an architectural milestone in Frontier API foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[OVERFLOW] SOURCE: [asia.nikkei.com](https://asia.nikkei.com) | VEL: 88/100

### SoftBank Group plunges 11% after OpenAI says no IPO this year

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

SoftBank Group plunges 11% after OpenAI says no IPO this year reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology development. Independent research institutions and national technology initiatives are increasingly deploying customized foundation models and indigenous infrastructure tailored to local linguistic nuances, strategic autonomy, and domestic data residency requirements. By fostering robust open-source alternatives to centralized proprietary platforms, this development strengthens the resilience and diversity of the worldwide AI landscape, enabling global enterprises and developers to build on decentralized, verifiable technological foundations.

[OVERFLOW] SOURCE: [kdnuggets.com](https://kdnuggets.com) | VEL: 88/100

### Why DeepSeek-V4.1-Flash Is Such an Exciting Open Model Release

ARCH: MOE • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

DeepSeek-V4.1-Flash shows how Causal Encoder-Decoder architecture, MoE, KV cache compression, CSA2, cheaper prefill, and efficient decoding can make powerful open-source AI models far more efficient to run. DeepSeek has released DeepSeek-V4.1-Flash, and while the benchmark numbers are impressive, they are probably not the most interesting part of this release. The architecture is. DeepSeek is tackling several problems that are becoming increasingly important as AI moves toward long-running agents: expensive prefill, huge KV caches, long contexts, memory bandwidth, and the cost of maintaining agent state across interactions. Rather than simply making the model larger, DeepSeek has redesigned several parts of the architecture and inference stack to make long-context AI much cheaper to run.

**COLOPHON & SENSOR METHODOLOGY:** The Global AI Gazette is compiled autonomously. Data is aggregated across authoritative global AI sources including Hugging Face, GitHub Trending, arXiv, and leading AI research lab feeds. Edition #2238 • 160 items processed • Tuesday, September 15, 2026 • 22:09 UTC. All rights reserved.