

THE GLOBAL AI GAZETTE

Worldwide AI Ecosystem Intelligence • Open Source • Research • Models • Tools • Sovereign AI

Thursday, September 10, 2026 • 16:05 UTC

EDITION NO. 2212

201 STORIES ANALYZED

■ LEAD STORY

ANOTHER researcher accuses OpenAI of training on conversations and then claiming a breakthrough

Velocity 96/100 • Impact 88/100 • Source: reddit.com

ANOTHER researcher accuses OpenAI of training on conversations and then claiming a breakthrough introduces key developments in machine learning foundation architectures. The release advances reasoning, inference efficiency, and model deployment across open-weight and frontier environments.

INNOVATION: Model: Frontier Foundation Weights & Open Checkpoint
ACTION: Deployment: Checkpoints accessible on Hugging Face; compatible with standard runtimes.

■ SECONDARY ANCHOR

Deepseek V4.1 Flash Release Video [Made with Deepseek V4.1 Flash]

I like to benchmark new models that come out on motion videos. So here's a test I did for deepseek v4.1 flash. And I have to say flash has probably graduated from being a Luna class model to nearly an Opus class model with this release, at least with motion videos. Prev.

■ TOP FLASH BULLETINS

[VEL 96] guide to using reasoning_effort on deepseek v4.1 flash — guide to using reasoning_effort on deepseek v4.1 flash introduces key developments in machine learning foundation architectures. The release advances reasoning, inference efficiency, and model deployment...

[VEL 96] Qwen3.8-27B-Uncensored-Genesis-V1-GGUF — Model available Qwen3.8-27B-Uncensored-Genesis-V1-MTP-GGUF This model is a practical realisation of things described in this paper, but adapted by me for machine learning: I am trying...

[VEL 88] HealthBench-Psych: A Mental Health Subset of OpenAI's HealthBench — General-purpose health benchmarks increasingly anchor claims about LLM medical performance, but they are not always resolved by clinical specialty, making domain-specific performance hard to isolate....

[VEL 88] Enhancing Virtual Agents through SLMs and Edge-Computing: An Exploratory Evaluation of Think and Memory Processes — Embodied intelligent virtual agents are expected to operate as persistent, adaptive, and context-aware entities within complex virtual and Metaverse worlds. However, implementing cognitively capable agents...

[VEL 88] Social Chain of Thought: A Multi-Agent Architecture Grounded in Medical Differential Diagnosis Methodology — Medical diagnostic reasoning is a high-impact use case for LLMs that carries significant implications for the health and wellbeing of users. When OpenAI (2026) reports...

■ SECTION II: EXECUTIVE AI STRATEGIC BRIEFING

Strategic Horizons & Action Directives for Enterprise Leadership

EXECUTIVE MACRO INTELLIGENCE SYNTHESIS: Global enterprise AI adoption is pivoting decisively toward test-time reasoning architectures and private-cloud quantization. Technology leadership must actively balance proprietary frontier model APIs with sovereign, open-weight deployments to reduce token expenditure while strictly sandboxing autonomous agent tool-calling boundaries.

VECTOR	DEVELOPMENT	SOURCE	DIRECTIVE
Model Sourcing & Licensing	Running Vision Qwen 3.8 27B on a 16GB Ca	reddit.com	Audit open-weights licensing vs proprietary APIs; evaluate reddit.com
Compute & Infrastructure CapEx	DeepSeek-V4.1-Flash surprised	reddit.com	Review GPU cluster allocation and power envelopes; benchmark reddit.co
Agentic Autonomy & Governance	Deepseek V4.1 Flash is 748B, not 552B	reddit.com	Implement deterministic sandboxes for autonomous tool execution, stric
Inference Latency & Quantization	Deepseek v4.1 flash finally has engrams,	reddit.com	Benchmark KV-cache compression (FP8/INT4/GGUF) and modern inference en
Open-Source Supply Chain	DeepSeek V4-1 Flash is out	reddit.com	Inspect upstream repository dependencies; audit tokenizer code, weight
Data Residency & Sovereignty	DeepSeek V4.1 Flash: Stronger, Faster, M	reddit.com	Verify compliance with sovereign AI frameworks and regional data resid

1. **Running Vision Qwen 3.8 27B on a 16GB Card, the config (45tks):** Audit open-weights licensing vs proprietary APIs; evaluate reddit.com parameter efficiency and commercial distribution terms.
2. **DeepSeek-V4.1-Flash surprised:** Review GPU cluster allocation and power envelopes; benchmark reddit.com hardware efficiency to optimize cost per token.
3. **Deepseek V4.1 Flash is 748B, not 552B:** Implement deterministic sandboxes for autonomous tool execution, strict rate limiting, and human-in-the-loop validation.
4. **Deepseek v4.1 flash finally has engrams, what do you expect from 4.1 pro?:** Benchmark KV-cache compression (FP8/INT4/GGUF) and modern inference engines (reddit.com) against TTFT SLAs.
5. **DeepSeek V4-1 Flash is out:** Inspect upstream repository dependencies; audit tokenizer code, weights provenance, and pinned runtime releases for reddit.com.
6. **DeepSeek V4.1 Flash: Stronger, Faster, More Accessible:** Verify compliance with sovereign AI frameworks and regional data residency requirements for reddit.com deployments.

■ SECTION III: TRENDING OPEN-SOURCE AI & GITHUB INNOVATIONS

Top GitHub Repositories, Developer Velocity & Architecture Dissections

PROJECT / REPO	SOURCE	STACK	VELOCITY	CORE ARCHITECTURAL FOCUS
Deploying Qwen3.8-2.4T-A95B on Amazon SageMak	aws.amazon.com	Python	88/100	On August 12, 2026, Alibaba's Qwen team released Qwen3.8-2.4T-A95B . This is the first time a Qwen-Max-class model.
Anthropic researcher quits with a warning: Se	arstechnica.com	Python	88/100	When a prominent researcher quits a job at a frontier AI lab these days, it's often to pursue.
Anthropic researcher believes more than 10% c	bbc.co.uk	Python	88/100	It is the latest in a series of increasing warnings about the safety threat posed by artificial intelligence.
OpenAI's deeper Samsung tie-up lands on a fou	digitimes.com	Python	88/100	OpenAI's deeper Samsung tie-up lands on a foundry that is running out of capacity delivers key capabilities for.
OpenAI says it cracked 90-year-old maths prob	bbc.co.uk	Python	88/100	OpenAI's claim that it solved parts of Navier-Stokes equations has quickly stirred controversy.
OpenAI Admits More AI Agents Went Astray in M	aibusiness.com	Python	88/100	The incident was the latest in a string of unauthorized actions by AI agents.

[TRENDING REPO] SOURCE: [aws.amazon.com](#) | VEL: 88/100
Deploying Qwen3.8-2.4T-A95B on Amazon SageMaker HyperPod with vLLM

On August 12, 2026, Alibaba's Qwen team released Qwen3.8-2.4T-A95B . This is the first time a Qwen-Max-class model has been made available as open weights. With 2.4 trillion total parameters (95 billion activated per token), a hybrid linear-plus-full-attention architecture, and native context up to 262K tokens (extensible to 1M), Qwen3.8 targets the most demanding agentic and reasoning workloads. These include multi-step coding, long-horizon planning, and autonomous tool use. Open weights models give you full control. Data stays within your infrastructure, inference behavior can be customized, and there are no per-token API fees at scale.

[TRENDING REPO] SOURCE: [arstechnica.com](#) | VEL: 88/100
Anthropic researcher quits with a warning: Self-improving AI could "kill us all"

When a prominent researcher quits a job at a frontier AI lab these days, it's often to pursue a new startup or protest a new business model. But AI researcher Jacob Coxon is using his departure from Anthropic to publicly warn that frontier AI companies are "gambling with our lives" with systems that they "earnestly believe... could kill us all by the end of the decade." In a social media thread Tuesday night, Coxon said that this existential risk is inherent not so much in today's models but.

[TRENDING REPO] SOURCE: [bbc.co.uk](#) | VEL: 88/100
Anthropic researcher believes more than 10% chance AI 'could kill all humans'

It is the latest in a series of increasing warnings about the safety threat posed by artificial intelligence.

[TRENDING REPO] SOURCE: [digitimes.com](#) | VEL: 88/100
OpenAI's deeper Samsung tie-up lands on a foundry that is running out of capacity

OpenAI's deeper Samsung tie-up lands on a foundry that is running out of capacity delivers key capabilities for AI software engineering and local execution. Engineered to enhance developer velocity, it streamlines model serving, evaluation, and pipeline orchestration.

[TRENDING REPO] SOURCE: [bbc.co.uk](#) | VEL: 88/100
OpenAI says it cracked 90-year-old maths problem in 88 hours

OpenAI's claim that it solved parts of Navier-Stokes equations has quickly stirred controversy.

■ SECTION IV: FRONTIER FOUNDATION MODELS & REASONING
BREAKTHROUGHS

Reasoning Architectures, Open Weights & Capability Benchmarks

FOUNDATION MODEL	SOURCE	SIZE / CTX	IMPACT	ARCHITECTURAL HIGHLIGHT
guide to using reasoning_effort on deepseek v	reddit.com	MoE / SOTA / 128K	88/100	guide to using reasoning_effort on deepseek v4.1 flash introduces key developments in machine learning foundation architectures.
Qwen3.8-27B-Uncensored-Genesis-V1-GGUF	reddit.com	27B / 128K	88/100	Model available Qwen3.8-27B-Uncensored-Genesis-V1-MTP-GGUF This model is a practical realisation of things described in this paper, but adapted by.
DeepSeek-V4.1-Flash surprised	reddit.com	50B / 128K	68/100	Hoping to see smartest medium size models soon & later with all available optimizations/architectures/etc.,. Thanks Deepseek!
Deepseek V4.1 Flash is 748B, not 552B	reddit.com	748B / 128K	68/100	People keep on getting confused about this, so I looked at the safetensors on hf.
Deepseek v4.1 flash finally has engrams, what	reddit.com	MoE / SOTA / 128K	68/100	If the ratio is the same, Maybe 1.6T -3.1T params plus .56T-1.06T engrams and fable 5.0 level performance?.
DeepSeek V4-1 Flash is out	reddit.com	MOE / 128K	68/100	Here we go again, DeepSeek is back again with a new model V4-1 Flash A multimodal Mixture-of-Experts (MoE).

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 88/100
guide to using reasoning_effort on deepseek v4.1 flash
guide to using reasoning_effort on deepseek v4.1 flash introduces key developments in machine learning foundation architectures. The release advances reasoning, inference efficiency, and model deployment across open-weight and frontier environments.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 88/100
Qwen3.8-27B-Uncensored-Genesis-V1-GGUF
Model available Qwen3.8-27B-Uncensored-Genesis-V1-MTP-GGUF This model is a practical realisation of things described in this paper, but adapted by me for machine learning: I am trying to solve the problem: why LLM models even for simple questions write walls of text during reasoning, and burn too much tokens instead of solving the task. And when number of parameters increase the problem became worse.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 68/100
DeepSeek-V4.1-Flash surprised
Hoping to see smartest medium size models soon & later with all available optimizations/architectures/etc.,. Thanks Deepseek! Ex 1: 30-50B MOE + 10-15B Engram + DeepSeek-V4.1-Flash type KVCache Ex 2: 15-30B Dense + 10-15B Engram + DeepSeek-V4.1-Flash type KVCache EDIT: Updated Engram to 10-15B from.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 68/100
Deepseek V4.1 Flash is 748B, not 552B
People keep on getting confused about this, so I looked at the safetensors on hf. The title should have been "Deepseek V4.1 Flash is 748B total/552B base, not 284B or 305B or 485B or 522B" The model is not 284B. The original Deepseek V4 Flash is 284B, but not the V4.1 Flash model The model is not 305B, despite what some people claim "So: ~305B real backbone + 203B engram = 508B total" This is incorrect.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 68/100
Deepseek v4.1 flash finally has engrams, what do you expect from 4.1 pro?
If the ratio is the same, Maybe 1.6T -3.1T params plus .56T-1.06T engrams and fable 5.0 level performance? Maybe v4.2 or 4.5 will have engram gradient modification?

■ SECTION V: TOP AI RESEARCH PAPERS & ARXIV BREAKTHROUGHS

Reasoning Paradigms, Multimodal Architectures & Algorithmic Design

[AI RESEARCH] [SOURCE: arxiv.org](#) | VEL: 88/100

HealthBench-Psych: A Mental Health Subset of OpenAI's HealthBench

General-purpose health benchmarks increasingly anchor claims about LLM medical performance, but they are not always resolved by clinical specialty, making domain-specific performance hard to isolate. Mental health is of acute public-health concern as millions of people turn to LLMs for psychological support, and most existing evaluations are bespoke academic benchmarks that are difficult to integrate into developer workflows. We introduce HealthBench-Psych and HealthBench-Psych-Hard.

[AI RESEARCH] [SOURCE: arxiv.org](#) | VEL: 88/100

RedKnot-MLA: Multi-Head Offline-Online Reuse for DeepSeek-V4 Long-Context Serving

Multi-head latent attention (MLA) exposes many logical query heads through one packed latent KV stream. This representation is memory efficient, but it removes the physical per-head cache boundary assumed by conventional head-wise reuse. We present our system, a DeepSeek-V4 realization of RedKnot's head-aware reuse principle. Each immutable document is processed offline at canonical position zero; certified Local-head contributions are retained as MLA-Off.

[AI RESEARCH] [SOURCE: arxiv.org](#) | VEL: 88/100

A Human Audit of OpenAIs AI-Generated Mathematical Proofs

We assess 18 chapter-specific reviews of the ten mathematical results announced by OpenAI on 1 August 2026, alongside review standards, Lean formalizations, subsequent research, and mathematical references. The article audits this review record without claiming a complete reconstruction of all ten proofs. No confirmed substantive mathematical error in a principal result remains in the examined assessments, although review depth varies and some dependencies remain partly checked.

[AI RESEARCH] [SOURCE: arxiv.org](#) | VEL: 72/100

Building Multilingual Bridges: Data Mixing as the Pillar of Generalization for In-Language Reasoning

Reasoning language models have made substantial advances on a variety of complex tasks, yet their capabilities remain overwhelmingly English-centric: models primarily reason in English regardless of the language they are prompted in. This is inaccessible for non-English-speaking users, risks losing the intent of the original question, and forgoes knowledge more readily expressed in the target language.

[AI RESEARCH] [SOURCE: news.mit.edu](#) | VEL: 72/100

MIT Schwarzman College of Computing launches pilot to help educators teach AI across disciplines

This summer, the MIT Schwarzman College of Computing welcomed faculty from colleges and universities across Greater Boston, South Carolina, West Virginia, and Texas to campus for the inaugural AI Educators Pilot, a weeklong workshop aimed at expanding how artificial intelligence is taught across disciplines and learning environments.

[AI RESEARCH] [SOURCE: arxiv.org](#) | VEL: 88/100

Qishi Engine on AstaBench E2E-Bench-Hard

This report analyzes Qishi Engine v0.8 across all 40 test tasks in AstaBench E2E-Bench-Hard, a benchmark that requires autonomous agents to carry a research question through experimental design, code implementation, actual execution, result analysis, and report delivery. Qishi Engine is model-configurable; this evaluation selected DeepSeek deepseek-v4pro-preview as the model backend. The official AstaBench leaderboard records a score of 0.816 and an average benchmark cost of USD 15.209 per.

■■ SECTION VI: DEVELOPER TOOLS, FRAMEWORKS & AI INFRASTRUCTURE

Inference Runtimes, Evaluation Frameworks, Vector DBs & Local Execution Engines

[AI TOOL] SOURCE: reddit.com | VEL: 80/100

I trained an audio model that can generate infinite one-shots for music production and turn text prompts into fully playable synths. I'm not only releasing the model but I've also released a video on exactly how I did it (and the inferencing pipeline to let others make text based synths.)

(hopefully this is okay to here - it seems like audio models and image / video modeals is allowed but yeah this is a bit different) So I've been doing independent audio research for a while now. The ultimate dream of this work was actually getting an AI to respond not only to instruments but also timbre itself as separate controllable things. Think a Grand Piano can sound both Warm / Gritty but also Cold / Sparkly. Its still a piano though.

[AI TOOL] SOURCE: mp.weixin.qq.com | VEL: 30/100

■■ Agent Loop ■■■■ RSAC2026 ■■■■ AIDR ■■■■■■■■■■

Multi-modal AIDR 2026-04-29 16:31 Beijing Achieves 95% asset visibility and second-level blocking, Lingjing AIDR protects the entire Agentic AI application landscape! "Security Paradigm Shift" from Moscone Center #March 23, 2026, Moscone Center, San Francisco. When Geordie AI won the RSAC2026 Innovation Sandbox Championship trophy, the applause from the audience could not conceal the anxiety of security people around the world: security has officially entered the no-man's land ruled by "Agentic AI". The most frightening thing is no longer the code, but the AI ■■■digital employees who make decisions independently on the corporate intranet but are in a "supervisory vacuum." 1■.

[AI TOOL] SOURCE: reddit.com | VEL: 96/100

Closed AI doesn't like biological research, user turns to open weight models

OpenAI has decided to fully shut down a protein design project I'm working on for a client. Needless to say, open weight models are the only way forward.

[AI TOOL] SOURCE: reddit.com | VEL: 96/100

Why the hell is LM Studio making LM Studio so difficult to download?

Who is the marketing genius at LM Studio that decided that going ALL IN on pushing their new Bionic Agent product meant they are going to make it a giant pain in the ass to find and download actual LM Studio. This is the dumbest marketing decision I've ever seen. I used to love LM Studio, it was the middle stepping stone in the logical progression of inference. Most OGs here likely started with Ollama, moved to LM Studio, on their way to vLLM. Now trying to go to LM Studio takes you to Bionic.

[AI TOOL] SOURCE: simonwillison.net | VEL: 96/100

On the Navier–Stokes Millennium Prize Problem

On the Navier–Stokes Millennium Prize Problem Impressive result from OpenAI, who used an unreleased model to produce a resolution to the Navier–Stokes existence and smoothness problem, one of the seven Millennium Prize Problems that have been subject to a \$1,000,000 prize since May 24th, 2000.

[AI TOOL] SOURCE: simonwillison.net | VEL: 96/100

Introducing ChatGPT Images 2.5

Introducing ChatGPT Images 2.5 OpenAI's image generation models are apparently used "more than 3 billion images across ChatGPT Images and the GPT■Image models in the API". This latest release improves their instruction-following ability across multiple turns, responds faster, and "is better at preserving the subjects in your reference photos". There are two new model IDs in the API: gpt-image-2.5-sunburst and gpt-image-2.5-flare.

■ SECTION VII: SOVEREIGN AI & WORLDWIDE REGIONAL ECOSYSTEMS

National AI Labs, Regional Models & Sovereign Compute Ecosystems

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](https://www.reddit.com) | VEL: 96/100

3060 12GB vs 4060 ti 16GB

I'm currently building my system around 3060s, but I might be able to get a 4060 for a nice deal. At first it seemed like a no brainer, but turns out the 4060 has lower memory bandwidth. In a system that already has 4x 3060 12GBs set up on a threadripper with tensor parallelism (mostly qwen3.8-27b), would it be worth having the 4060 ti 16GB around for the extra 4GB and occasional gaming, or is it just going to slow the rest of the setup down for AI?

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](https://www.reddit.com) | VEL: 96/100

Harness doesn't matter

I was not aware that the harness makes such a big difference. DeepSeek V4.1.

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](https://www.reddit.com) | VEL: 96/100

Mention if a "new model" is a finetune

A few posts tagged with "new model" present models that are finetunes. My opinion : I'd rather have the "new model" tag reserved for new "major" releases, like a new Qwen model, Deepseek V4 -> Deepseek V4.1, etc., that involved a new pretrain or intensive post-training (in opposition to a small finetune). Otherwise, maybe prepend "[Finetune]" to the title to indicate that the new model is "less of a big news", a use a "new finetune" tag, to differentiate between the two kinds of new models.

[SOVEREIGN [CN]] SOURCE: arxiv.org | VEL: 88/100

How Does mHC Use Its Residual Streams? Selective Routing and Near-Identity Mixing

Hyper-Connections and their manifold-constrained variant mHC widen a residual pathway from one stream to n, yet how trained models use this capacity remains unclear: how broadly blocks read and write, how strongly the residual pathway mixes streams, and whether the streams carry distinct representations. We examine these properties in the four-stream residual pathway of DeepSeek-V4-Flash using effective stream counts, cross-stream residual weights, and inter-stream cosine similarity.

[SOVEREIGN [CN]] SOURCE: arxiv.org | VEL: 88/100

'Ghaib in Translation' aka Unseen Harm: Measuring Cross-Script Safety Inconsistency with 'Missed-in-Urdu' Scores in LLM Hate Speech Detection

Urdu, the world's tenth most spoken language with 246 million speakers, remains almost entirely absent from mainstream LLM safety evaluation and nine years of WOAHP proceedings. To investigate whether this absence has measurable consequences for content moderation reliability, five large language models, GPT-4o, Claude Sonnet 4.5, Gemini 2.5 Flash, Qwen-2.5, and Llama-3.1, were tested across six datasets spanning Nastaliq Urdu, Roman Urdu, English, and code-switched Urdu-English.

[SOVEREIGN [FR]] SOURCE: arxiv.org | VEL: 88/100

Explanatory Engagement Under Rare Anomalous Failure: Asymptotic Rarity in Model Behavior (or: The Asymptotic AI)

Prior work on LLM behavior under anomalous conditions asks whether a model notices anomalies. We ask a narrower question: once a model sits in a workflow with a low, controllable failure rate, does its explanatory engagement - length, specificity, self-reported confidence - change as failure grows asymptotically rarer? We built a local, zero-cost harness on three open-weight models (qwen3:8b, llama3.1:8b, mistral:7b) running a repeated tool-call task where one call fails at probability p.

SECTION VIII: AI HARDWARE, COMPUTE CLUSTERS & SILICON

GPU Clusters, Custom Accelerators, Datacenter Infrastructure & Silicon

SILICON / SYSTEM	SOURCE	COMPUTE SPECS	VELOCITY	BENCHMARK / TELEMETRY
Running Vision Qwen 3.8 27B on a 16GB Card, t	reddit.com	High Velocity	96/100	I am just sharing my config for Qwen 3.8 27b that fits on a 5060TI, what is cool.
Qwen3.8-Flash-Next on 2x3090 + DDR4, part 4:	reddit.com	29 t/s	96/100	. 17 -> 25-29 t/s with the expert cache PR, 37-41 t/s after switching to UD-Q4_K_XL and stacking.
Qwen3.8-Flash-Next on MLX-serve, 1m context i	reddit.com	40 tok/s	96/100	Hi, I'm the co-creator of this Qwen3.8-Flash-Next engine support in MLX-serve. I've been tuning this one to run.
Substrate-Aware AI Agents: Execution Context	arxiv.org	High Velocity	88/100	Autonomous AI agents increasingly select actions in environments whose memory, execution-time, runtime, compute, and operational constraints determine what.
SAT-Edge-Agent: Hardware-in-the-Loop Edge-Age	arxiv.org	High Velocity	88/100	Onboard satellite intelligence requires a task layer that translates mission intent into local tool calls, exposes execution state.
Best Open source TTS right now for narration?	reddit.com	High Velocity	80/100	I run these models on Kaggle notebook, so not all TTS models, such as the ones that use.

[AI HARDWARE] SOURCE: reddit.com | VEL: 96/100
Running Vision Qwen 3.8 27B on a 16GB Card, the config (45tks).

I am just sharing my config for Qwen 3.8 27b that fits on a 5060TI, what is cool about this is that you can even get vision! and a 85K context (I have 1.5gb of headroom for more context or a better quant) Model: IQ3_XXS-mtp from Using beellama Config used: [*] model = ..\llm-models\Qwen3.8-27B-GSQ-RCO-IQ3_XXS-mtp.gguf mmproj = ..\llm-models\mmproj-Qwen3.8-27B-BF16.gguf image-min-tokens = 256.

[AI HARDWARE] SOURCE: reddit.com | VEL: 96/100
Qwen3.8-Flash-Next on 2x3090 + DDR4, part 4: 2.2-2.5x faster prefill by kicking the expert cache off the GPU while the prompt runs . 17 -> 25-29 t/s with the expert cache PR, 37-41 t/s after switching to UD-Q4_K_XL and stacking MTP on the cache, the top-k fallback that was sorting more than it needed to. This one is all about prefill, which was honestly the weak spot the whole time. 80+ seconds before the first token on an 8k prompt, and 24 minutes on a 119k one...I know . Box is still 2x 3090, dual Broadwell Xeon, llama.cpp, UD-Q4_K_XL with the Q8 MTP head on the.

[AI HARDWARE] SOURCE: reddit.com | VEL: 96/100
Qwen3.8-Flash-Next on MLX-serve, 1m context is released!
Hi, I'm the co-creator of this Qwen3.8-Flash-Next engine support in MLX-serve. I've been tuning this one to run both fast, efficient and correct up 1m context using kv cache 8 bits in M5 Max 128GB. Qwen is working well at very long context as showed in the video (a snapshot at ~760k context), I let it build MLX Serve Monitor plugin that you've seen on the right side of Opencode2's app. The generation sustain through 1m context at around 40 tok/s on prose and 75 tok/s on coding.

[AI HARDWARE] SOURCE: arxiv.org | VEL: 88/100
Substrate-Aware AI Agents: Execution Context as a First-Class Input

Autonomous AI agents increasingly select actions in environments whose memory, execution-time, runtime, compute, and operational constraints determine what counts as a suitable plan. We call the absence of this execution context from an agent's planning state substrate blindness. We test this general proposition through numerical code generation, where selected implementation choices and operational consequences are directly observable.

[AI HARDWARE] SOURCE: arxiv.org | VEL: 88/100
SAT-Edge-Agent: Hardware-in-the-Loop Edge-Agent Orchestration for Onboard Satellite Intelligence

Onboard satellite intelligence requires a task layer that translates mission intent into local tool calls, exposes execution state, and returns machine-consumable artifacts under communication and power constraints. We present SAT-Edge-Agent, a hardware-in-the-loop (HIL) edge-agent system deployed on a commercial off-the-shelf ARM-based heterogeneous edge system-on-chip.

■ SECTION IX: AUTONOMOUS AGENTS, MULTI-AGENT SWARMS & ROBOTICS

Agent Orchestration, Computer-Use Primitives, Action Models & Robotics

AGENT / FRAMEWORK	SOURCE	PROTOCOL	VELOCITY	CORE CAPABILITY DOMAIN
Enhancing Virtual Agents through SLMs and Edg	arxiv.org	PyTorch / ONNX	88/100	Embodied intelligent virtual agents are expected to operate as persistent, adaptive, and context-aware entities within complex virtual and.
Social Chain of Thought: A Multi-Agent Archit	arxiv.org	PyTorch / ONNX	88/100	Medical diagnostic reasoning is a high-impact use case for LLMs that carries significant implications for the health and.
When LLM Agents Negotiate: Private Informatio	arxiv.org	PyTorch / ONNX	88/100	As LLM agents move from decision support to autonomous procurement, firms need to know whether delegated negotiators create.
Research acceleration: The view inside OpenAI	simonwillison.ne	PyTorch / ONNX	96/100	Research acceleration: The view inside OpenAI Apparently today is RSI day at OpenAI, for Recursive Self-Improvement - I.
Scanning the Harness: An Empirical Study of S	arxiv.org	PyTorch / ONNX	88/100	AI coding agents such as Claude Code, Cursor, GitHub Copilot, and OpenAI Codex are configured through artifacts developers.
What Does Multi-Harness RL Learn? Credit Assi	arxiv.org	PyTorch / ONNX	88/100	Agent reinforcement learning (RL) increasingly runs through full execution harnesses, and a multi-harness recipe mixes two choices: exposing.

[AGENT SYSTEM] SOURCE: [arxiv.org](#) | VEL: 88/100

Enhancing Virtual Agents through SLMs and Edge-Computing: An Exploratory Evaluation of Think and Memory Processes

Embodied intelligent virtual agents are expected to operate as persistent, adaptive, and context-aware entities within complex virtual and Metaverse worlds. However, implementing cognitively capable agents in such environments is conceptually and technologically challenging.

[AGENT SYSTEM] SOURCE: [arxiv.org](#) | VEL: 88/100

Social Chain of Thought: A Multi-Agent Architecture Grounded in Medical Differential Diagnosis Methodology

Medical diagnostic reasoning is a high-impact use case for LLMs that carries significant implications for the health and wellbeing of users. When OpenAI (2026) reports that more than 5% of ChatGPT messages globally are healthcare-related, the transparency of these systems becomes a serious design concern. This is especially true for complex cases, where differential diagnosis often requires integrating multiple forms of specialist reasoning.

[AGENT SYSTEM] SOURCE: [arxiv.org](#) | VEL: 88/100

When LLM Agents Negotiate: Private Information and Dynamic Bargaining in Supply Chains

As LLM agents move from decision support to autonomous procurement, firms need to know whether delegated negotiators create value, divide it predictably, and avoid money-losing contracts. We study this in a canonical supply chain bargaining problem: a buyer with private demand information negotiates a quantity-payment contract with an uninformed seller. We benchmark nine LLMs from OpenAI, Google, and Alibaba against a validated Perfect Bayesian Equilibrium across 9,840 LLM-to-LLM negotiations.

[AGENT SYSTEM] SOURCE: [simonwillison.net](#) | VEL: 96/100

Research acceleration: The view inside OpenAI

Research acceleration: The view inside OpenAI Apparently today is RSI day at OpenAI, for Recursive Self-Improvement - I think it's their new AGI. Both this piece and the new essay An Alien Mind (by Chief Scientist Jakub Pachocki) talk about it, and this one doesn't even bother to expand the acronym. Included are details on how OpenAI's own research team are using coding agents.

[AGENT SYSTEM] SOURCE: [arxiv.org](#) | VEL: 88/100

Scanning the Harness: An Empirical Study of Supply-Chain Defects in AI Coding-Agent Configurations

AI coding agents such as Claude Code, Cursor, GitHub Copilot, and OpenAI Codex are configured through artifacts developers write and share: instruction files, skills, hooks, MCP server declarations, subagents. This harness is a dependency layer installed from marketplaces and public repositories, running with the developer's privileges, with no lockfile, no install-time check, and no vocabulary for what a component may do.

■ SECTION X: GLOBAL AI COMMUNITY WIRE & OVERFLOW DIGEST

High-velocity AI stories from today's sweep that didn't fit earlier sections

[OVERFLOW] [SOURCE: arxiv.org](#) | VEL: 88/100

The Truth Stays in the Family: Enhancing Contextual Grounding via Inherited Truthful Heads in Model Lineages

Recent advances in large language models (LLMs) have produced many specialized multimodal LLMs (MLLMs) that share common foundational LLMs, forming distinct model lineages. It remains unclear whether a fundamental behavioral link exists between the foundational LLMs and downstream variants. We investigate this question by quantifying head-level context-truthfulness scores. Across diverse LLM and MLLM lineages, including Vicuna-, Qwen2.5-, LLaMA2-, and Mistral-based models, we find that Truth.

[OVERFLOW] [SOURCE: aws.amazon.com](#) | VEL: 88/100

Take on your most ambitious work with GPT-6 Astra on Amazon Bedrock

GPT-6 Astra from OpenAI brings greater depth and judgment to your most demanding tasks and runs on the Amazon Bedrock inference engine built for high performance, security, and scale. Organizations are already running AI agents that write code, analyze data, and automate complex workflows at production scale on Amazon Bedrock. GPT-6 Astra raises the potential of what those agents can deliver.

[OVERFLOW] [SOURCE: aibusiness.com](#) | VEL: 88/100

How Mistral's New Funding is a Bridge to Sovereign AI

After starting as an open-weight startup, Mistral has shifted toward sovereign AI, given the European market it operates in.

[OVERFLOW] [SOURCE: aibusiness.com](#) | VEL: 88/100

AI Beyond the Hype: Insights From Google

By establishing key foundations, businesses can create effective AI workflows that drive meaningful outcomes.

[OVERFLOW] [SOURCE: bitsoffreedom.nl](#) | VEL: 88/100

Don't let chatbot developers avoid their responsibility any longer

Large tech companies behind generative AI software do not take sufficient responsibility for the output of their tools, even if it helps perpetrators of abuse or violence. That must change, either through stricter rules or through lawsuits. OpenAI, the company behind AI chatbot ChatGPT, has already faced several lawsuits.

[OVERFLOW] [SOURCE: aisingapore.org](#) | VEL: 88/100

Breaking Barriers, Building Bridges: Increasing Language Representation in Southeast Asia

Held outside of Singapore for the very first time, the third Languages Summit was co-hosted by AI Singapore, Google and VISTEC in Bangkok, Thailand and brought together a passionate community of AI experts and researchers from all corners of the SEA region. This event, dedicated to building a more inclusive AI future, provided a platform for exchanging discussions and insights around efficient model training, obtaining and sharing high-quality data, and regional updates on AI.

[OVERFLOW] [SOURCE: aisingapore.org](#) | VEL: 88/100

Learning to Generate Unbounded 3D Scenes from Image Collections

Introduction Scene generation has raised considerable attention in recent years, addressing the growing need for 3D creative tools in the metaverse. At the core of 3D content creation is inverse graphics, which aims to recover 3D representations from 2D observations. Given the cost and labor for creating 3D assets, the ultimate goal of 3D content creation would be learning a generative model from in-the-wild 2D images.

[OVERFLOW] [SOURCE: reddit.com](#) | VEL: 80/100

Apple wants to give me \$1175 for a Mac Mini M4 Pro? And would you sell for a DGX Spark or M5-based Studio (which?)

I thought Trade-in value offered by Apple was only ever close to reasonable (for not having to go through the extra work of selling it yourself) if you bought the base model and did not upgrade anything. And you would get less than half of what you paid. For example: The base price of the M4 Pro Mac Mini was \$1399. On Apple's trade in page for the Mac Mini it says "Up to \$620". So their offer retains 44% of the value. But I upgraded the GPU, RAM, and SSD pushing the price to \$2099.

COLOPHON & SENSOR METHODOLOGY: The Global AI Gazette is compiled autonomously. Data is aggregated across authoritative global AI sources including Hugging Face, GitHub Trending, arXiv, and leading AI research lab feeds. Edition #2212 • 201 items processed • Thursday, September 10, 2026 • 16:05 UTC. All rights reserved.