

THE GLOBAL AI GAZETTE

Worldwide AI Ecosystem Intelligence • Open Source • Research • Models • Tools • Sovereign AI

Thursday, September 10, 2026 • 20:52 UTC

EDITION NO. 2213

165 STORIES ANALYZED

■ LEAD STORY

ANOTHER researcher accuses OpenAI of training on conversations and then claiming a breakthrough

Velocity 96/100 • Impact 88/100 • Source: reddit.com

ANOTHER researcher accuses OpenAI of training on conversations and then claiming a breakthrough reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology development. Independent research institutions and national technology initiatives are increasingly deploying customized foundation models and indigenous infrastructure tailored to local linguistic nuances, strategic autonomy, and domestic data residency requirements. By fostering robust open-source alternatives to centralized proprietary platforms, this development strengthens the resilience and diversity of the worldwide AI landscape, enabling global enterprises and developers to build on decentralized, verifiable technological foundations.

INNOVATION: Model: Frontier Foundation Weights & Open Checkpoint
ACTION: Deployment: Checkpoints accessible on Hugging Face; compatible with standard runtimes.

■ SECONDARY ANCHOR

Deepseek V4.1 Flash Release Video [Made with Deepseek V4.1 Flash]

I like to benchmark new models that come out on motion videos. So here's a test I did for deepseek v4.1 flash. And I have to say flash has probably graduated from being a Luna class model to nearly an Opus class model with this release, at least with motion videos. Prev.

■ TOP FLASH BULLETINS

[VEL 96] [guide to using reasoning_effort on deepseek v4.1 flash](#) — guide to using reasoning_effort on deepseek v4.1 flash marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation....

[VEL 96] [Qwen3.8-27B-Uncensored-Genesis-V1-GGUF](#) — Model available Qwen3.8-27B-Uncensored-Genesis-V1-MTP-GGUF This model is a practical realisation of things described in this paper, but adapted by me for machine learning: I am trying...

[VEL 88] [LangChain AI Application Framework: langchain-openai==1.6.2](#) — LangChain is a framework for building agents and LLM-powered applications. It helps you chain together interoperable components and third-party integrations to simplify AI application development...

[VEL 96] [DeepSeek V4.1 Flash is available in HuggingChat](#) — DeepSeek V4.1 Flash is available in HuggingChat marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The...

[VEL 96] [Running Vision Qwen 3.8 27B on a 16GB Card, the config \(45tks\)](#). — I am just sharing my config for Qwen 3.8 27b that fits on a 5060TI, what is cool about this is that you can even...

SECTION II: EXECUTIVE AI STRATEGIC BRIEFING

Strategic Horizons & Action Directives for Enterprise Leadership

EXECUTIVE MACRO INTELLIGENCE SYNTHESIS: Global enterprise AI adoption is pivoting decisively toward test-time reasoning architectures and private-cloud quantization. Technology leadership must actively balance proprietary frontier model APIs with sovereign, open-weight deployments to reduce token expenditure while strictly sandboxing autonomous agent tool-calling boundaries.

VECTOR	DEVELOPMENT	SOURCE	DIRECTIVE
Model Sourcing & Licensing	Deepseek V4.1 Flash is 748B, not 552B	reddit.com	Audit open-weights licensing vs proprietary APIs; evaluate Deepseek pa
Compute & Infrastructure CapEx	Deepseek v4.1 flash finally has engrams,	reddit.com	Review GPU cluster allocation and power envelopes; benchmark Deepseek
Agentic Autonomy & Governance	DeepSeek V4-1 Flash is out	reddit.com	Implement deterministic sandboxes for DeepSeek autonomous tool executi
Inference Latency & Quantization	DeepSeek V4.1 Flash: Stronger, Faster, M	reddit.com	Benchmark KV-cache compression (FP8/INT4/GGUF) and modern inference en
Open-Source Supply Chain	deepseek-ai/DeepSeek-V4.1-Flash · Huggin	reddit.com	Inspect upstream repository dependencies; audit tokenizer code, weight
Data Residency & Sovereignty	Qwen3.8-Flash-Next on 2x3090 + DDR4, par	reddit.com	Verify compliance with sovereign AI frameworks and regional data resid

1. **Deepseek V4.1 Flash is 748B, not 552B:** Audit open-weights licensing vs proprietary APIs; evaluate Deepseek parameter efficiency and commercial distribution terms.
2. **Deepseek v4.1 flash finally has engrams, what do you expect from 4.1 pro?:** Review GPU cluster allocation and power envelopes; benchmark Deepseek hardware efficiency to optimize cost per token.
3. **DeepSeek V4-1 Flash is out:** Implement deterministic sandboxes for DeepSeek autonomous tool execution, strict rate limiting, and human-in-the-loop validation.
4. **DeepSeek V4.1 Flash: Stronger, Faster, More Accessible:** Benchmark KV-cache compression (FP8/INT4/GGUF) and modern inference engines for DeepSeek against TTFT SLAs.
5. **deepseek-ai/DeepSeek-V4.1-Flash · Hugging Face:** Inspect upstream repository dependencies; audit tokenizer code, weights provenance, and pinned runtime releases for deepseek.
6. **Qwen3.8-Flash-Next on 2x3090 + DDR4, part 4: 2.2-2.5x faster prefill by kicking the expert cache off the GPU while the prompt runs:** Verify compliance with sovereign AI frameworks and regional data residency requirements for Qwen3.8-Flash-Next on 2x3090 deployments.

SECTION III: TRENDING OPEN-SOURCE AI & GITHUB INNOVATIONS

Top GitHub Repositories, Developer Velocity & Architecture Dissections

PROJECT / REPO	SOURCE	STACK	VELOCITY	CORE ARCHITECTURAL FOCUS
LangChain AI Application Framework: langchain	github.com	Python	88/100	LangChain is a framework for building agents and LLM-powered applications. It helps you chain together interoperable components and.
Hugging Face Transformers: Release 5.17.0	github.com	Python	72/100	Release v5.17.0 New Model additions HYV4 Hy4-Preview is a 780B-parameter mixture-of-experts language model that activates 49B parameters per.
open-edge-platform/anomalib: An anomaly detec	github.com	Python	80/100	A library for benchmarking, developing and deploying deep learning anomaly detection algorithms** > ■ **Announcing v2.6.1 Release!**■.
deepinsight/insightface: State-of-the-art 2D	github.com	Python	80/100	InsightFace project is mainly maintained by Jia Guo and Jiankang Deng . InsightFace 2.0 adds PrivateFrame for local.
apache/seatunnel: SeaTunnel is a multimodal,	github.com	Python	80/100	## Overview SeaTunnel is a multimodal, high-performance, distributed data integration tool, capable of synchronizing vast amounts of data.
steven2358/awesome-generative-ai: A curated l	github.com	Python	80/100	> A curated list of modern Generative Artificial Intelligence projects and services. Generative Artificial Intelligence is a technology.

[TRENDING REPO] SOURCE: [github.com](#) | VEL: 88/100

LangChain AI Application Framework: langchain-openai==1.6.2

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

LangChain is a framework for building agents and LLM-powered applications. It helps you chain together interoperable components and third-party integrations to simplify AI application development — all while future-proofing decisions as the underlying technology evolves. > [!TIP] > Just getting started? Check out **Deep Agents** — a higher-level package built on LangChain for agents that have built-in capabilities for common usage patterns such as planning, subagents, file system usage, and more. If you're looking for more advanced customization or agent orchestration, check out LangGraph, our framework for building controllable agent workflows.

[TRENDING REPO] SOURCE: [github.com](#) | VEL: 72/100

Hugging Face Transformers: Release 5.17.0

ARCH: 780B • CTX: 1M • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Release v5.17.0 New Model additions HYV4 Hy4-Preview is a 780B-parameter mixture-of-experts language model that activates 49B parameters per token. Each MoE layer holds 256 routed experts plus one always-active shared expert and routes every token to 8 of them. The context window is 1M tokens. The architecture combines four features: Multi-head Latent Attention (MLA) compresses keys and values into a low-rank latent (kv_lora_rank) that kv_b_proj expands back to one key/value per query head.

[TRENDING REPO] SOURCE: [github.com](#) | VEL: 80/100

open-edge-platform/anomalib: An anomaly detection library comprising state-of-the-art algorithms and features...

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

A library for benchmarking, developing and deploying deep learning anomaly detection algorithms** > ■ **Announcing v2.6.1 Release!**■ > > This patch release migrates the Kaputt dataset to Hugging Face and fixes several data-handling and metric bugs. > > Key Changes > > - **Kaputt** on Hugging Face**: Dataset now downloads from the official Hugging Face repository instead of requiring a manual request form. > - **Security fix****: Archive extraction now blocks tar/zip path traversal on all supported Python versions. > - **Bug fixes****: Random tiling now honours the configured tile width, `seed=0` is treated as a valid seed in `random\_split`, and the binary classification curve metric is faster. > > We value your input!.

[TRENDING REPO] SOURCE: [github.com](#) | VEL: 80/100

deepinsight/insightface: State-of-the-art 2D and 3D Face Analysis Project...

ARCH: MoE / SOTA • CTX: 50M • QUANT: INT8 • STACK: CUDA

InsightFace project is mainly maintained by Jia Guo and Jiankang Deng . InsightFace 2.0 adds PrivateFrame for local video face blur and mosaic, reference-photo selection, and editable analysis JSON through desktop, CLI, and Python workflows, plus optional RGB liveness detection for FaceAnalysis and InsightFace Server. The Python package also supports raccoon_s / raccoon_l model packages, automatic CoreML/CUDA/CPU provider selection, and reusable CoreML compilation caches. PrivateFrame defaults to Fast mode (target 15 analysis FPS) . The new InsightFace Server provides a simple self-hosted Web UI, snake_case REST API, and Python client for detection, optional liveness checks, comparison, registration, and Person search.

[TRENDING REPO] SOURCE: [github.com](#) | VEL: 80/100

apache/seatunnel: SeaTunnel is a multimodal, high-performance, distributed, massive data integrated...

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Overview SeaTunnel is a multimodal, high-performance, distributed data integration tool, capable of synchronizing vast amounts of data daily. It's trusted by numerous companies for its efficiency and stability. ## Why Choose SeaTunnel SeaTunnel addresses common data integration challenges: - **Diverse Data Sources****: Seamlessly integrates with hundreds of evolving data sources. - **Multimodal Data Integration****: Supports the integration of video, images, binary files, structured and unstructured text data. - **Complex Synchronization Scenarios****: Supports various synchronization methods, including real-time, CDC, and full database synchronization. - **Resource Efficiency****: Minimizes computing resources and JDBC connections for real-time synchronization. - **Quality and Monitoring****: Provides data quality and monitoring to prevent data loss or duplication.

SECTION IV: FRONTIER FOUNDATION MODELS & REASONING
BREAKTHROUGHS

Reasoning Architectures, Open Weights & Capability Benchmarks

FOUNDATION MODEL	SOURCE	SIZE / CTX	IMPACT	ARCHITECTURAL HIGHLIGHT
guide to using reasoning_effort on deepseek v	reddit.com	MOE / 128K	88/100	guide to using reasoning_effort on deepseek v4.1 flash marks an architectural milestone in Open-Weights foundation modeling, engineered for.
Qwen3.8-27B-Uncensored-Genesis-V1-GGUF	reddit.com	27B / 128K	88/100	Model available Qwen3.8-27B-Uncensored-Genesis-V1-MTP-GGUF This model is a practical realisation of things described in this paper, but adapted by.
DeepSeek V4.1 Flash is available in HuggingCh	reddit.com	MOE / 128K	68/100	DeepSeek V4.1 Flash is available in HuggingChat marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput.
DeepSeek-V4.1-Flash surprised	reddit.com	50B / 128K	68/100	Hoping to see smartest medium size models soon & later with all available optimizations/architectures/etc.,. Thanks Deepseek!
Deepseek V4.1 Flash is 748B, not 552B	reddit.com	748B / 128K	68/100	People keep on getting confused about this, so I looked at the safetensors on hf.
Deepseek v4.1 flash finally has engrams, what	reddit.com	MoE / SOTA / 128K	68/100	If the ratio is the same, Maybe 1.6T -3.1T params plus .56T-1.06T engrams and fable 5.0 level performance?.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 88/100
guide to using reasoning_effort on deepseek v4.1 flash
ARCH: MOE • CTX: 128K • QUANT: FP8 • STACK: vLLM

guide to using reasoning_effort on deepseek v4.1 flash marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 88/100
Qwen3.8-27B-Uncensored-Genesis-V1-GGUF
ARCH: 27B • CTX: 128K • QUANT: GGUF • STACK: PyTorch / ONNX

Model available Qwen3.8-27B-Uncensored-Genesis-V1-MTP-GGUF This model is a practical realisation of things described in this paper, but adapted by me for machine learning: I am trying to solve the problem: why LLM models even for simple questions write walls of text during reasoning, and burn too much tokens instead of solving the task. And when number of parameters increase the problem became worse.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 68/100
DeepSeek V4.1 Flash is available in HuggingChat
ARCH: MOE • CTX: 128K • QUANT: FP8 • STACK: vLLM

DeepSeek V4.1 Flash is available in HuggingChat marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 68/100
DeepSeek-V4.1-Flash surprised
ARCH: 50B • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Hoping to see smartest medium size models soon & later with all available optimizations/architectures/etc.,. Thanks Deepseek! Ex 1: 30-50B MOE + 10-15B Engram + DeepSeek-V4.1-Flash type KVCache Ex 2: 15-30B Dense + 10-15B Engram + DeepSeek-V4.1-Flash type KVCache EDIT: Updated Engram to 10-15B from.

[FRONTIER MODEL] SOURCE: reddit.com | IMPACT: 68/100
Deepseek V4.1 Flash is 748B, not 552B
ARCH: 748B • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

People keep on getting confused about this, so I looked at the safetensors on hf. The title should have been "Deepseek V4.1 Flash is 748B total/552B base, not 284B or 305B or 485B or 522B" The model is not 284B. The original Deepseek V4 Flash is 284B, but not the V4.1 Flash model The model is not 305B, despite what some people claim "So: ~305B real backbone + 203B engram = 508B total" This is incorrect.

■ SECTION V: TOP AI RESEARCH PAPERS & ARXIV BREAKTHROUGHS

Reasoning Paradigms, Multimodal Architectures & Algorithmic Design

[AI RESEARCH] SOURCE: [arxiv.org](#) | VEL: 88/100
RedKnot-MLA: Multi-Head Offline-Online Reuse for DeepSeek-V4 Long-Context Serving
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Multi-head latent attention (MLA) exposes many logical query heads through one packed latent KV stream. This representation is memory efficient, but it removes the physical per-head cache boundary assumed by conventional head-wise reuse. We present our system, a DeepSeek-V4 realization of RedKnot's head-aware reuse principle. Each immutable document is processed offline at canonical position zero; certified Local-head contributions are retained as MLA-Off.

[AI RESEARCH] SOURCE: [arxiv.org](#) | VEL: 72/100
Building Multilingual Bridges: Data Mixing as the Pillar of Generalization for In-Language Reasoning
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Reasoning language models have made substantial advances on a variety of complex tasks, yet their capabilities remain overwhelmingly English-centric: models primarily reason in English regardless of the language they are prompted in. This is inaccessible for non-English-speaking users, risks losing the intent of the original question, and forgoes knowledge more readily expressed in the target language.

[AI RESEARCH] SOURCE: [news.mit.edu](#) | VEL: 72/100
MIT Schwarzman College of Computing launches pilot to help educators teach AI across disciplines
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

This summer, the MIT Schwarzman College of Computing welcomed faculty from colleges and universities across Greater Boston, South Carolina, West Virginia, and Texas to campus for the inaugural AI Educators Pilot, a weeklong workshop aimed at expanding how artificial intelligence is taught across disciplines and learning environments.

[AI RESEARCH] SOURCE: [arxiv.org](#) | VEL: 88/100
Watermarks Without Verification: AI Text Watermarking After the EU AI Act
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

On August 2, 2026, the obligations of Article 50 of the EU AI Act took effect, requiring generative AI providers to mark the content their systems produce and ensure it can be detected as AI-generated. Days later, Anthropic disclosed that every Claude model released after that date embeds a watermark based on SynthID-Text in all generated text, enabled by default with no user opt-out; Google has deployed SynthID-Text in Gemini since 2024.

[AI RESEARCH] SOURCE: [arxiv.org](#) | VEL: 88/100
A Tool-Augmented, GPT-4 Chatbot for Real-Time Repository Data Analysis
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Software repositories contain vast amounts of data on code contributions, bug reports, and project activities, yet this information remains challenging for non-technical stakeholders and developers to access due to limited expertise in querying repositories. To address this, we introduce a novel chatbot architecture leveraging OpenAI's GPT-4 model for automated extraction and analysis of repository data.

[AI RESEARCH] SOURCE: [arxiv.org](#) | VEL: 88/100
How Does mHC Use Its Residual Streams? Selective Routing and Near-Identity Mixing
ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Hyper-Connections and their manifold-constrained variant mHC widen a residual pathway from one stream to n, yet how trained models use this capacity remains unclear: how broadly blocks read and write, how strongly the residual pathway mixes streams, and whether the streams carry distinct representations. We examine these properties in the four-stream residual pathway of DeepSeek-V4-Flash using effective stream counts, cross-stream residual weights, and inter-stream cosine similarity.

■■ SECTION VI: DEVELOPER TOOLS, FRAMEWORKS & AI INFRASTRUCTURE

Inference Runtimes, Evaluation Frameworks, Vector DBs & Local Execution Engines

[AI TOOL] SOURCE: [reddit.com](#) | VEL: 96/100

Closed AI doesn't like biological research, user turns to open weight models

ARCH: MoE • CTX: 128K • QUANT: FP8 • STACK: vLLM

Closed AI doesn't like biological research, user turns to open weight models marks an architectural milestone in Frontier API foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[AI TOOL] SOURCE: [reddit.com](#) | VEL: 96/100

Don't let FOMO win if you're interested in local llm from a hobby/learning aspect

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: vllm

Just a reminder for those out there itching to get into local llms - don't let FOMO or "gear acquisition syndrom" take over. No matter the hobby, it's so easy to get stuck in a trap where we buy more trying to do more only to realize we've lost the fun in it all or even the notion of learning. Obviously, if you're into writing llama or vllm or hardware drivers or whatever - you got to do what you got to do.

[AI TOOL] SOURCE: [simonwillison.net](#) | VEL: 96/100

On the Navier–Stokes Millennium Prize Problem

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

On the Navier–Stokes Millennium Prize Problem Impressive result from OpenAI, who used an unreleased model to produce a resolution to the Navier–Stokes existence and smoothness problem, one of the seven Millennium Prize Problems that have been subject to a \$1,000,000 prize since May 24th, 2000.

[AI TOOL] SOURCE: [simonwillison.net](#) | VEL: 96/100

Introducing ChatGPT Images 2.5

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Introducing ChatGPT Images 2.5 OpenAI's image generation models are apparently used "more than 3 billion images across ChatGPT Images and the GPT-4o Image models in the API". This latest release improves their instruction-following ability across multiple turns, responds faster, and "is better at preserving the subjects in your reference photos". There are two new model IDs in the API: gpt-image-2.5-sunburst and gpt-image-2.5-flare.

[AI TOOL] SOURCE: [reddit.com](#) | VEL: 80/100

Doesn't it seem strange to you that AA changes twice in one week?

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Doesn't it seem strange to you that AA changes twice in one week? reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology development. Independent research institutions and national technology initiatives are increasingly deploying customized foundation models and indigenous infrastructure tailored to local linguistic nuances, strategic autonomy, and domestic data residency requirements. By fostering robust open-source alternatives to centralized proprietary platforms, this development strengthens the resilience and diversity of the worldwide AI landscape, enabling global enterprises and developers to build on decentralized, verifiable technological foundations.

[AI TOOL] SOURCE: [reddit.com](#) | VEL: 80/100

OUI-1: a model that generates bespoke UI elements

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

so i saw that openui.com released OUI-1, a model fine-tuned on DiffusionGemma. the training dataset uses OpenUI-Lang, a custom DSL (domain-specific language), instead of plain HTML, Markdown, or React code. what makes it interesting is that you can already get a regular LLM to use OpenUI-Lang through a system prompt, but that eats up a lot of the context window.

■ SECTION VII: SOVEREIGN AI & WORLDWIDE REGIONAL ECOSYSTEMS

National AI Labs, Regional Models & Sovereign Compute Ecosystems

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

Deepseek Has Soft Retired Deepseek V4 Pro

ARCH: MoE • CTX: 128K • QUANT: FP8 • STACK: vLLM

Deepseek Has Soft Retired Deepseek V4 Pro marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

3060 12GB vs 4060 ti 16GB

ARCH: 27B • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

I'm currently building my system around 3060s, but I might be able to get a 4060 for a nice deal. At first it seemed like a no brainer, but turns out the 4060 has lower memory bandwidth. In a system that already has 4x 3060 12GBs set up on a threadripper with tensor parallelism (mostly qwen3.8-27b), would it be worth having the 4060 ti 16GB around for the extra 4GB and occasional gaming, or is it just going to slow the rest of the setup down for AI?

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

Harness doesn't matter

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Harness doesn't matter reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology development. Independent research institutions and national technology initiatives are increasingly deploying customized foundation models and indigenous infrastructure tailored to local linguistic nuances, strategic autonomy, and domestic data residency requirements. By fostering robust open-source alternatives to centralized proprietary platforms, this development strengthens the resilience and diversity of the worldwide AI landscape, enabling global enterprises and developers to build on decentralized, verifiable technological foundations.

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 96/100

Mention if a "new model" is a finetune

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

A few posts tagged with "new model" present models that are finetunes. My opinion : I'd rather have the "new model" tag reserved for new "major" releases, like a new Qwen model, Deepseek V4 -> Deepseek V4.1, etc., that involved a new pretrain or intensive post-training (in opposition to a small finetune). Otherwise, maybe prepend "[Finetune]" to the title to indicate that the new model is "less of a big news", a use a "new finetune" tag, to differentiate between the two kinds of new models.

[SOVEREIGN [GLOBAL]] SOURCE: [simonwillison.net](#) | VEL: 80/100

.blend URL Viewer

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Tool: .blend URL Viewer I'm continuing to have a lot of fun with GPT-6 Astra and Blender (see my TIL). As a big fan of the Imperial Fabergé Easter eggs, I've always thought it would be fun to make some new ones that celebrate popular culture. Yesterday I decided to try out the new ChatGPT Images 2.5 by running this prompt: Generate a photo of a faberge egg that's themed after the TV show Pluribus - research first It gave me this - honestly not bad for a first attempt!

[SOVEREIGN [GLOBAL]] SOURCE: [reddit.com](#) | VEL: 80/100

Apple A20 Pro debuts with 7-core GPU, 32-core Neural Engine and 50% more memory bandwidth (~115 GB/s)

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

It seems to use a 96-bit LPDDR5X memory bus, instead of the previous 64-bit wide busses. Considering it's on 2nm, that's expensive silicon. That should result in around 115 GB/s memory bandwidth. A20 Pro also doubles the size of Apple's dedicated Neural Engine (from 16 to 32 cores total).

SECTION VIII: AI HARDWARE, COMPUTE CLUSTERS & SILICON

GPU Clusters, Custom Accelerators, Datacenter Infrastructure & Silicon

SILICON / SYSTEM	SOURCE	COMPUTE SPECS	VELOCITY	BENCHMARK / TELEMETRY
Running Vision Qwen 3.8 27B on a 16GB Card, t	reddit.com	High Velocity	96/100	I am just sharing my config for Qwen 3.8 27b that fits on a 5060TI, what is cool.
Qwen3.8-Flash-Next on 2x3090 + DDR4, part 4:	reddit.com	29 t/s	96/100	. 17 -> 25-29 t/s with the expert cache PR, 37-41 t/s after switching to UD-Q4_K_XL and stacking.
Qwen3.8-Flash-Next on MLX-serve, 1m context i	reddit.com	40 tok/s	96/100	Hi, I'm the co-creator of this Qwen3.8-Flash-Next engine support in MLX-serve. I've been tuning this one to run.
llama.cpp High-Performance LLM Engine: b10899	github.com	High Velocity	88/100	vulkan: small M matrix optimizations for qwen (#28457) vulkan: optimize m=1 mul_mat by swapping A/B vulkan: Improve small.
Substrate-Aware AI Agents: Execution Context	arxiv.org	High Velocity	88/100	Autonomous AI agents increasingly select actions in environments whose memory, execution-time, runtime, compute, and operational constraints determine what.
NVIDIA/Megatron-LM: Ongoing research training	github.com	High Velocity	80/100	This repository contains two components: Megatron-LM and Megatron Core . Megatron-LM is a reference example that includes Megatron.

[AI HARDWARE] SOURCE: reddit.com | VEL: 96/100

Running Vision Qwen 3.8 27B on a 16GB Card, the config (45tks).

ARCH: 27B • CTX: 85K • QUANT: GGUF • STACK: PyTorch / ONNX

I am just sharing my config for Qwen 3.8 27b that fits on a 5060TI, what is cool about this is that you can even get vision! and a 85K context (I have 1.5gb of headroom for more context or a better quant) Model: IQ3_XXS-mtp from Using beellama Config used: [*] model = ..\llm-models\Qwen3.8-27B-GSQ-RCO-IQ3_XXS-mtp.gguf mmproj = ..\llm-models\mmproj-Qwen3.8-27B-BF16.gguf image-min-tokens = 256.

[AI HARDWARE] SOURCE: reddit.com | VEL: 96/100

Qwen3.8-Flash-Next on 2x3090 + DDR4, part 4: 2.2-2.5x faster prefill by kicking the expert cache off the GPU while the prompt runs

ARCH: MoE / SOTA • CTX: 8K • QUANT: UD-Q4_K_XL • STACK: llama.cpp

. 17 -> 25-29 t/s with the expert cache PR, 37-41 t/s after switching to UD-Q4_K_XL and stacking MTP on the cache, the top-k fallback that was sorting more than it needed to. This one is all about prefill, which was honestly the weak spot the whole time. 80+ seconds before the first token on an 8k prompt, and 24 minutes on a 119k one...I know . Box is still 2x 3090, dual Broadwell Xeon, llama.cpp, UD-Q4_K_XL with the Q8 MTP head on the.

[AI HARDWARE] SOURCE: reddit.com | VEL: 96/100

Qwen3.8-Flash-Next on MLX-serve, 1m context is released!

ARCH: MoE / SOTA • CTX: 1M • QUANT: Native / FP16 • STACK: MLX

Hi, I'm the co-creator of this Qwen3.8-Flash-Next engine support in MLX-serve. I've been tuning this one to run both fast, efficient and correct up 1m context using kv cache 8 bits in M5 Max 128GB. Qwen is working well at very long context as showed in the video (a snapshot at ~760k context), I let it build MLX Serve Monitor plugin that you've seen on the right side of Opencode2's app. The generation sustain through 1m context at around 40 tok/s on prose and 75 tok/s on coding.

[AI HARDWARE] SOURCE: github.com | VEL: 88/100

llama.cpp High-Performance LLM Engine: b10899

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: llama.cpp

vulkan: small M matrix optimizations for qwen (#28457) vulkan: optimize m=1 mul_mat by swapping A/B vulkan: Improve small M perf Allow split_k with small M. Make small vs med tile selection (for coopmat2) depend on M, not just N.

[AI HARDWARE] SOURCE: arxiv.org | VEL: 88/100

Substrate-Aware AI Agents: Execution Context as a First-Class Input

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Autonomous AI agents increasingly select actions in environments whose memory, execution-time, runtime, compute, and operational constraints determine what counts as a suitable plan. We call the absence of this execution context from an agent's planning state substrate blindness. We test this general proposition through numerical code generation, where selected implementation choices and operational consequences are directly observable.

SECTION IX: AUTONOMOUS AGENTS, MULTI-AGENT SWARMS & ROBOTICS

Agent Orchestration, Computer-Use Primitives, Action Models & Robotics

AGENT / FRAMEWORK	SOURCE	PROTOCOL	VELOCITY	CORE CAPABILITY DOMAIN
Research acceleration: The view inside OpenAI	simonwillison.ne	PyTorch / ONNX	96/100	Research acceleration: The view inside OpenAI Apparently today is RSI day at OpenAI, for Recursive Self-Improvement - I.
LangChain AI Application Framework: langchain	github.com	PyTorch / ONNX	88/100	LangChain is a framework for building agents and LLM-powered applications. It helps you chain together interoperable components and.
Scanning the Harness: An Empirical Study of S	arxiv.org	PyTorch / ONNX	88/100	AI coding agents such as Claude Code, Cursor, GitHub Copilot, and OpenAI Codex are configured through artifacts developers.
What Does Multi-Harness RL Learn? Credit Assi	arxiv.org	PyTorch / ONNX	88/100	Agent reinforcement learning (RL) increasingly runs through full execution harnesses, and a multi-harness recipe mixes two choices: exposing.
JarvisGUI: Towards Cross-Device GUI Agents wi	arxiv.org	PyTorch / ONNX	72/100	Real-world GUI usage frequently involves workflows that span multiple devices and platforms, requiring the transfer of intermediate results.
memgraph/memgraph: High-performance open-sour	github.com	PyTorch / ONNX	80/100	Memgraph is a high-performance, in-memory graph database that powers real-time AI context and graph analytics.

[AGENT SYSTEM] SOURCE: [simonwillison.net](#) | VEL: 96/100

Research acceleration: The view inside OpenAI

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Research acceleration: The view inside OpenAI Apparently today is RSI day at OpenAI, for Recursive Self-Improvement - I think it's their new AGI. Both this piece and the new essay An Alien Mind (by Chief Scientist Jakub Pachocki) talk about it, and this one doesn't even bother to expand the acronym. Included are details on how OpenAI's own research team are using coding agents.

[AGENT SYSTEM] SOURCE: [github.com](#) | VEL: 88/100

LangChain AI Application Framework: langchain-anthropic==1.7.2

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

LangChain is a framework for building agents and LLM-powered applications. It helps you chain together interoperable components and third-party integrations to simplify AI application development — all while future-proofing decisions as the underlying technology evolves. > [!TIP] > Just getting started? Check out ****Deep Agents**** — a higher-level package built on LangChain for agents that have built-in capabilities for common usage patterns such as planning, subagents, file system usage, and more. If you're looking for more advanced customization or agent orchestration, check out LangGraph, our framework for building controllable agent workflows.

[AGENT SYSTEM] SOURCE: [arxiv.org](#) | VEL: 88/100

Scanning the Harness: An Empirical Study of Supply-Chain Defects in AI Coding-Agent Configurations

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

AI coding agents such as Claude Code, Cursor, GitHub Copilot, and OpenAI Codex are configured through artifacts developers write and share: instruction files, skills, hooks, MCP server declarations, subagents. This harness is a dependency layer installed from marketplaces and public repositories, running with the developer's privileges, with no lockfile, no install-time check, and no vocabulary for what a component may do.

[AGENT SYSTEM] SOURCE: [arxiv.org](#) | VEL: 88/100

What Does Multi-Harness RL Learn? Credit Assignment and Portability in Coding Agents

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Agent reinforcement learning (RL) increasingly runs through full execution harnesses, and a multi-harness recipe mixes two choices: exposing the policy to several harnesses, and comparing their rewards inside one relative-advantage group. We isolate the second choice in repository-level coding.

[AGENT SYSTEM] SOURCE: [arxiv.org](#) | VEL: 72/100

JarvisGUI: Towards Cross-Device GUI Agents with Dynamic Task Composition

ARCH: MoE / SOTA • CTX: 128K • QUANT: Native / FP16 • STACK: PyTorch / ONNX

Real-world GUI usage frequently involves workflows that span multiple devices and platforms, requiring the transfer of intermediate results, maintenance of shared state, and coordination across heterogeneous environments. However, existing GUI benchmarks overwhelmingly evaluate agents on single-device, statically defined tasks, thus leaving such cross-device capabilities largely unexamined, resulting in an overly optimistic assessment of agents' readiness for real-world usage.

■ SECTION X: GLOBAL AI COMMUNITY WIRE & OVERFLOW DIGEST

High-velocity AI stories from today's sweep that didn't fit earlier sections

[OVERFLOW] **SOURCE:** pandaily.com | **VEL:** 88/100

DeepSeek Ships V4.1 Flash GA With Causal-Encoder-Decoder MoE as V4 Pro Retires

ARCH: MOE • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

DeepSeek has released DeepSeek V4.1 Flash for general availability, promoting the model from a short limited beta into a production SKU on a new Causal-Encoder-Decoder Mixture-of-Experts architecture. The company presents V4.1 Flash as the smallest member of that structure family and says it now surpasses DeepSeek V4 Pro on capability, cost, speed and end-to-end task time, clearing the path for an orderly Pro retirement later this week. Architecturally, V4.1 Flash is a 552-billion-parameter.

[OVERFLOW] **SOURCE:** marktechpost.com | **VEL:** 88/100

DeepSeek AI Released DeepSeek-V4.1-Flash with 1M Context, FP4 KV Cache, and Cross-Layer Attention Reuse

ARCH: 552B • **CTX:** 1M • **QUANT:** FP4 • **STACK:** PyTorch / ONNX

Long-horizon agents have turned LLM serving into an input-heavy workload. Repeated prefills and million-token contexts leave KV caches that strain HBM, SSD capacity, and bandwidth. DeepSeek AI built its newest release around that exact bottleneck. DeepSeek-V4.1-Flash is a multimodal Mixture-of-Experts model with 552B backbone parameters, 196B additional Engram parameters, and a 1M-token context window. It activates 8B parameters per token during prefill and 16B during decode.

[OVERFLOW] **SOURCE:** pandaily.com | **VEL:** 88/100

Alibaba Opens Qwen3.8-2.4T-A95B Weights as First Qwen-Max-Class MoE

ARCH: MOE • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Alibaba's Qwen team has published open weights for Qwen3.8-2.4T-A95B, describing it as the first Qwen-Max-class model released for download rather than API-only access. The sparse mixture-of-experts checkpoint holds about 2.4 trillion total parameters and activates roughly 95 billion per token across 512 experts, with 10 routed experts plus one shared expert used at each step. The release is distinct from earlier Qwen open drops such as Qwen-Drive and smaller Qwen3.8 coding snapshots, positionin.

[OVERFLOW] **SOURCE:** digitimes.com | **VEL:** 88/100

DeepSeek reportedly advances Shanghai IPO plans as it cuts Flash model API prices

ARCH: MOE • **CTX:** 128K • **QUANT:** FP8 • **STACK:** vLLM

DeepSeek reportedly advances Shanghai IPO plans as it cuts Flash model API prices marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[OVERFLOW] **SOURCE:** digitimes.com | **VEL:** 88/100

Samsung SDS partners with OpenAI and Anthropic in AI push

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Samsung SDS partners with OpenAI and Anthropic in AI push reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology development. Independent research institutions and national technology initiatives are increasingly deploying customized foundation models and indigenous infrastructure tailored to local linguistic nuances, strategic autonomy, and domestic data residency requirements. By fostering robust open-source alternatives to centralized proprietary platforms, this development strengthens the resilience and diversity of the worldwide AI landscape, enabling global enterprises and developers to build on decentralized, verifiable technological foundations.

[OVERFLOW] **SOURCE:** digitimes.com | **VEL:** 88/100

DeepSeek's next AI test is not the model; it's everything around it

ARCH: MOE • **CTX:** 128K • **QUANT:** FP8 • **STACK:** vLLM

DeepSeek's next AI test is not the model; it's everything around it marks an architectural milestone in Open-Weights foundation modeling, engineered for high-throughput reasoning and featuring extended context evaluation. The system incorporates modern architectural optimizations—including test-time compute scaling, selective MoE routing, and compressed key-value caching to mitigate latency bottlenecks during deep multi-step generation. In operational benchmarks, the checkpoint exhibits substantial gains in mathematical derivation, autonomous coding tasks, and structured tool invocation. Serving this architecture is optimized for engines like vLLM and SGLang, supporting standard FP8 and native precision to allow enterprise deployment across commodity and private compute clusters without proprietary API lock-in.

[OVERFLOW] **SOURCE:** digitimes.com | **VEL:** 88/100

Firmus, OpenAI strike compute deal for Malaysia AI factories

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Firmus, OpenAI strike compute deal for Malaysia AI factories reflects the rapid acceleration of sovereign artificial intelligence ecosystems and decentralized technology development. Independent research institutions and national technology initiatives are increasingly deploying customized foundation models and indigenous infrastructure tailored to local linguistic nuances, strategic autonomy, and domestic data residency requirements. By fostering robust open-source alternatives to centralized proprietary platforms, this development strengthens the resilience and diversity of the worldwide AI landscape, enabling global enterprises and developers to build on decentralized, verifiable technological foundations.

[OVERFLOW] **SOURCE:** wired.com | **VEL:** 88/100

The AI Researcher Who Just Quit Anthropic Says It's 'Crunch Time for Humanity'

ARCH: MoE / SOTA • **CTX:** 128K • **QUANT:** Native / FP16 • **STACK:** PyTorch / ONNX

Artificial intelligence researcher Jacob Coxon sent shock waves through Silicon Valley and beyond on Tuesday by announcing his resignation from Anthropic and delivering a grave warning that the AI race is putting all of our lives at risk. In his post on X, which now has more than 100 million views, Coxon wrote that many of the people building AI share his views and believe time is running out to ensure AI systems are built safely. "The consensus is that the next year or two is crunch time for humanity," Coxon, who worked on the pretraining stage of AI development, said in an interview with WIRED.

COLOPHON & SENSOR METHODOLOGY: The Global AI Gazette is compiled autonomously. Data is aggregated across authoritative global AI sources including Hugging Face, GitHub Trending, arXiv, and leading AI research lab feeds. Edition #2213 • 165 items processed • Thursday, September 10, 2026 • 20:52 UTC. All rights reserved.