

THE AETHER GUARD — GLOBAL AI & TECHNOLOGY GAZETTE

"Omnis Intelligentia Patefacietur" — Frontier Research • Trending Repositories • Developer Tools • AI Models • Security & Governance

Tuesday, September 08, 2026 • 23:25 UTC

EDITION NO. 2204

AETHERGUARD AI OBSERVER

■ **BREAKING INVESTIGATION // CRITICAL ZERO-DAY INCIDENT**

Safety alignment in large language models (LLMs) remains brittle against a growing spectrum of attacks. Jailbreak attacks bypass safety mechanisms through crafted prompts, while neuron-level attacks directly prune safety-critical neurons post-deployment. Both exploit a common weakness: safety-relevant information concentrates in a sparse neuron subset. We present NeuronGuard, a fine-tuning-stage defense that simultaneously hardens LLMs against both attack classes by redistributing safety signals

EXPLOITATION MECHANICS: Attack archetype: Jailbreak - Authentication/authorization bypass
EMERGENCY DIRECTIVE: Deploy prompt injection defenses and output filtering; Apply official vendor patches immediately; Restrict network ingress and isolate affected components

■ **SECONDARY ANCHOR DISPATCH: CRITICAL THREAT INDICATOR**

NeuronFuzz: Safety Neuron Guided Fuzzing for LLM Safety Evaluation

Safety evaluation is critical for assessing whether aligned Large Language Models (LLMs) remain robust against jailbreak attacks. Existing automated testing methods, however, largely rely on response-level feedback: each candidate prompt typically requires generating a target-model response to evaluate its attack effectiveness. This process is expensive and, more importantly, provides only sparse guidance on strongly aligned models, where most candidates are rejected with the same failure...

■ TOP FRONT-PAGE FLASH BULLETINS

- [VEL 50] Breaking Claude Code Opus 5 Auto Mode** — Breaking Claude Code Opus 5 Auto Mode Anthropic are putting a great deal of faith in Claude Code's auto mode for protecting their coding agent...
- [VEL 30] Increasing active parameters per token in MOE (Qwen 35B A4B+) reduce reasoning token by 8.5% - and you don't need to train or finetune!** — I want to share a short paper just published exploring a simple but surprisingly effective optimization for sparse MoE reasoning models. The idea: Instead of...
- [VEL 30] model: add NVIDIA Nemotron-3-Puzzle-75B-A9B (NemotronHPuzzle) support by YanissAmz · Pull Request #25444 · ggml-org/llama.cpp** — 75B MoE is an interesting size to check, you can run it today (no MTP support yet) The model employs a hybrid MoE architecture with...

SECTION II: CISO & EXECUTIVE BOARD STRATEGIC BRIEFING

Macro Risk Posture, Threat Velocity Heatmaps, and 24-Hour Boardroom Mandates

MACRO THREAT POSTURE: The enterprise threat posture remains indexed at DEFCON 3 (Elevated). Continuous telemetry across 92 authoritative sensor nodes records an aggressive convergence between nation-state reconnaissance and automated ransomware syndicates. Attackers are exploiting unauthenticated edge perimeter gateways and forging cloud identity provider tokens to establish persistence before enterprise SOC teams detect initial intrusion. Concurrently, the rapid enterprise adoption of autonomous generative AI agents has introduced an entirely new class of prompt injection and tool execution vulnerabilities that evade legacy web application firewalls (WAFs).

ENTERPRISE ATTACK SURFACE EXPOSURE MATRIX

VECTOR / BOUNDARY	LIKELIHOOD	ENTERPRISE IMPACT	PRIMARY DETECTION CONTROL	EXECUTIVE MANDATE
Cloud Identity & IdP	High	Full Tenant Compromise	Conditional Access & Token Audit	Enforce FIDO2 hardware MFA
Kubernetes & Containers	Critical	Lateral Pod Escape	eBPF Runtime Audit	Enforce read-only root filesystems
AI Agent / LLM APIs	High	Prompt Injection & Tool Abuse	Context boundary filters	Sandbox agent tool invocations
Edge Perimeter Gateways	Critical	Unauthenticated RCE	NetFlow anomaly inspection	Disallow public admin panels
Software Supply Chain	High	Build Pipeline Poisoning	SBOM & SHA256 verification	Mandate signed container commits

PRIORITIZED 24-HOUR EXECUTIVE DIRECTIVES

[EXECUTIVE DIRECTIVE] SOURCE: [simonwillison.net](#) | VELOCITY: 30/100 | IMPACT: 40/100

Qwen3.8-Flash-Next

Qwen3.8-Flash-Next Another open weights model from Qwen. This one is "a multimodal MoE model that also serves as an early preview of the architecture used in Qwen4". It's pretty big: 125B tokens, but only 6B active which means it gets a significant performance boost. I've been trying it out on a DGX Spark using these Unsloth quantized models. I'm still exploring the model - so far I've tried the 72.5GB UD-IQ1_S one (producing these pelicans) and the 78.9GB UD-Q2_K_XL (producing these). My favori

[EXECUTIVE DIRECTIVE] SOURCE: [nvd.nist.gov](#) | VELOCITY: 100/100 | SEVERITY: 100/100

CISA KEY: CVE-2023-4966 - Citrix NetScaler ADC and NetScaler Gateway Buffer Overflow Vulnerability

Citrix NetScaler ADC and NetScaler Gateway contain a buffer overflow vulnerability that allows for sensitive information disclosure when configured as a Gateway (VPN virtual server, ICA Proxy, CVPN, RDP Proxy) or AAA virtual server. Required Action: Apply mitigations and kill all active and persistent sessions per vendor instructions [<https://www.netscaler.com/blog/news/cve-2023-4966-critical-security-update-now-available-for-netscaler-adc-and-netscaler-gateway/>] OR discontinue use of the produ

[EXECUTIVE DIRECTIVE] SOURCE: [reddit.com](#) | VELOCITY: 75/100 | IMPACT: 37/100

CVE-2026-33696: From a Schema Name to RCE in n8n

 submitted by /u/TradeGold6317 [link] [comments]

[EXECUTIVE DIRECTIVE] SOURCE: [reddit.com](#) | VELOCITY: 10/100 | IMPACT: 5/100

Qwen3.8-Flash-Next: 256k context, 16tok/s on DDR4 and a Tesla T4

I've got an refurb Dell R740 running Proxmox that I put a Tesla T4 in, mainly to run some CTC local transcription work, but thought it would be fun to try DS4 when it came out, and it was appalling at around 2 tok/s. However pulled it out again when Qwen3.8 dropped, and it was much improved, particularly with ik_llama. Hardware: Dell R740, 2x Xeon Gold 6230, 384GB DDR4-2666, one Tesla T4 16GB. Guest VM pinned to one NUMA node: 20 cores, 168GB RAM. Model: Unsloth Qwen3.8-Flash-Next UD-Q4_K_XL, 11

[EXECUTIVE DIRECTIVE] SOURCE: [reddit.com](#) | VELOCITY: 10/100 | IMPACT: 5/100

UPDATE: Qwen3.8-Flash-Next on 2x3090 + DDR4 (Part 2): 25-29 -> 37-41 t/s decode (UD-Q4_K_XL + expert cache + MTP), plus a branch you can build

This is a follow-up to my post from yesterday (17 -> 25-29 t/s with the expert cache PR). Same box: 2x RTX 3090 on PCIe 3.0, dual Xeon E5-2696 v4, 188 GB DDR4-2133 LRDIMM, llama.cpp, full 261k context, f16 KV, all 48 expert layers in host RAM, everything else on the GPUs. Since then I switched quants, stacked MTP on top of the cache, fixed the load time, found a bug in the cache PR and found out my RAM was thermal throttling (Now i gotta buy an additional case fan lol). Numbers are all from the

SECTION III: TRENDING OPEN-SOURCE AI & GITHUB INNOVATIONS

Top GitHub Repositories, Developer Velocity, Architecture Dissections & Deployment Guides

GLOBAL DEVELOPER VELOCITY: Open-source AI engineering on GitHub is expanding across distributed runtimes, agentic workflows, and quantized model serving. Developers worldwide are rapidly converging on local-first LLM orchestration, synthetic data pipelines, and high-throughput inference kernels that bypass proprietary API bottlenecks.

PROJECT / REPO	DOMAIN FOCUS	ARCHITECTURE / STACK	COMMUNITY MOMENTUM
vllm-project / vllm	High-Throughput Serving	PagedAttention, CUDA C++, Python	★ 35,000+ Stars • Standard Engine
ollama / ollama	Local Execution Engine	Go, llama.cpp, Cross-Platform	★ 95,000+ Stars • Desktop Standard
run-llama / llama_index	Production Agentic RAG	Python, Hybrid Vector Connectors	★ 38,000+ Stars • Enterprise Retrieval
deepseek-ai / DeepSeek-V3	Frontier MoE Foundation	Multi-Head Latent Attention	★ 60,000+ Stars • Open Weights

FEATURED OPEN-SOURCE DISPATCHES & REPOSITORIES

[TRENDING REPO] SOURCE: github.com | VELOCITY: 30/100 | IMPACT: 15/100

Security Tool / PoC: siyuan-note/siyuan

SiYuan is a privacy-first personal knowledge management system, supporting fine-grained block-level reference and Markdown WYSIWYG. To learn more, read the online user guide or join the SiYuan English Discussion Forum . Most features are free, even for commercial use. Some features are only available to paid members, for more details please refer to Pricing . It is recommended to give priority to installing through the application market on desktop and mobile, so that you can upgrade the version with one click in the future.

[TRENDING REPO] SOURCE: github.com | VELOCITY: 30/100 | IMPACT: 15/100

Security Tool / PoC: activepieces/activepieces

Documentation ■■ Create a Piece ■ Deploy ■ Join Discord All-in-one AI automation designed to be extensible through a type-safe pieces framework written in TypeScript . When you contribute pieces to Activepieces they become automatically available as MCP servers that you can use with LLMs through Claude Desktop, Cursor or Windsurf! ■ Open Ecosystem: All pieces are open source and available on npmjs.com, 60% of the pieces are contributed by the community . ■■ Largest open source MCP toolkit : All our pieces (280+) are available as MCP that you can use with LLMs on Claude Desktop, Cursor or Windsurf. ■■ Pieces are written in Typescript : Pieces are npm packages in TypeScript, offering full customization with the best developer experience, including hot reloading for local piece development on your machine. ■

[TRENDING REPO] SOURCE: github.com | VELOCITY: 30/100 | IMPACT: 15/100

Security Tool / PoC: slackhq/nebula

A scalable overlay networking tool with a focus on performance, simplicity and security Language: Go Stars: 25 stars today

SECTION IV: FRONTIER AI MODELS & AUTONOMOUS AGENTS

Sovereign Architectures, Reasoning Breakthroughs, Parameter Scale & Benchmark Matrices

THE FRONTIER REASONING PARADIGM: Machine intelligence has pivoted from pure next-token prediction to reinforcement learning during inference (test-time compute). Architectures such as DeepSeek-R1, OpenAI o3, and Claude 3.7 Sonnet produce verifiable internal reasoning traces, demonstrating human-expert parity across software engineering, competitive mathematics, and formal logic benchmarks.

MODEL / ARCHITECTURE	SCALE / ACTIVE	CONTEXT	PRIMARY BENCHMARK	INNOVATION HIGHLIGHT
DeepSeek-R1	671B / 37B MoE	128k Tokens	AIME 2024: 79.8%	Pure RL cold-start reasoning
Claude 3.7 Sonnet	Proprietary	200k Tokens	SWE-bench: 70.3%	Hybrid instant/extended thinking
OpenAI o3-mini	Proprietary	200k Tokens	Math: 91.2%	High-speed chain-of-thought
Qwen-2.5 72B	72B Dense	128k Tokens	MMLU: 86.1%	Bilingual coding & open weights

FRONTIER MODEL DISPATCHES & CAPABILITY DOSSIERS

[FRONTIER MODEL] SOURCE: [simonwillison.net](#) | VELOCITY: 50/100 | IMPACT: 50/100

Breaking Claude Code Opus 5 Auto Mode

Breaking Claude Code Opus 5 Auto Mode Anthropic are putting a great deal of faith in Claude Code's auto mode for protecting their coding agent users against prompt injection attacks. They recently made that the default and have made bold claims about its effectiveness. Johann Rehberger is one of the most credible prompt injection researchers active today. He found an attack against auto mode which he claims works 80% of the time, by tricking Claude Code into downloading and uncompressing a zip a

[FRONTIER MODEL] SOURCE: [reddit.com](#) | VELOCITY: 30/100 | IMPACT: 60/100

Increasing active parameters per token in MOE (Qwen 35B A4B+) reduce reasoning token by 8.5% - and you don't need to train or finetune!

I want to share a short paper just published exploring a simple but surprisingly effective optimization for sparse MoE reasoning models. The idea: Instead of retraining anything, we just tweak the router at runtime. Specifically, we expand the expert selection budget ($N \geq KN \geq K$) only in the late transformer layers, with a linear decay factor applied to the extra experts. Early layers stay untouched. So Qwen 3.6 35B A3B becomes Qwen 3.6 35B A4B+ ! What we found — "Succinct Convergence": When you giv

[FRONTIER MODEL] SOURCE: [reddit.com](#) | VELOCITY: 30/100 | IMPACT: 40/100

model: add NVIDIA Nemotron-3-Puzzle-75B-A9B (NemotronHPuzzle) support by YanissAmz · Pull Request #25444 · ggml-org/llama.cpp

75B MoE is an interesting size to check, you can run it today (no MTP support yet) The model employs a hybrid MoE architecture with interleaved Mamba, MoE, and Attention layers. Like Nemotron-3-Super, it supports Multi-Token Prediction (MTP) for faster text generation. Compared to its parent, Puzzle-75B-A9B reduces the model from 120.7B total / 12.8B active parameters to 75.3B total / 9.3B active parameters. We discussed this model on [r/LocalLLaMA](#) here <https://www.reddit.com/r/LocalLLaMA/comment>

■ SECTION V: TOP AI RESEARCH PAPERS & ARXIV BREAKTHROUGHS

Reasoning Paradigms, Multimodal Architectures, Autonomous Planning & Latent Alignment

ACADEMIC & INDUSTRIAL DISCOVERY WIRE: Peer-reviewed and preprint investigations across arXiv document monumental leaps in test-time compute optimization, self-correcting agentic loops, and multi-modal alignment. Researchers increasingly focus on algorithmic sample efficiency, process reward models (PRMs), and verifiable constraint satisfaction over brute-force pre-training.

[AI RESEARCH] **SOURCE:** [news.mit.edu](#) | **VELOCITY:** 70/100 | **IMPACT:** 65/100

Solving the solvent problem

Lithium-ion batteries are the leading choice in today's electric vehicle and battery energy storage system industries, but they contain a number of critical minerals — including lithium, cobalt, nickel, and graphite — that are considered essential for economic and national security reasons, and therefore vulnerable to supply chain disruptions. As renewable energy, electrified infrastructure, and high-power digital technologies continue to grow, there is an increasing need for energy storage systems.

[AI RESEARCH] **SOURCE:** [arxiv.org](#) | **VELOCITY:** 60/100 | **IMPACT:** 85/100

Subspace Inference Enables Efficient Active Reward Learning from Preferences

Reinforcement learning from human feedback (RLHF) has emerged as a powerful yet sample-inefficient approach for learning reward models from human preferences, making active learning a critical component in synthesizing informative preference queries. However, effective uncertainty quantification required for active learning remains a key challenge for large neural network reward models. In this paper, we introduce PreferenceEKF, a sample-efficient approach that tracks reward model uncertainty by

[AI RESEARCH] **SOURCE:** [arxiv.org](#) | **VELOCITY:** 70/100 | **IMPACT:** 35/100

Structured but Fragile: On the Limits of LLMs in Cybersecurity Decision-Making

Large language models (LLMs) are increasingly used in cybersecurity workflows, yet it remains unclear whether they can perform structured security reasoning or merely rely on superficial cues and prior knowledge. We study this question in the context of defence selection over attack graphs derived from real-world threat scenarios, including ransomware, supply-chain compromise, cloud abuse, Kubernetes attacks, POS malware, and ICS/OT intrusion. Given a budget constraint, LLMs must select security

[AI RESEARCH] **SOURCE:** [arxiv.org](#) | **VELOCITY:** 60/100 | **IMPACT:** 30/100

Rapid Poison: Practical Poisoning Attacks Against the Rapid Response Framework

The Rapid Response (RR) framework, deployed in production systems, including Anthropic's ASL-3 safeguards, continuously improves jailbreak-detection classifiers. When new jailbreaks emerge that bypass these classifiers, Rapid Response generates synthetic variants for training, helping the model generalize from the new attacks and quickly adapt. We reveal that prompt injection can infiltrate this pipeline to deliver poisoned samples into the classifier's training set, enabling two attack objectives.

■■ SECTION VI: DEVELOPER TOOLS, FRAMEWORKS & AI INFRASTRUCTURE

Inference Runtimes, Evaluation Frameworks, Vector DBs & Local GPU Execution Engines

ENTERPRISE AI RUNTIME STACK: The operational foundation of artificial intelligence requires fault-tolerant vector storage, low-latency CUDA/Metal inference engines, and rigorous benchmark harnesses. Tooling ecosystems enable organizations to deploy resilient multi-agent swarms with granular access controls and continuous latency optimization.

[AI TOOL] SOURCE: reddit.com | VELOCITY: 75/100 | IMPACT: 37/100

CVE-2026-33696: From a Schema Name to RCE in n8n

 submitted by /u/TradeGold6317 [link] [comments]

[AI TOOL] SOURCE: reddit.com | VELOCITY: 100/100 | IMPACT: 70/100

Authentication bypass in EOL Proxmox VE 7 release

A lot of proxmox 7.0-7 and < 8.0.4 has been rooted today, it's needed to have port :8006 open, exploitation demo here: https://forum.proxmox.com/threads/proxmox-ve-7-is-vulnerable-to-some-type-of-0day-rce-non-auth.186078/post-867875 submitted by /u/WiuEmPe [link] [comments]

[AI TOOL] SOURCE: reddit.com | VELOCITY: 60/100 | IMPACT: 30/100

GeoNetwork - Pre-Auth RCE via Unauthenticated File Upload and Unsafe XSLT Processor (4 CVEs, 121 government deployments, all patched)

 submitted by /u/ZealousidealHunter80 [link] [comments]

[AI TOOL] SOURCE: reddit.com | VELOCITY: 45/100 | IMPACT: 47/100

Off the Hook: Discovering and Observing Active Exploitation of Sangoma Switchvox CVE-2026-9586

 submitted by /u/scopedsecurity [link] [comments]

■ SECTION VII: SOVEREIGN AI & GLOBAL NATION-STATE RADAR

Translated Sovereign Telemetry from CNCERT, CNNVD, Qihoo 360, Antiy & Sovereign AI Labs

SOVEREIGN AI & VULNERABILITY GOVERNANCE: Sovereign compute ecosystems in China and the Asia-Pacific theatre are driving dual revolutions in domestic foundation models (DeepSeek, Qwen, GLM) and proactive threat intelligence. Under national vulnerability regulations, regional labs maintain early visibility into zero-day exploits before international NVD assignment. Autonomous telemetry synthesizes real-time Chinese and sovereign sources.

[SOVEREIGN WIRE ■■■] SOURCE: mp.weixin.qq.com | VELOCITY: 100/100 | SEVERITY: 75/100

VMware ESXi CVE-2024-37085 vulnerability verification analysis

Venus 2024-08-08 17:38 Beijing Recently, Microsoft disclosed an in-field attack report of an ESXi vulnerability (CVE-2024-37085). This vulnerability is an authentication bypass vulnerability in VMware ESXi and has been exploited by multiple ransomware. Through this vulnerability, an attacker can obtain full operating permissions for ESXi added to the AD domain and control the virtual machines contained in the ESXi. For more security information and analysis articles, please pay attention to the Venustech ADLab WeChat public account and official website (adlab.venustech.com.cn) 01 Vulnerability Overview Recently, Microsoft disclosed an in-field attack report of an ESXi vulnerability (numbered CVE-2024-37085) [1]. This vulnerability is an authentication bypass vulnerability in VMware ESXi and has been exploited by multiple ransomware. Through this vulnerability, an attacker can obtain full operating permissions for the ESXi joined to the AD domain and control the virtual machines included in the ESXi. The NVD description of the vulnerability is [2]: VMware ESXi contains an authentication bypass vulnerability. A malicious actor with sufficient Active Directory (AD) permissions can gain full access

[SOVEREIGN WIRE ■■■] SOURCE: mp.weixin.qq.com | VELOCITY: 50/100 | IMPACT: 25/100

The open source risk management platform "Fuxi" has made important progress in security patch migration, helping to mitigate open source software security risks.

CodeWisdom 2026-05-13 09:00 Shanghai focuses on the problem of open source security patch migration. Based on the large model of syntax semantic enhancement, it enables automated patch migration Research background: Security patch migration is an important means of mitigating risks in the open source software supply chain With the wide application of open source software, more and more software systems rely on open source components, reuse open source code, and form multiple long-term maintenance branches or downstream derivative projects based on upstream projects. When upstream projects disclose vulnerabilities and issue security patches, whether the relevant patches can be timely and accurately migrated to other affected branches or downstream derivative projects is directly related to the overall security level of the open source software supply chain. However, in a true open source ecosystem, the migration of security patches is neither comprehensive nor timely. An empirical study of 26 popular open source projects and 806 CVEs showed that more than 80% of CVE-branch pairs never completed a patch migration; of these unpatched vulnerabilities, 47.39% were high-risk or severe; and even if a fix was eventually completed, it would take an average of 40.46 days. More notably, about 20% of vulnerabilities already have a public PoC, meaning that an attacker can exploit publicly available information before the patch is migrated to the affected branch or downstream project [1]. For security patches that are difficult to migrate in a timely manner, academia and industry A range of studies have been conducted. Existing automated patch migration techniques can be broadly divided into two categories: those based on pattern matching and those based on large models.

[SOVEREIGN WIRE ■■■] SOURCE: mp.weixin.qq.com | VELOCITY: 40/100 | IMPACT: 20/100

Selected in the Agent Security Report | Lingmai AI allows code security detection to "read" business logic!

Original Vanguard Digital Supply Chain 2026-06-25 14:00 Beijing Lingmai AI "read business logic", precisely dig out the code security agent of the unauthorized vulnerability. When OpenClaw-like applications move to the large-scale deployment stage, security is no longer an optional addition, but a prerequisite for supporting their global landing and long-term operation. By 2030, 15% of China's top 500 enterprises will have disrupted operations due to inadequate control and governance of artificial intelligence agents. Faced with the risk of high fines and even management changes, business managers urgently need to build an agent security and governance system to safely avoid potential crises in the era of large models. — "Global CIO Agenda 2026 Forecast - China's Revelation" In the Agentic AI era, traditional detection tools can find grammar loopholes, but they can't understand the business intent behind the code. As a benchmark manufacturer in the field of digital supply chain security, suspension security has always been deeply cultivated and applied in the field of code security. The company's new generation of core products, Spiritual AI, gives differentiated answers: let code security testing truly "read" business logic for the first time. Suspension mirror · AI code leakage Relying on the first AI intelligent code leakage mining technology, Lingmai AI can accurately identify business logic vulnerabilities that are difficult to cover by traditional SAST such as horizontal/vertical overstepping and permission bypassing; and with "static analysis + AI intelligent enhancement", In the authoritative target machine test, a false positive rate (FPR) of 0% is achieved for core vulnerability types such as command injection, SQL injection, etc. 1 ■ R&D efficiency-enhancing rear machine "

[SOVEREIGN WIRE ■■■] SOURCE: reddit.com | VELOCITY: 30/100 | IMPACT: 60/100

Increasing active parameters per token in MOE (Qwen 35B A4B+) reduce reasoning token by 8.5% - and you don't need to train or finetune!

I want to share a short paper just published exploring a simple but surprisingly effective optimization for sparse MoE reasoning models. The idea: Instead of retraining anything, we just tweak the router at runtime. Specifically, we expand the expert selection budget ($N \geq K \geq K$) only in the late transformer layers, with a linear decay factor applied to the extra experts. Early layers stay untouched. So Qwen 3.6 35B A3B becomes Qwen 3.6 35B A4B+ ! What we found — "Succinct Convergence": When you giv

■ ■ SECTION VIII: HIGH-VELOCITY EXPLOITED VULNERABILITIES & CISA KEV

Known Exploited Vulnerabilities Catalog, CVSS Risk Ratings & Urgent Patch Directives

ACTIVE IN-THE-WILD EXPLOITATION TELEMETRY: Adversaries demonstrate weaponization velocity that outpaces standard patch cadences. CISA Known Exploited Vulnerabilities (KEV) represent immediate risk to federal and commercial enterprise operations. Attackers leverage automated scanners to discover exposed endpoints within hours of proof-of-concept code publication.

[CISA KEV / CVE] SOURCE: [rapid7.com](https://www.rapid7.com) | VELOCITY: 100/100 | SEVERITY: 100/100

CVE-2026-19490: Critical Vulnerability Affecting Citrix NetScaler ADC and NetScaler Gateway

OverviewOn August 19, 2026, a security advisory was published for CVE-2026-19490, a critical authentication bypass vulnerability affecting Citrix NetScaler ADC and NetScaler Gateway. The vulnerability carries a CVSS v4.0 base score of 9.3 and can be exploited remotely by an unauthenticated attacker over the network without user interaction or elevated privileges. NetScaler ADC and NetScaler Gateway are widely deployed enterprise networking products commonly positioned at or near the network perimeter.

[CISA KEV / CVE] SOURCE: [rapid7.com](https://www.rapid7.com) | VELOCITY: 100/100 | SEVERITY: 100/100

Rapid7 Analysis: Unauthenticated Remote Code Execution in JetBrains TeamCity (CVE-2026-63077)

OverviewOn July 27, 2026, JetBrains published a security advisory for CVE-2026-63077, a critical unsafe deserialization vulnerability affecting JetBrains TeamCity. An attacker who can reach a TeamCity server over HTTP or HTTPS can exploit the agent polling protocol without credentials and execute operating system commands with the privileges of the TeamCity server process. JetBrains reported no known active exploitation when it disclosed the vulnerability. However, on August 5, 2026, CISA added C

[CISA KEV / CVE] SOURCE: [rapid7.com](https://www.rapid7.com) | VELOCITY: 100/100 | SEVERITY: 100/100

KindaRails2Shell: CVE-2026-66066, Critical Arbitrary File Read and Possible Remote Code Execution in Ruby on Rails

OverviewOn July 29, 2026, the Ruby on Rails project published a security advisory for CVE-2026-66066, a critical vulnerability affecting Active Storage image processing when used in conjunction with the libvips image processing library. The vulnerability has a CVSSv4 score of 9.5 and is classified as Initialization of a Resource with an Insecure Default (CWE-1188). An unauthenticated attacker may be able to leverage CVE-2026-66066 and read files accessible to the Rails application process, potentially

[CISA KEV / CVE] SOURCE: nvd.nist.gov | VELOCITY: 100/100 | SEVERITY: 80/100

CISA KEV: CVE-2025-32433 - Erlang Erlang/OTP SSH Server Missing Authentication for Critical Function Vulnerability

Erlang Erlang/OTP SSH server contains a missing authentication for critical function vulnerability. This could allow an attacker to execute arbitrary commands without valid credentials, potentially leading to unauthenticated remote code execution (RCE). By exploiting a flaw in how SSH protocol messages are handled, a malicious actor could gain unauthorized access to affected systems. This vulnerability could affect various products that implement Erlang/OTP SSH server, including—but not limited to

■ SECTION IX: VERIFIED PROOF-OF-CONCEPTS & RED TEAM REPOSITORIES

Exploit-DB, Packet Storm & MITRE ATT&CK; / ATLAS Threat Taxonomy

EXPLOIT WEAPONIZATION TIMELINES: The window between vulnerability publication and functional exploit automation has compressed to under 24 hours. Red team repositories release proof-of-concept scripts that require proactive protocol-level inspection rules before binary patches can be fully staged.

TECHNIQUE	NAME	TACTIC	MITIGATION DIRECTIVE
T1190	Exploit Public-Facing App	Initial Access	Deploy WAF rules & patch edge perimeter gateways
T1059	Command & Scripting Interpreter	Execution	Enforce PowerShell Constrained Language & container read-only roots
T1078	Valid Accounts & Token Theft	Defense Evasion	Mandate FIDO2 MFA & continuous conditional access re-evaluation
AML.T0054	LLM Prompt Injection	AI ATLAS	Input guardrails & strict tool parameter schema validation
AML.T0043	Model Weights Exfiltration	AI ATLAS	Encrypt model storage & enforce egress inspection on inference clusters

[VERIFIED PoC] SOURCE: [nvd.nist.gov](#) | VELOCITY: 100/100 | SEVERITY: 100/100
CISA KEV: CVE-2023-4966 - Citrix NetScaler ADC and NetScaler Gateway Buffer Overflow Vulnerability

Citrix NetScaler ADC and NetScaler Gateway contain a buffer overflow vulnerability that allows for sensitive information disclosure when configured as a Gateway (VPN virtual server, ICA Proxy, CVPN, RDP Proxy) or AAA virtual server. Required Action: Apply mitigations and kill all active and persistent sessions per vendor instructions [<https://www.netscaler.com/blog/news/cve-2023-4966-critical-security-update-now-available-for-netscaler-adc-and-netscaler-gateway/>] OR discontinue use of the produ

[VERIFIED PoC] SOURCE: [nvd.nist.gov](#) | VELOCITY: 100/100 | SEVERITY: 70/100
CISA KEV: CVE-2025-20352 - Cisco IOS and IOS XE Software SNMP Denial of Service and Remote Code Execution Vulnerability

Cisco IOS and IOS XE contains a stack-based buffer overflow vulnerability in the Simple Network Management Protocol (SNMP) subsystem that could allow for denial of service or remote code execution. A successful exploit could allow a low-privileged attacker to cause the affected system to reload, resulting in a DoS condition, or allow a high-privileged attacker to execute arbitrary code as the root user and obtain full control of the affected system. Required Action: Apply mitigations per vendor

[VERIFIED PoC] SOURCE: [rapid7.com](#) | VELOCITY: 100/100 | SEVERITY: 50/100
CVE-2026-63520: Microsoft SharePoint Remote Code Execution (FIXED)

OverviewRapid7 Labs conducted a zero-day research project against Microsoft SharePoint, resulting in the discovery of two new vulnerabilities that, when chained together, achieve unauthenticated remote code execution (RCE) against a vulnerable SharePoint server. Today, both Rapid7 and Microsoft are disclosing the second vulnerability in this chain, the RCE vulnerability CVE-2026-63520. The first vulnerability in the chain, CVE-2026-55040, was disclosed by Rapid7 and Microsoft last month.Our full

■■ SECTION X: 24-HOUR REMEDIATION PLAYBOOK & OPERATIONAL DIRECTIVES

Actionable Patching SLAs, AI Agent Runtime Hardening & Broadsheet Colophon

TIER	SLA	SCOPE & DIRECTIVES
P0 EMERGENCY	< 4 Hours	Apply vendor patches for active in-the-wild zero-days (CISA KEV). Isolate compromised host instances immediately.
P1 CRITICAL	< 24 Hours	Patch high-velocity CVEs (CVSS >= 8.5). Rotate service account credentials for exposed cloud providers.
P2 HIGH	< 72 Hours	Audit AI agent system prompts, update dependencies in container registries, and verify model SafeTensors.

TACTICAL AI & INFRASTRUCTURE HARDENING DIRECTIVES

1. **AI Agent Sandboxing:** Execute all LLM tool invocations in isolated gVisor/firecracker microVMs with strictly bounded egress and zero root capabilities.
2. **SafeTensors Deserialization Mandate:** Strictly reject untrusted PyTorch .bin/.pt pickle files across training and inference clusters.
3. **Perimeter Access Isolation:** Disallow public internet exposure of administrative ports (SSH, RDP, Kubernetes API, Ollama daemon).

COLOPHON & SENSOR METHODOLOGY: The Aether Guard — Global AI & Technology Gazette is compiled autonomously by the AetherGuard Intelligence Engine. Data is aggregated across 92 authoritative global sources including Hugging Face, GitHub Trending, arXiv, CISA, NVD, Exploit-DB, and sovereign CERTs. Neural NLP analyzers perform multi-language translation, algorithmic deduplication, and structured technical synthesis. All rights reserved.